

CONSTRAINTS ON THE EMERGENCE AND
INTERPRETABILITY OF LINGUISTIC CONVENTIONS

A DISSERTATION
SUBMITTED TO THE DEPARTMENT OF PSYCHOLOGY
AND THE COMMITTEE ON GRADUATE STUDIES
OF STANFORD UNIVERSITY
IN PARTIAL FULFILLMENT OF THE REQUIREMENTS
FOR THE DEGREE OF
DOCTOR OF PHILOSOPHY

Veronica Boyce

July 2025

© Copyright by Veronica Boyce 2025
All Rights Reserved

I certify that I have read this dissertation and that, in my opinion, it is fully adequate in scope and quality as a dissertation for the degree of Doctor of Philosophy.

(Michael C. Frank) Principal Adviser

I certify that I have read this dissertation and that, in my opinion, it is fully adequate in scope and quality as a dissertation for the degree of Doctor of Philosophy.

(Laura Gwilliams)

I certify that I have read this dissertation and that, in my opinion, it is fully adequate in scope and quality as a dissertation for the degree of Doctor of Philosophy.

(Judith Fan)

Approved for the Stanford University Committee on Graduate Studies

Abstract

Human communication takes place across many circumstances, and humans are adept at flexibly using language to achieve goals even in constrained conditions. The formation of linguistic conventions is often studied in iterated reference games, where over repeated reference to the same targets, a describer–matcher pair establishes partner-specific shorthand names for targets. In this dissertation, we expand the range of empirical evidence about when conventionalization occurs and what types of conventions are formed, all within the framework of iterated reference games. In Chapter 2, we ask which properties of a group’s interaction structure facilitate successful communication, using a repeated reference game paradigm where we manipulated several key constraints on the group’s interaction, including group size, the amount of feedback that matchers could give to describers, and the availability of peer interactions between matchers. In Chapter 3, we explore the developing pragmatic and referential communicative abilities of young children using a simplified, child-friendly paradigm. In Chapter 4, we consider the nature of conventions in iterated reference games, in particular to what extent the partner-specificity of these linguistic conventions causes them to be opaque to outsiders. Overall, we find that some signatures of convention formation can occur across a wider variety of circumstances than previously known, and that the descriptions used in iterated reference games are relatively comprehensible to outsiders.

Dedication

To those who supported me in this journey.

Acknowledgments

The path my PhD and my PhD years took was shaped by a number of people to whom I owe thanks for their support.

First, I want to thank my parents for supporting me throughout my entire life. They cultivated intellectualism and knowledge-seeking in me from a young age. They also showed me that going to graduate school (at Stanford) was just a thing one does; the Stanford part is mostly coincidence.

On the research side, I'm grateful to my mentors, collaborators, and lab mates. I'd like to thank my mentors and committee members (past and present) for their thoughtful feedback and support on projects.

I'm extremely grateful to Mike for all of his support and wisdom. I heard a lot of praise about him when I was choosing a graduate program, and I think it's all true and then some. I've learned a lot, both from direct pedagogy and from observing what Mike does. And not just about research and academia, but also (and this is I think especially by example) how to be a good person, balance life and work, and deal with the ups and downs of life in academia. (Sorry about writing somewhat sappy things Mike, but I really do appreciate your being there for me throughout my PhD.)

Mike has cultivated a really great community in LangCog. I'm grateful to all the members of LangCog (and the extended LangCog universe) who I've overlapped with. It's been a really great culture of people and, at least for me, the group has cultivated a great sense of belonging in spite of the diverse research topics.

I also owe thanks to my collaborators, and in particular, to the many research assistants who helped out collecting data for the developmental project. That project would not have been possible without them, and I'm grateful that they were willing to deal with preschoolers when I was not. I also owe thanks to Bobby for dealing

with all the RA and Bing and iPad logistics for that project.

Alvin and I started out in LangCog at the same time and have gone from colleagues to collaborators and friends. I'm grateful to Alvin for a variety of different interactions – pushing me to do some projects, being board game buddies, having an equal dedication to doing puzzle hunts (including at work during work hours), and giving feedback on this dissertation's introduction.

While often not people I know personally, some people who have causally shaped my PhD journey are those who have made and maintained the software packages, open science resources, and academic computing tools that I've used. I don't know how this PhD would have gone if I didn't have tidyverse and the Rmarkdown to LaTeX to pdf pipeline, but it surely would have been worse and more filled with paranoid checking of numbers.

I've been very lucky to live at Ithaka for four years of my PhD. Thanks to everyone I've overlapped with at Ithaka for making our shared house a home and keeping me fed and providing me with people to hang out with when I wanted to socialize. The feeling of coming home to dinner and people to have dinner with provided a sense of needed stability for me for the past four years.

To certain Ithakans, a deeper thanks is due.

Kelley and Isabel were both examples of how it's possible to do a Stanford PhD while being a real person who wanted boba, and didn't always want to work, and to do all of this without going too crazy. Your survivals of the final year of grad school provided evidence that it was doable.

Isaac and Daniel are in my mind pillars of Ithaka and cornerstones of my friend circle within Ithaka. I'm very glad for a certain pizza outing that led me to get to be their friends. Their senses of humor and willingness to humor my frequent desires for ice cream and gossip walks are appreciated.

Walks actually turns out to be a common thread in a number of my interactions – I might be afraid of forgetting specific individuals, but I think walk companions covers most potential omissions. So, to everyone who I went on walks (or more commonly “walkies”) with during my PhD, I'm appreciative of our walks and co-occurring conversations. The motivations for walks varied greatly, and the conversations had many tones and were useful in a range of situations.

Moving cross-country to start graduate school in 2020 was not an easy time for me, and I'm grateful to Andrea and the SGs for their important roles getting me through that rough patch. They have continued to be great supportive friends after that first year.

Finally, to Kevin and Ben, my “other significant other” and my “significant other”: each of you has been my go-to person for part of my PhD, and you've both done many walkies and coworkies with me, not to mention baking and board games and biking and everything else that doesn't start with the letter “b”. I love you both dearly.

Contents

Abstract	v
Dedication	vi
Acknowledgments	vii
1 Introduction: Computational psycholinguistic approaches applied to understanding convention formation in iterated reference games	1
1.1 Motivations and definitions	1
1.2 Iterated reference games as a model system	4
1.3 Frameworks	6
1.3.1 Efficiency	6
1.3.2 Bayesian pragmatics	9
1.3.3 Information Theory	12
1.4 Phenomena	13
1.4.1 Mentalizing and non-mentalizing approaches	14
1.4.2 Partner specificity	17
1.4.3 Audience design & effort sharing	19
1.4.4 Convention formation	22
1.4.5 Development of pragmatic communication skills	23
1.4.6 Incremental approaches to processing	25
1.4.7 Production constraints	27
1.5 Where does this leave us?	29
1.5.1 Efficiency	30
1.5.2 Bayesian pragmatics	31

1.5.3	Information theory	32
1.5.4	Summary	33
1.6	Outline of the dissertation	34
2	Interaction structure constrains the emergence of conventions in group communication	36
2.1	Introduction	36
2.2	Results	38
2.2.1	Overview of experiments	40
2.2.2	Smaller and higher-coherence groups are more accurate	41
2.2.3	Smaller and higher-coherence groups are more efficient	43
2.2.4	Larger groups make greater use of matcher contributions . . .	46
2.2.5	Descriptions converge faster in groups with thicker channels .	46
2.2.6	Games with thicker channels diverge from one another more quickly	50
2.3	Discussion	51
2.4	Materials and Methods	54
2.4.1	Participants	54
2.4.2	Materials	55
2.4.3	Procedure	55
2.4.4	Data pre-processing and exclusions	57
2.4.5	Modelling strategy	58
3	Preschoolers can form conventional pacts with each other to communicate about novel referents	60
3.1	Introduction	60
3.2	Experiment 1	63
3.2.1	Methods	63
3.2.2	Results	65
3.2.3	Discussion	69
3.3	Experiment 2	70
3.3.1	Methods	70
3.3.2	Results	71

3.4	Joint analysis	72
3.5	General discussion	74
4	Idiosyncratic but not opaque: Linguistic conventions formed in reference games are interpretable by naïve humans and vision–language models	76
4.1	Introduction	76
4.2	Task setup	80
4.2.1	Materials	80
4.2.2	Experimental procedure	80
4.2.3	Computational models	81
4.3	Experiment 1	82
4.3.1	Methods	83
4.3.2	Results and discussion	83
4.4	Experiment 2	84
4.4.1	Methods	85
4.4.2	Results	86
4.4.3	Model results	88
4.4.4	Discussion	89
4.5	Experiment 3	89
4.5.1	Methods	89
4.5.2	Results and discussion	91
4.6	General Discussion	94
	Conclusion	97
	Appendix A	104
	Feedback	104
	Semantic content of descriptions	104
	Additional game transcripts	107
	Recruitment and number of games	108
	Sensitivity analyses	111
	Matcher contributions	114

Reaction times	118
Group level variation in models	121
Additional measures of convergence	123
Distinctiveness of tangrams	124
Summaries of model outputs	125
Accuracy models	125
Reduction models	127
Primary reduction model	127
Extra reduction model	129
Log reduction model	130
Matcher reduction models	130
Initial utterance reduction model	132
Linguistic content models	134
Convergence within games: comparison to last round	134
Divergence across games	135
Divergence across tangrams	137
Convergence to next	138
Divergence from first	140
Exploratory Mega-analytic models	142
Sensitivity analysis models	144
Accuracy	145
Reduction	146
Convergence within games	148
Divergence across games	149

References	151
-------------------	------------

List of Tables

2.1	Examples from 6-player groups in Experiment 3 of successful descriptions for the same image across repetitions. Describers are indicated with an asterisk. More example descriptions are in Appendix Tables A1 and A2.	45
3.1	Example descriptions children successfully used to identify different target images in Experiment 2.	69
4.1	Cross-validated accuracies for classifiers. Standard deviations in accuracy across the 10 folds are shown in parentheses. Best performance within each model class is underlined, and best overall performance is bolded.	82

List of Figures

2.1	(A) Participants played a repeated reference game in groups of size 2 to 6. On each trial, a describer described the target image to the group of matchers. Each image appeared once per block for six blocks. (B) Experiments varied along 3 dimensions: Group size, group coherence, and matcher contributions. (C) Experiment 1 (pink) varied group size from 2 to 6 players while holding group coherence and matcher contributions constant. Experiment 2 (blue) held group size constant at 6 and manipulated the other dimensions. Experiment 3 (green) tested 4 corners of the space, crossing group size (2 vs. 6 players) with the thickness of interaction structure (high vs. low coherence and matcher contributions).	39
2.2	Behavioral results across all three experiments. (A-C). Matcher accuracy at selecting the target image. (D-F). Number of words produced by the describer each trial. For all, small dots are per game, per block means, and smooth lines are predictions from model fixed effects with 95% credible intervals. Y-axes are truncated, and a few outliers points are not visible.	42

2.3	Example utterances describing the shown tangram figure produced by two 3-player games in Experiment 1. To measure convergence within a game (blue), we measured the cosine similarity between SBERT embeddings of descriptions and the embedding of the round 6 utterance (taken to be the convention). Higher cosine similarity indicates more similar meaning. To measure divergence between games (green), we measured the similarity between embeddings of utterances from the same round across games.	47
2.4	Language similarity results measured with pairwise cosine similarity between embeddings of two utterances. (A-C). Convergence of descriptions within games as measured by similarity between an utterance from block 1-5 to the block 6 utterance in the same game for the same image. (D-F). Divergence of descriptions across games as measured by the similarity between two utterances produced for the same image by different groups in the same block. For all, small dots are per game, per block means, and smooth lines are predictions from model fixed effects with 95% credible intervals. Y-axes are truncated, and a few outliers points are not visible.	49
3.1	Experimental setup and procedure. Panel A shows the experimental setup. Panel B shows the 4 possible targets; names for targets are for cross-reference with later figures only. Panel C shows the procedure for Experiment 1; within critical blocks targets were ordered randomly. Panel D shows the procedure for Experiment 2.	62
3.2	Children’s accuracy at selecting the correct target over time. Error bars are bootstrapped 95% CIs with a linear trend line overlaid. . .	66
3.3	Length of description produced by the teller each trial. Grey dots are individual data points, colored dots are per trial means with bootstrapped 95% CIs.	67
3.4	Semantic similarity between pairs of descriptions from different sources. Dots are means and lines are bootstrapped 95% CIs.	68

3.5	Semantic similarity between descriptions from earlier blocks (1-3) and the last block in Experiment 2. Heavy dots are means with bootstrapped 95% CIs; light dots are individual values.	73
4.1	Experimental setup. Naïve matchers read transcripts from trials in reference games from Boyce et al. (2024) and selected which image they thought was being described. Matchers recieved bonus payments for correct selections.	80
4.2	Correlation between human accuracy and CLIP-MLP probability of target in Experiment 1. Small points are individual descriptions, colored by decile of CLIP-MLP probability, large points and error bars are the bootstrapped mean and 95% CI across descriptions for each decile.	83
4.3	Accuracies for naïve human matchers and the CLIP-MLP model for Experiments 2a and 2b, grouped by the source of the referential description. Facets are the communication thickness of the original game and x-axis is when in the game the transcript came from. Point estimates and 95% CrI are predictions from the fixed effects of logistic and beta regressions. Bootstrapped mean accuracy from the original matchers is included as a ceiling, and random chance as a baseline. .	84
4.4	Accuracies for naïve human matchers and the CLIP-MLP model for Experiments 2a and 2b, split out by target image. Point estimates and 95% CI are predictions from the fixed effects and by-tangram random effects of logistic and beta regressions, bootstrapped across conditions. Bootstrapped mean accuracy from the original matchers is included as a ceiling, and random chance as a baseline.	85
4.5	Accuracies for Experiment 3. Error bars are bootstrapped 95% CIs.	90
4.6	Schematic of informativity predictors used in the incremental reading time analysis.	92
4.7	Estimates and 95% credible intervals for coefficients in the incremental processing model. Word-level predictors are shown in panel A, colored by spillover position. Condition effects are shown in panel B.	93

Chapter 1

Introduction: Computational psycholinguistic approaches applied to understanding convention formation in iterated reference games

1.1 Motivations and definitions

The use of language is a feature that distinguishes humans from other species. Language enables coordination between individuals, enforcing social norms via gossip, and communicating complex and abstract ideas, which in turn allows for the cultural accumulation of technologies beyond what any individual could invent. Paired with the technological innovation of writing, language allows for communicating across long distances of space and time.

Humans and other social species can communicate in a variety of modalities, but linguistic structures make language more expressive than other species' communication systems. Humans learn languages in early childhood and then are able to use the shared set of symbol-meaning mappings and combinatorial rules in order to communicate with other humans. Even young children realize that language is a very

powerful system for transmitting ideas from mind to mind, as they demonstrate when they use it to make demands upon their caregivers. As a literate human, language is ubiquitous to (at least my) daily life, but the more one considers language, the more there is to explain and the more impressive our linguistic capabilities seem.

The word “language” can refer to a suite of related concepts all of which are targets of study in related domains. There’s the abstract idea of a specific language, for instance, English, as a set of grammatical rules and word meanings. This is an abstract idealization as attempts to specify formal rules that cleanly generate the accepted utterances of a language run up against variability across communities. From a statistical perspective, one can aggregate across variability to study language as it is used by a certain community, not a static entity but something that evolves slowly over time. However, this is also an abstraction, as individuals each have separate grammars that may not fully align even within a community, so meanings must be treated as somewhat probabilistic or distributional. The idea of an individual’s grammar is an abstraction over the process of language learning from individual exposures to speech acts (and texts) and the psychological processes and neural architectures that produce and process language.

Separate from the scale at which we study language, there are other lenses that language is studied through. Language is preserved in artifacts of text and recordings that are studied separately from the human context and intentionality that produced them. Language can also be studied as a capacity for learning a system for expressing ideas and then expressing and understanding ideas in a system. From an information-theoretic lens, language can be viewed as a particular communication channel, subject to the same bounds as any other signal-transmission method. From the opposite perspective, languages often have substantial social and cultural relevance to their users. Generalizing across purposes, language is a tool that humans use to communicate with each other, and we can use both the combinatorial structure of language and pragmatic variations on the structure to accomplish a variety of communicative goals via the meaning conveyed and other features that are implied.

These different aspects of language reveal the complexities of studying language and lead to disagreements around what is studied in linguistics or the “language

sciences”. In the dissertation, I focus primarily on language as a tool for communication, but I will also consider other aspects of language including artifacts of language, language as a shared set of expectations, and psychological processing of language.

Another consideration when studying language use is what counts as language? Some scholarly traditions concentrate on conversational face-to-face language use for communicative purposes, and exclude derivatives like wordplay, phonetic codes, or language models as non-linguistic things that happen to use linguistic symbols. In this work, I do not assign *a priori* primacy to any particular parts or modalities of language. It may be that some types of language use are easier to process than others, or that some uses of language rely on domain general inferences external to the language faculties, but these are empirical questions whose results I will not assume.

While not as polysemous as the word *language*, *pragmatics* also has a number of partially overlapping meanings. When I use *pragmatics* I mean pragmatics expansively as something along the lines of meaning in context or context as a component to inferring meaning (Bohn et al., 2019).

This dissertation focuses on how humans use language to communicate around the edges of expressivity – how we extend our lexica to accommodate communicative needs, and how we coordinate with each other in these atypical uses of language. These linguistic stretches can be ad hoc and temporary, but if they fill appropriate communicative niches, what are initially atypical usages can spread and lead to the adoption of new words, new meanings, and other forms of language change.

As previewed above in discussion of the many angles from which language can be studied, there are trade-offs between studying language in more realistic situations and studying language in more tractable and controlled environments. In this work, I focus on language use *in context* because it is both tractable enough that experiments can usefully inform models and theories and analogous enough to real-world situations to have some explanatory power.

In the rest of the introduction, I first motivate the use of iterated reference games as a model system. I then discuss three main approaches from psycholinguistics that can be applied to convention formation in iterated reference games: efficiency, Bayesian pragmatics, and information theory. Then I review findings from the conversational and iterated reference game literatures and consider how these findings

might be subsumed into larger psycholinguistic frameworks. Finally, I close with avenues for further inquiry and preview the few avenues that are addressed in the rest of the dissertation.

1.2 Iterated reference games as a model system

Iterated reference games are a commonly used paradigm for studying adaptive language use. A *reference game* consist of a set of people communicating with each other about a set of potential referents. One person, the *describer*, has a specific designated target (usually pre-identified as the target by the experimenter) that they seek to communicate to the *matcher* (or matchers). This act of reference requires coordination between describer and matcher on a symbol–meaning alignment. *Iterated* reference games involve repeated reference to the same targets within the same group of people, so that a certain conversational history builds up that could shape the construction and understanding of later symbols. Crucially, this task uses referents that are somehow difficult to name in context, usually by a combination of low codeability or high similarity to each other.

As described, this task is agnostic to communication modality, and iterated reference games are sometimes conducted using gestures or drawing (Bohn & Frank, 2019; Garrod et al., 2007; Hawkins, Sano, et al., 2023). Here, I focus on how this task interacts with the rich baseline of symbol–meaning mappings afforded by natural language.

People sometimes use words in creative ways to express meanings. For instance, even if a person doesn’t know the name for an object or is having a tip-of-the-tongue event, they can still communicate about the object by using a more circuitous label. If this communicative need occurs repeatedly, an initially ad hoc referring expression can calcify into a settled name for the object. Communicative acts serve a variety of purposes, but a simple, yet central one is reference: connecting a symbol with a meaning to pick out a target. This is a subgoal needed for more complex acts – it’s hard to tell a story or debate someone without agreement on what different words or phrases refer to.

Combining these observations – that language can be adapted to fill communicative needs, that repeated need leads to changes in referring form, and that reference is an easy to elicit communicative act that is central to other communicative acts – suggests the utility of iterated reference games.

Iterated reference games are a commonly used task. The earliest studies we are aware of are from the 1960s (Krauss & Bricker, 1967; Krauss & Weinheimer, 1964, 1966), and iterated reference games seem to have become more popular starting in the 1980s (Clark & Wilkes-Gibbs, 1986). Iterated reference games have been used to study a wide variety of communicative and linguistic phenomena. The primary targets of study are coordination and the formation of conventions or conceptual pacts (Brennan & Clark, 1996; Dahan, 2025; Eliav et al., 2023; Ghaleb et al., 2024; Hawkins, Frank, et al., 2020; Metzing & Brennan, 2003, among others). Some studies focus particularly on iterated reference games as a way of studying collaboration in natural conversations (Clark & Wilkes-Gibbs, 1986; Dahan, 2024; Schober & Clark, 1989; Wilkes-Gibbs & Clark, 1992). Other studies focus on other aspects of coordination, including the negotiation of referring expressions (Bangerter et al., 2020; Knutsen & Le Bigot, 2018) and the role of feedback (Fox Tree & Clark, 2013; Krauss & Glucksberg, 1977; Krauss & Weinheimer, 1966). Iterated reference games with multiple partners are used to study how speakers design their utterances for their specific partners (Horton & Gerrig, 2002, 2005; Turner & Knutsen, 2021; S. O. Yoon & Brown-Schmidt, 2014, 2019a, 2019b) and how speakers might generalize patterns of social conventions across interlocutors (Hawkins, Goodman, et al., 2020; Hawkins et al., 2021). Iterated reference games are used as the dialogue condition in studies on the differences between monologue and dialogue (H. P. Branigan et al., 2011; Fox Tree, 1999; Murfitt & McAllister, 2001; Tolins et al., 2017). Iterated reference games have also been used to study conversation and memory (McKinley et al., 2017; Vanlangendonck et al., 2018), model cultural conflicts between merging teams (Weber & Camerer, 2003), examine speech-gesture trade-offs (de Ruiter et al., 2012) and study the effects of alcohol on social dynamics (Garrison et al., 2024).

In addition to studies on typical populations, iterated reference games are also used to study the abilities of special populations. Iterated reference games are used with children and teenagers (H. P. Branigan et al., 2016; Leung et al., 2024; Turner

& Knutsen, 2021) and with older adults (Horton & Spieler, 2007; Hupet et al., 1993; Lysander & Horton, 2012) to study cognitive and linguistic abilities across the lifespan. Iterated reference games are used to study types of memory in people with amnesia (Duff et al., 2005; Duff, Gupta, et al., 2011; Duff, Warren, et al., 2011; S. O. Yoon et al., 2017) and to study various populations with neurological and psychological conditions (Champagne-Lavau et al., 2009; Chantraine et al., 1998; Gupta et al., 2012; Hengst, 2003).

1.3 Frameworks

Just as one might hope that the results observed in iterated reference games generalize to more realistic situations, one also might hope that the results of iterated reference games could be explained by broad coverage theories of language use. I review three broad approaches that are common within computational psycholinguistics: efficiency, Bayesian pragmatics, and information theory.

1.3.1 Efficiency

One unifying framework in psycholinguistics is efficiency, the idea that language and language use is under pressure to support efficient communication by maximizing the ratio of relevant information transmitted to effort. Efficiency is a high level framework, and as such, requires linking assumptions to enable testing against data. Comparisons between attested languages and counter-factual languages bound what types of linking hypotheses are needed for efficiency to be a parsimonious framework.

Efficiency arises from trade-offs between communicative expressivity and some combination of learnability and ease of production (Kirby et al., 2015; Piantadosi et al., 2012). At any time, a language is thus in a dynamic equilibrium between the usage pressure to communicate many ideas (pushing complexity of language) and learnability and ease-of-production pressures (pushing simplicity of language).

Features of language are distributed more closely to the Pareto frontier of the expressivity-simplicity trade-off than would be expected by chance. A well-known observation is that word frequencies follow a power-law distribution, which Zipf (1949) explains in terms of a “principle of least effort”, where words that are more frequent

are also shorter. This type of power-law distribution is common across domains and generated by a variety of processes, so while well known, it is not strong evidence for efficiency as a driver of linguistic distributions (Piantadosi, 2014). Stronger evidence comes from the lexical partitioning of sub-domains such as color, number, and kinship terms, where the distribution of systems falls on the Pareto frontier between complexity (number of terms) and informativity (how many bits each term provides) (Gibson et al., 2019; Kemp et al., 2018; Zaslavsky et al., 2018). Syntactic features of language such as harmonic word order and dependency length also appear to be optimized for increased expressivity with minimized processing effort (Gibson et al., 2019; J. Hawkins, 1995).

Arguments for efficiency are based on the language artifacts of grammars and transcripts, but efficiency pressures act on language use as a process, not language as a static code (Gibson et al., 2019). Efficiency imposes a joint constraint on the entire communicative process: minimizing the total time and effort involved in transmitting an idea in one person’s mind to a sufficiently close idea in another person’s mind. While shorter utterances are usually considered more efficient, utterance length in syllables or clock-time is not actually under pressure, but rather the time to produce or parse is under pressure. While length and processing time often correlate, under situations where a verbose alternative is more accessible and comprehensible, it would be favored compared to a succinct but hard-to-find wording.

As with other high level frameworks, the parsimony of the efficiency framework is entirely based on its auxiliary linking theories (Gershman, 2018). A way of judging the explanatory adequacy of efficiency is by determining what auxiliary assumptions are needed to connect a given phenomena under the efficiency umbrella and assessing how reasonable those assumptions are.

As an illustration, I consider what it would take for efficiency to explain ambiguity and redundancy in referring expressions.

Ambiguity and redundancy in referring expressions

How do we reconcile the idea that language is efficient with violations of efficiency in language use (“look at the yellow banana”) and the ubiquity of ambiguous utterances?

Terms such as “redundant” and “over-informative” are commonly used to describe

situations where people produce modifiers that do not restrict the extension of a noun phrase, e.g. “blue cup” when only one cup is salient. Here, the substring “cup” is an alternative utterance that could uniquely identify the target (Rubio-Fernandez et al., 2021). People do produce these non-restricting modifiers in reference tasks, especially for color, but this behavior seems to run counter to the idea of efficient language use.

Formalizing claims of redundancy or over-informativity requires a definition of what would be minimally informative. Defining what is minimally informative requires commitment to a fully specified semantic-pragmatic system, at least for the set of alternative utterances. In the general case, this is difficult to formally specify: for instance, if specificity implicatures are within the option space, are those calculated before or after informativeness is measured (Bergen et al., 2016)? In limited cases, one can sidestep issues of formal theory by empirically measuring the information content of different utterances (Degen et al., 2020), but this does not scale up well in part due to issues with compositionality.

The flip side of “redundancy” is ambiguity: many, many utterances are ambiguous. In general, strong contextual factors render ambiguity a non-issue (Piantadosi et al., 2012), but then the physical and social context must be considered when judging whether language is efficient or redundant. Determining what is efficient language use in-context requires not just analyzing phrases and their alternatives, but also processing and production costs, which may be highly contingent on contextual factors and conversational history.

The idea that utterances should have “just enough” information (Grice, 1989) has inspired a wealth of empirical research into what utterances people produce and what utterances people comprehend. By comparing these two halves of language use, one could determine how calibrated production and comprehension are. Tight alignment between production and comprehension would support an efficiency framework with minimal linking assumptions. In contrast, discrepancies between what is produced and what is comprehensible would require auxiliary assumptions, such as bottlenecks on production, or communicative goals other than informativity. One would then have to judge whether these auxiliary assumptions are parsimonious.

1.3.2 Bayesian pragmatics

A broad coverage framework in cognitive science is modeling humans as Bayesian agents making resource rational choices. Linguistic actions can be modeled using Bayesian pragmatic models, where (boundedly) rational Bayesian agents reason about each other to determine the meanings of utterances.

A prominent model within Bayesian pragmatics is the Rational Speech Acts (RSA) framework, an information-theoretic, computational framework for making quantitative predictions about pragmatic inferences in contexts (Frank & Goodman, 2012; Goodman & Frank, 2016). The RSA family of models involves two interlocutors recursively reasoning about how the other would produce or interpret utterances, grounding out in a listener (or speaker) who behaves in a pre-specified “literal” way.

Computational frameworks such as RSA provide a way to factor together different trade-offs and determine their relative weights in a model (Goodman & Frank, 2016). A softmax is taken over the scores of the options to produce a distribution of interpretations and utterances, with some parameters such as the degree of optimality fit based on data. RSA models often use only one or two levels of recursion, which produces a reasonable fit to human experimental judgments, consistent with findings that most people reason pragmatically at a low recursion depth (Franke & Degen, 2016).

RSA models have been used to predict ad hoc pragmatics, conventionalized pragmatic implicatures, and children’s word learning (Bergen et al., 2016; Bohn et al., 2022; Degen et al., 2020; Goodman & Stuhlmüller, 2013; Ohmer et al., 2022). RSA models generally include a utility or informativity term representing how well an utterance resolves uncertainty in favor of the target referent. Other commonly included factors include a saliency prior and some cost function penalizing longer and more complex utterances (Goodman & Frank, 2016). Some models go beyond informativity, incorporating options to infer the question-under-discussion (Kao, 2014; Qing et al., 2016) or balance informativity with politeness (E. J. Yoon et al., 2020).

A full RSA model would incorporate all of these components and infer their weights. However, for tractability, usually only those features that are considered relevant to the domain of interest are included in the model. Because of the flexible framework, RSA can be used to model many sources and levels of uncertainty,

and then integrate out that uncertainty to make predictions, but also update on the sources of uncertainty in response to input.

The Challenge of Semantics

A significant challenge of RSA models is how to define the “literal” listener or speaker. In toy domains, the set of possible utterances is small, and it is feasible to enumerate a set of meanings (Bergen et al., 2016; Frank & Goodman, 2012; Goodman & Stuhlmüller, 2013). A soft or continuous semantics can be used to represent that some dimensions of meaning might be more strongly informative than others, which supports the prediction of patterns of “redundant” color adjective use in referring expressions (Degen et al., 2020).

Soft semantics can conflict with compositionality. Soft semantics can exist without explicit compositionality, provided there is a way to determine match between expression and meaning. At a small scale, this can be determined empirically and fully hand-specified. One option for scaling-up would be building neural network systems that embed utterances into semantic space and then taking distances to the target meaning, although this requires making assumptions about what “literal” meaning is. The system might have implicit compositionality, but there are no guarantees about the relationships between the embeddings of “red”, “apple”, and “red apple” in these systems. This may be more severe for expressions like “green apple on the red box” and “red apple on the green box” where syntactic relations may not be fully preserved by sentence-level embeddings (Erk, 2022). Depending on the domain, embedding systems may or may not have satisfactory match to human semantic judgments.

Alternatively, one could design an explicitly compositional system, where meanings are defined on components and there is a process for mapping the meanings of pieces into the meanings of whole phrases. The process could be deterministic rules or an algorithm or neural model implementation. In formal semantics, a system of logical relations and lambda calculus is used, which can assign that “red apple” is the intersection of “red” and “apple”. However, this formal system would then predict that “apple” is equally true for members of “red apple” and “blue apple”, which doesn’t match people’s judgments due to typicality effects (Degen et al., 2020). The challenge to both a holistic soft semantics and a compositional approach is that language both

has strong compositional structure and also cases where the same structure composes meanings differently.

One reason to want both compositionality and soft semantics is to be able to learn from interactions. Given that an expression–meaning pairing was used successfully (or unsuccessfully), one should be able to locally update the lexicon to reflect this additional information, which requires a way to appropriately propagate the updated meaning to component pieces and related phrases.

There is not yet a satisfactory solution for initial or literal meanings in non-toy domains: some situations can be handled by empirically measuring likelihoods in an exhaustive ways, but this holistic approach is not compatible with incremental RSA (Cohn-Gordon et al., 2018) or larger sets of utterances that require compositionality to be defined. Some work has used grounded neural language models for alternative generation and likelihood measures, which has scaling potential in domains that have the right grounded datasets, but not all domains have these datasets (Cohn-Gordon et al., 2018; Monroe & Potts, 2015; White et al., 2020). In order to extend RSA towards more realistic and open-ended scenarios, an important issue is to find a meaning representation that (even at a computational level) will appropriately support pragmatic reasoning.

RSA with history: CHAI

One key obstacle in going from RSA-style models to explaining iterated reference games is the need to model the process of transitioning from a successful long, multi-part description to a different, and usually shorter, one. One RSA-style model that can qualitatively generate this pattern of behavior is CHAI, a framework to bridge different levels of convention formation (Hawkins, Franke, et al., 2023).

CHAI uses parameterized hierarchical uncertainty over lexica so that initially, there is variability in how words will be interpreted, which a speaker must integrate over. At the outset, listeners’ lexica are not fully known which represents potential background differences that may be shared by a sub-population and differences from communication history that may be individual-specific. After receiving feedback on the success of the utterance, the new information can be propagated to both the individual and group lexical representations, resulting in a new prior for future rounds.

CHAI can qualitatively account for various convention-formation phenomena in toy systems (Hawkins, Franke, et al., 2023).

CHAI’s explanation for reduction is consistent with the ambiguity arguments of Piantadosi et al. (2012): initially the context is not sufficiently constraining to resolve the ambiguity of a short utterance, so multiple descriptive pieces are needed to triangulate the meaning. As conversational history accrues, the context is sufficient to disambiguate the shorter description.

In toy models of interlocutors playing a reference game with soft semantics, initial utterances use multiple properties to collectively increase the degree of certainty in the target. This successful reference then shapes the prior over word meanings, until the degree of certainty afforded by only one word is sufficient to pick out a referent.

The same approach CHAI uses to generalize from individuals to groups and populations could also be used to generalize across stimuli. Then, a successful reference, like “man with arms up”, would not only update the meaning of that phrase and its components, but could also lend weight to the over-hypothesis about using body-part-based descriptions, and thus update the speaker’s beliefs about the listener’s semantics for the words “head” and “leg”.

As with other RSA models, a primary challenge is how to scale the model and add open-ended semantics to allow for comparison against experimental data using natural language and natural image stimuli.

1.3.3 Information Theory

Information theory is a formalism used in computational psycholinguistics to link abstract ideas to concrete measures and predictions. Information theoretic concepts such as entropy, surprisal, and rate distortion theory have all been key to quantitative theories of language.

Information theory provides a link from functionalist accounts of language usage to the artifacts of language systems, by asking what the constraints are that would render the observed language systems an optimal solution to the functional question of communication (Futrell & Hahn, 2022). One advantage of information theory is that it provides ways to calculate measures of complexity in a relatively theory-neutral way (Futrell & Hahn, 2022). Measures derived from information theory such

as surprisal, conditional entropy, and mutual information have been very productive in language science both for unifying aspects of language typology (ex. dependency locality) (Futrell & Levy, 2017) and for explaining patterns in language processing (ex. surprisal theory and entropy reduction) (Hale, 2001, 2016; Levy, 2008). The ideas of rate distortion and messages compressing meaning through limited channels have been key to testing efficiency claims related to syntax, word order, and grammatical markers (Futrell & Levy, 2017; Mollica et al., 2021). Information theory based theoretical proofs and toy models have been applied to domains such as language evolution (Plotkin & Nowak, 2000). RSA formulations are closely related to information-theoretic solutions to joint optimization between utility and effort (Zaslavsky et al., 2020).

One difficulty with applying information theoretic measures is the need for a probability distribution over possible meanings and messages; thus, similarly to other frameworks, a potential sticking point is choosing an appropriate distributional language model or datasets to capture the baseline language.

1.4 Phenomena

The communication and conversation literature provides descriptive characterizations of how people communicate in naturalistic and experimental settings. For iterated reference games, the empirical literature provides a fairly consistent and reliable set of results around reduction, convention, and partner-specificity. While there are consistently observable trends in the conversation literature broadly and in iterated reference games specifically, the effects tend to be characterized with verbal theories. Research thus far has not rigorously identified necessary and sufficient conditions or dose-response relationships between circumstances (or experimental settings) and the size of the effects. Quantitative accounts for the phenomena are also lacking.

The verbal theories often fail to make explicit predictions about what would happen under conditions other than those tested or what experiments would adjudicate between theories, which makes testing these theories difficult. Verbal theories often explain their results narrowly as the product of specific situations that are being

tested, rather than as part of broader phenomena. As I review the empirical findings and associated theories, I will consider the available evidence, problematize the verbal theories, and consider how broader-coverage approaches from computational psycholinguistics might be used to more formally model or subsume the empirical results.

1.4.1 Mentalizing and non-mentalizing approaches

One historically important distinction in the study of conversation is whether agents track other agents’ internal states. While an intuitively appealing distinction, it would be difficult to measure. One could build, for instance, model-based and model-free algorithms and then compare which better captures an observed phenomenon. This is reminiscent of fights between associationist versus rule-based accounts in linguistics, which don’t seem as relevant now that symbolic and non-symbolic approaches are seen more as different possible implementations than as different theories. We know now that what looks rule-based could be instantiated in a non-symbolic, associationist way, and a rational action symbolic approach could look heuristic under performance load.

Mentalizing tradition and “common ground”

The “mentalizing” tradition treats humans as representing other humans as agents with internal states that include knowledge and goals. Within this umbrella category, approaches vary in how these representations are implemented, how information gets added or modified, what exactly is tracked, and when representations (versus heuristics) are used.

A commonly occurring phrase in the mentalizing tradition is the term “common ground” used to refer to knowledge that two agents share. In some cases, it is used in a pre-theoretic way to mean roughly “things you think another person will understand and won’t be surprised if you reference” (Garrison et al., 2024; Leung et al., 2024). For instance, Hanna et al. (2003) defines common ground as the “mutual knowledge, beliefs, and assumptions” held by the interlocutors. This meaning is roughly comparably to “givenness” in other domains (Fay et al., 2010). However,

“common ground” also has a more specific and theoretically loaded meaning, originating from the privileged versus mutual versus common knowledge framework (Clark & Wilkes-Gibbs, 1986). Under this usage, “common ground” is defined via infinite recursion in knowing that the other person knows that the first person knows that.... This technical usage is used in formal semantics where entities referred to with a definite article must be introduced to common ground via accommodation, which could itself be seen as a kludge to explain how people get away with making statements whose presuppositions aren’t known to be shared (Horton & Keysar, 1996; Pickering & Garrod, 2004). Empirical work suggests that humans generally don’t engage in more than a couple layers of recursion in their pragmatic reasoning (Franke & Degen, 2016). Thus, it is generally not necessary to distinguish knowledge types at deeper recursion levels than mutual knowledge that both people know to be mutual.

Even if we accept that common ground has grown away from its formal semantic roots to become a synonym of “mutual knowledge”, there is still a question of how agents determine what is mutually known. A common answer is to lay out a deterministic set of rules that supposedly characterize how an agent knows a piece of information is mutually known (Brown-Schmidt, 2012; Clark & Wilkes-Gibbs, 1986; Horton & Keysar, 1996). Some of these rules are generally useful, such as the idea that items in shared visual presence are mutually known, but there are always going to be exceptions and uncertainty, which calls into question the utility of a deterministic flow-chart approach.

The idea of mutual knowledge can be easily subsumed within a Bayesian pragmatics framework. While the “mentalizing” approach places special emphasis on common ground, it is not clear whether there is a reason to model shared knowledge differently from cases where one models another’s knowledge that one does not personally have. For instance, in deciding who to learn from, someone might determine that another agent knows what’s behind a door and then use this information to, for instance, interpret a quantifier like “some” (Goodman & Stuhlmüller, 2013).

Another advantage of the Bayesian approach is that one can avoid deterministic rules and instead reason under uncertainty. Modeling uncertainty could be useful for explaining hedging behavior, and joint updating could be used to model situations

where a piece of knowledge was originally thought to be shared, but new information decreases this likelihood. The framework of flexible inference and probabilistic knowledge states is a powerful tool.

Interactive alignment theory

In contrast to mentalizing approaches, “interactive alignment theory” claims to explain how people can successfully collaborate on reference tasks without reasoning about each others’ mental states (Gandolfi et al., 2022; Pickering & Garrod, 2004). Accounts of naïve ego-centrism suggest that mentalizing is difficult, which provides motivation for an account of collaborative behavior without reference to mental states. Recent work suggests that, contra naïve ego-centrism, humans reason socially about each other readily and from a young age (Rakoczy, 2022). If mentalizing is not in fact prohibitively costly, the need for a non-mentalizing theory is much reduced. However, some scientists have believed that explicit reasoning about other minds was too computationally expensive and slow to adequately explain behavior, and Pickering & Garrod (2004) tried to account for observed alignment between interlocutors by extending ideas from structural priming.

Structural priming occurs when processing a grammatical structure (for instance, a double object construction) increases the likelihood of expressing an idea using that structure rather than an alternative (producing a double object rather than a prepositional object construction) (Tooley & Traxler, 2010). In production, structural priming is observed even when there is not lexical overlap (ex. different verbs), but is increased when there is lexical overlap. Structural priming in comprehension has been observed only with the same lexical head item (Tooley & Traxler, 2010). The mechanisms of structural priming are an active area of theorizing, without any consensus. The tendency to reuse constructions, styles, or words to express ideas can be observed in non-experimental, naturalistic settings (Doyle et al., 2016; Doyle & Frank, 2016; Pickering & Ferreira, 2008).

Pickering & Garrod (2004) claims that the alignment occurs via “priming” and is “resource-free and automatic”, without providing an explanation of how this would work in terms of memory or processing. Unfortunately, as expressed, interactive alignment theory doesn’t provide an algorithmic or mechanistic theory. While inspired

from structural priming, interaction alignment goes much further. Structural priming makes narrow claims around why people tend to say “the ball that’s red” after their partner says “the block that’s blue”, whereas interactive alignment claims there is trickle-up and trickle-down effects across linguistic levels that reinforce alignment and align people’s “situation models”. This is a strong claim that is not well supported by a testable algorithm, and does not make risky predictions. The phenomena of interest, repetition, could be explained by some sort of “priming”, or it could be explained by explicit reasoning or production-side ease-of-use constraints (MacDonald, 2013).

Another issue with interactive alignment theory is that it does not account for one of the key features of iterated reference games: convergence to conventionalized referring expressions that evolve and shorten. Interactive alignment does not have an explanation for this pattern of change, rather than mere repetition. Interactive alignment theory generally uses symmetric designs where participants both have partial information, so it may be built as an explanation for a different type of result than the asymmetric describer/matcher designs common to the mentalizing tradition.

While conversational alignment is a worthwhile area of study, the theories of conversational alignment do not provide a satisfying account of the phenomena of iterated reference games. The concern that full inference is not cognitively feasible may be more fruitfully explored using approximate inference in a resource-rational framework or by placing information theoretic limits on the capacity to represent features of individual partners.

1.4.2 Partner specificity

In iterated reference games, describers typically revert to longer descriptions when they interact with a new matcher. This change in behavior is described as “partner-specificity”, as the conventional names developed with one partner are specific to that partnership and not transferred over to other new partners (Brennan & Clark, 1996; Hawkins et al., 2021; Metzing & Brennan, 2003). Partner-specificity is more generally also used to refer to how different pairs diverge to different names for the targets (Hawkins, Frank, et al., 2020). Partner specificity assumes that speakers are mentally representing their partners and their partners’ knowledge states, but is agnostic to the exact form of representation, and so could be compatible with heuristic

or distributed representations in addition to symbolic.

Empirical evidence from experiments where one director talks with multiple partners suggests that people do “partial pooling” over their partners (Hawkins, Goodman, et al., 2020; S. O. Yoon & Brown-Schmidt, 2019a). A speaker A may vary their descriptions when talking to partner B versus partner C, but A talking with B will be more like A talking with C than D talking to E, as there is some shared generalization that A makes across partners B and C. Paired with the tendency to reduce descriptions over time, partial pooling leads to a jagged pattern of reference length: when switching to a new partner, speakers use longer utterances, but not as long as their initial utterance with their first partner (S. O. Yoon & Brown-Schmidt, 2019a).

A reversion to more elaborate descriptions for new matchers could be efficient if the context from the initial utterances is requisite to understanding the shorter versions. If the elaborated and shortened forms of the concepts (ex “arms up, like a goalie diving for a ball” versus “goalie”) are considered synonyms (in the same way chimp and chimpanzee are synonyms), then the switching could be explained under a uniform information density model, provided that the longer form occurs in a context when it is more surprising and informative (Mahowald et al., 2013).

Partial pooling and partner specificity can be instantiated in a hierarchical model of Bayesian pragmatics, such as CHAI or lexical uncertainty models (Bergen et al., 2016; Hawkins, Franke, et al., 2023; Potts et al., 2015; Schuster & Degen, 2020). These models provide theoretic coverage, but testing these frameworks against openlexicon natural language data is not yet computationally tractable.

From an information-theoretic lens, the mutual information would include the conversation history, the current utterance, and the visual context of targets. Within a partnership, the conversation history can serve as an encoder that allows later conversational elements to use a more efficient coding to map between language and referent. Partner specificity effects could arise when, conditioned on conversation history, the target and convention predict each other, but target and convention do not predict each other given no history or a different history.

1.4.3 Audience design & effort sharing

Audience design is the idea that speakers appear sensitive to the knowledge state of their listener and say things that are easy for the listeners to comprehend. Broadly understood, audience design is occurring when a person speaks differently to a 5-year old than to an adult, or to a colleague than to a lay person. However, despite the intuition that audience design occurs in everyday interactions, it is difficult to determine if intentional audience design occurs in iterated reference games.

The phrase “audience design” sometimes implies intention on the part of the speaker (Horton & Gerrig, 2002, 2005), and sometimes is used when utterances are constructed based on what’s easy for the speaker, and listener ease is a side effect (Horton & Keysar, 1996; MacDonald, 2013; Rogers et al., 2013). For instance, giving an elaborated description to a naïve listener could be an intentional design choice (that contrasts with a shorter describer for a non-naïve listener), but a speaker could also start descriptions with given material because it is more accessible to them, and only incidentally convenient for the listener. Audience design is sometimes used even more broadly to refer to production choices that make utterances easier to comprehend (V. S. Ferreira, 2019). In a very broad sense, efficiency pressures for uniform information density could be construed as audience design.

Speakers and listeners often share recent context, find the same things salient, and linguistically what is easier to produce is often easier to process. Thus, disentangling ease of production from ease of comprehension requires careful experimental settings where these two processes diverge (F. Ferreira, 2004).

Audience design relates to the broader topic of how interlocutors split the communicative burden with one another. Depending on the task and the communication modality, speakers and listeners have different options for how to balance the communicative load (Clark & Wilkes-Gibbs, 1986; Fay et al., 2010; Fox Tree & Clark, 2013). For instance, a listener could describe what options they see or otherwise prompt the speaker. Load splitting may vary based on the capacities of the interlocutors (ex. a speaker might craft their utterances more when talking to a child versus an adult) and the capacities of the channels (ex. speakers may use different approaches if listeners can interrupt).

Multi-way conversations complicate the theories of audience design and partner

specificity by introducing a larger audience of more partners. Two main questions are whether speakers “aim low” or “aim high” in balancing the needs of the listeners and whether speakers track listeners as individuals or an aggregate (S. O. Yoon & Brown-Schmidt, 2017, 2019a).

Empirical results indicate that speakers are sensitive to the individual knowledge states of listeners. Describers seem to track the knowledge states of different matchers independently when they alternate between addressing two matchers (S. O. Yoon & Brown-Schmidt, 2017). When simultaneously addressing a group of matchers with disparate knowledge levels, speakers give descriptions that are fairly similar to what they give naïve listeners (S. O. Yoon & Brown-Schmidt, 2017). However, there is also some gradient, where descriptions vary based on the proportion of naïve listeners in groups of up to 4 addressees (S. O. Yoon & Brown-Schmidt, 2019a). With pairs of listeners who differ both in their background knowledge and their visual context, speakers produce descriptions that are sensitive to the interaction between background knowledge and distinctiveness of the target in context (S. O. Yoon & Brown-Schmidt, 2019b). Speakers in group contexts sometimes produce utterances to support multiple listeners, such as referring to a target with both the established name and an elaborated description (S. O. Yoon & Brown-Schmidt, 2019a). Research on balancing multiple listeners has not extended to larger groups, so it is unknown how the tracking of partner’s knowledge states might degrade in groups of more than 5 people.

Larger group communication is instead studied more in network structures, where (non-linguistic) convention formation has been studied on networks of up to 50 people (Guilbeault et al., 2021). In referring to targets sampled from a continuous space, large groups tended to develop fairly consistent category boundaries across independent networks, while small networks had more idiosyncratic category boundaries. This work suggests there may be group-size and group/network-structure dependencies on what sorts of conventions arise, but network games on traditional iterated reference games have focused on small networks of four people (Hawkins, Goodman, et al., 2020; Hawkins et al., 2021).

Partner specificity and audience design accounts implicitly assume that the speaker’s

behavior is close to optimal. Generally, humans are fairly pragmatic in their communication, so we might expect speakers to be well-calibrated on what listeners will understand. Speaker behavior certainly reflects the notion that different descriptions are appropriate depending on the listener’s knowledge state and their shared conversation history.

However, direct tests of whether descriptions are optimally tailored require some sort of counterfactual. One option is to test how well naïve matchers do given descriptions from different circumstances. Naïve matchers do better the more their observation history resembles that of the original conversation – when hearing descriptions in order instead of in reverse order (Murfitt & McAllister, 2001), when listening to the entire game instead of starting in the third round (Schober & Clark, 1989), and when seeing yoked trials from a single game rather than trials sampled across 10 games (Hawkins, Sano, et al., 2023). Descriptions produced in dialogue are easier to understand than those produced in monologue (Murfitt & McAllister, 2001). The gradient of performance suggests that some changes in descriptions may be adaptive to listeners who share more background. However, naïve matchers still have far above chance performance indicating that descriptions are not exclusively understandable by the knowledgeable listeners they were intended for. To the extent that new people can understand later utterances without context, a counterfactual more rapid adoption of short forms could achieve a certain level of accuracy just as well as the observed pattern. Thus, from a purely audience design listener-focused perspective, speakers are clearly not efficient. However, we do not know what the optimal utterance sequences would be for listeners.

Before we are forced to conclude that humans are mis-calibrated in their linguistic utterances, we can consider some potential auxiliary hypotheses that could render observed behavior efficient. One possibility is that the communicative bottleneck is firmly on the production side: speakers are doing the best they can to rapidly produce descriptions, and producing more succinct descriptions earlier would be too effortful for speakers. Another possibility is that social norms and task dynamics dissuade the use of short form descriptions even if they are likely to be understood: perhaps the initial descriptions are more polite, or perhaps people are (overly) cautious in their efficiency/accuracy trade-offs. Both of these hypotheses could be tested empirically.

Bayesian pragmatic models could be extended to cover communicating with multiple listeners at once. Rather than just communicate pragmatically with one person, balancing multiple people’s needs would require balancing different goals. A lexical uncertainty Bayesian pragmatics model could be used to capture the naïve listener behavior in eavesdropper scenarios. Effort splitting could in theory be represented by changing RSA parameters such as depth of reasoning, starting lexical semantics, or cost functions.

1.4.4 Convention formation

A key observation from iterated reference games is that dyads tend to form shared “conventions” (also called “conversational pacts”) about how to refer to the initially ambiguous targets. These conventions tend to be partner- and context-specific: changes in the speaker, audience members, or context can all license the use of a new description (Ibarra & Tanenhaus, 2016; Metzing & Brennan, 2003; S. O. Yoon & Brown-Schmidt, 2014). Absent these changes, pairs tend to stick near the semantic space of successful descriptions.

Convention formation and conceptual pacts can occur on different granularities of semantic alignment. Partners could align on the lexical level, by both calling a figure “ballerina”. They could align on conceptualizing the figure as a ballet dancer with a tutu, manifesting in descriptions with semantic, but not lexical, overlap such as “ballerina” and “dancing in a tutu”. Partners could agree on a general paradigm for how to describe figures, such as in terms as humans in different postures. Horton & Gerrig (2002) distinguish between “lexical entrainment” when the same words are reused, and “conceptual similarity” when there is broader similarity that does not repeat the same words. These levels often co-occur, but in order to computationally quantify the phenomenon of convention formation, we have to make choices about how to operationalize the similarity: is it word overlap, lemma overlap, distance in semantic space, or something else?

In addition to convention formation, repeated references also tend to change from descriptions of a set of features to more holistic labels. Typically, in convention formation, holistic or analogic descriptions are the ones that stick (Clark & Wilkes-Gibbs,

1986), but almost any type of label can sometimes crystallize into the convention, including ones that pick up on low-level or meta-level features. The range of successful options makes explaining convention formation difficult, and suggests that there is strong path-dependency in the evolution of reduced utterances, potentially influenced by hard to quantify factors such as relationships and humor value.

An utterance’s length and meaning are not *a priori* related, but these two features tend to correlate in iterated reference games (Hawkins, Frank, et al., 2020). The reduction of utterances is sometimes used as a measurement proxy for the semantic changes (Clark & Wilkes-Gibbs, 1986; Garrison et al., 2024; Hawkins, Franke, et al., 2023), and Bovet et al. (2023) considers it a measure of common ground. Convention and reduction are also conflated with partner-specificity, presumably again because of co-occurrence (Clark & Wilkes-Gibbs, 1986; Wilkes-Gibbs & Clark, 1992). Better tools for measuring each of these constructs more directly and independently will be needed to address the theoretical issues around how interrelated these concepts are. It remains an empirical question whether the shortening of utterances, convention formation, and partner-specificity of descriptions are inseparable or merely occur together in the experimental paradigms considered in the literature.

From an information-theoretic perspective, convention formation is convergence towards an encoding that is more efficient at representing the specific set of target meanings than the baseline language. These meanings are not commonly needed and so start with low codeability, but over local usage, a (temporary) warp of the meanings can be achieved that is locally more optimized to express the target meanings.

1.4.5 Development of pragmatic communication skills

One perspective on linguistic phenomena is to consider their developmental trajectory to understand when the phenomena arises and how it changes across development. Iterated references games have been studied only occasionally with children.

Glucksberg et al. (1966) suggested that 4-5-year-old preschoolers struggle with child-child referential communication. In their paradigm, one child was given an ordered set of 6 blocks and tasked with describing the image on each block so their partner could pick out the corresponding block from the partner’s unordered set of blocks. As children described and selected blocks, they stacked them on pegs,

and the ordering of blocks on pegs was used to assess children’s accuracy. While 4-5-year-old children succeeded on practice trials with familiar shapes and visual access to each other’s blocks, children failed on critical trials where the blocks had abstract drawings and there was no visual access. Even after multiple rounds with the same images, children were not able to correctly order the blocks. Glucksberg et al. (1966) attributed children’s communicative failures to their production of ego-centric descriptions that did not account for the other child’s perspective. Similar experiments with older children indicated a gradual improvement through adolescence both for initial accuracy and for the increase in accuracy across repetitions. Still, even the 9th grade sample was noticeably worse than the adult college student sample (Glucksberg & Krauss, 1967; Krauss & Glucksberg, 1969).

Since Glucksberg et al. (1966), few studies have revisited the question of whether children can successfully communicate in an iterated reference game. H. P. Branigan et al. (2016) found that 8-10-year-olds exhibited adult-like patterns of increasing accuracy, increasing speed, and shorter descriptions across repetitions, but the children’s accuracy was still far below adult performance and highly variable between dyads. In another study, 4-6-year-olds successfully used conventionalized gestures with their partners in an iterated reference game where children could only use gestures to communicate (Bohn & Frank, 2019). In a recent study, 4-8-year-old children succeeded at playing an iterated reference game with a parent using a simple, child-friendly tablet-based task (Leung et al., 2024). In this task, even 4-year-olds had an initial accuracy above 80%, which rose to above 90% in later repetitions (Leung et al., 2024).

In developmental work, it is crucial to assess whether paradigms are measuring the abilities of interest or whether task demands are masking children’s nascent abilities. Across different tasks in developmental science, carefully designed studies that use methods that require fewer additional cognitive resources have been revealing evidence of children having earlier emerging abilities than previously documented. As the task demands in Glucksberg et al. (1966) are substantial, their negative results could reflect either children’s ego-centrism or children’s difficulty with the task due to limited domain general cognitive capacity.

How could psycholinguistics frameworks add to our understanding of developmental pragmatics? If we had a computational model of adult performance, one could frame the question of development in terms of what changes or restrictions to the model would approximate child behavior. In an RSA-type model, children’s data might be best approximated with a lower optimization parameter or sampling fewer alternatives. A model of child pragmatics might also be modeled with a smaller and differently initialized lexicon. One question would be whether the optimal update size after each interaction in a CHAI model would be larger or smaller when fitting children’s data. The RSA modeling formalism suggests different ways that children and adults might differ from each other, that may be testable even outside the framework.

From an efficiency perspective, one question is whether children have the same trade-off between expressivity and effort as adults do. Cognitive capacity constraints or reduced executive control may make some utterances much more effortful for children, which could shift the entire curve to favor less-expressive and less-effortful utterances. Child – adult interactions, where one participant has much more cognitive capacity, may have yet other trade-off curves.

From an information theory lens, communication channels involving children could be modeled as extra noisy or having a more limited channel capacity (Yurovsky & Frank, 2017).

1.4.6 Incremental approaches to processing

Comprehension and production are both processes that unfold over time as meanings are converted to and from visual or acoustic signals. One open question in comprehension is how early in processing contextual effects such as pragmatic implicatures or speaker identity are integrated.

A common approach to studying comprehension in real time is via visual world eye-tracking where looks are interpreted as belief that the visual target is the intended referent. Visual world paradigms are used to investigate whether there is an initial stage of interpretation that differs from the final resolution. For instance, the literal meaning of an utterance might be available sooner than pragmatic or contextualized interpretations. Different sources of information that influence processing at different times are sometimes referred to as “top-down” for the use of context and “bottom-up”

for just the primary signal.

The relative balance and time course of top-down and bottom-up influences is a long-standing topic of interest in psycholinguistics and neurolinguistics (Gwilliams & Davis, 2022; Horton & Gerrig, 2005; Horton & Keysar, 1996; Tanenhaus et al., 1995). One debate in particular concerns the role of ego-centrism in the early stages of linguistic processing. In Egocentric Anchoring and Adjustment theory, Keysar et al. (2000) argues for an initially egocentric perspective on the basis that people often initially look at objects that are good matches to a description even if the object is not mutually visible to the speaker (see also Horton & Keysar, 1996; Keysar et al., 2003). This theory posits a two-stage model with an initial completely egocentric stage. This theory contrasts with the more broadly-held idea that one of the sources of information for linguistic processing is egocentric perspective but it can be incorporated in parallel to theory of mind information (Rubio-Fernández, 2008). Egocentric Anchoring and Adjustment rests on two dubious linking hypotheses: 1) if people consider an interlocutor’s perspective, their prior is that the interlocutor only refers to mutually known objects; and 2) looks to an object indicate that it is considered a potential referent (eye-tracking data is widely interpreted this way, but there is evidence that this proxy is only approximate, Degen et al., 2021). Studies supporting Egocentric Anchoring and Adjustment always have the egocentric-only lure be extremely salient and by far the best match to the spoken description.

The counterpoint presented by Hanna et al. (2003) is a constraint-based theory where many factors can play into comprehension to varying extents, and on differing time courses. These studies use egocentric lures that are equivalently good to the shared visibility target and find evidence for early perspective taking, which is confirmed by other experiments in different paradigms (Hanna et al., 2003; Nadig & Sedivy, 2002).

The larger picture of language processing and neuroscience strongly suggests that contextual information (at least from linguistic context) is used to predict upcoming linguistic material, and thus that “top-down” information can affect even early stages of linguistic processing such as phoneme identification and lexical processing (Gwilliams et al., 2023; Gwilliams & Davis, 2022). Predictive coding is a theory with good coverage of neural evidence in many domains, including language, with language

processing regions seemingly sensitive to both lexical and structural information in predicting upcoming words (Shain et al., 2020). Questions remain around whether earlier behavioral readouts are more influenced by perceptual information and less by pragmatics and whether cognitive load influences the time course of integrating different information sources.

1.4.7 Production constraints

Language production is less well characterized than language processing because of the increased difficulty in studying language production experimentally. There is agreement that production requires conceptualization, formulation, and articulation, but that is about all that is agreed upon (Bürki, 2018). The phonological realization of words that are more frequent or usually appear in more predictable contexts is reduced relative to words that usually convey more information (Bürki, 2018). Phonological reduction is consistent with both production facilitation or audience design as speaker and listener language statistics are so similar their effects are hard to distinguish.

Beyond single words, there are open questions about what the planning looks like, how the time courses of different stages of production may overlap, and what size and types of units are stored in memory (Bürki, 2018). One approach is interference studies that can determine how far ahead different sentence elements are planned, but studying utterance planning is challenging (Momma et al., 2016). The mechanics of production are uncertain even in situations with less context-sensitivity than iterated reference games.

As with behavioral measures, neuroscience measures are also used more for processing than for production, and within production, articulatory processes are better studied than higher-level planning processes. fMRI studies indicate that language production activates the same regions as language processing, but these regions are especially active during phrase-structure building, suggesting that production may have higher costs than comprehension (Hu et al., 2022). The increased load from production is consistent with accounts that suggest that utterance design (and language evolution) are substantially shaped by biases to ease production. MacDonald (2013) points out three production biases: easy first, plan reuse, and reduce interference that could explain some patterns of utterance production, including in reference

games.

One approach to identifying speaker-side difficulties in iterated reference games is to track disfluencies (S. O. Yoon & Brown-Schmidt, 2014). Intuitively, producing descriptions of abstract figures is difficult and cognitive load may lead to non-optimal productions (ex using “square head” as part of a description when it is not diagnostic) that are reminiscent of young children’s tendency to describe objects in similar ways regardless of what level of detail is needed to distinguish them in context (Leung et al., 2024).

Time-course locked measures such as eye-tracking and disfluency analysis could be used to understand speaker’s utterance planning. Carefully controlled comparisons to monologue conditions could also separate out speaker-side constraints from audience design. Varying time pressure or adding pre-articulation planning time could also distinguish what behaviors are the result of satisficing under limited resources and what behaviors are designed to be optimal (Horton & Keysar, 1996). As larger questions around production indicate, understanding language production in reference games is difficult and requires very careful and clever experiment designs.

Production-side constraints are posited to drive linguistic efficiency, but efficiency frameworks do not provide detailed predictions about what speakers would or wouldn’t say. In general, RSA operates on a level of abstraction that only considers entire utterances. However, given that most iterated reference games allow the speaker to keep talking, one could in theory adapt incremental RSA (Cohn-Gordon et al., 2018) to explain potential order effects around what parts of descriptions are produced first.

One information-theoretic computational-level model of incremental language production comes from Futrell (2023). This model applies the Rate Distortion Theory of Control to production, suggesting a trade-off between incremental steps that are good for the communicative goal (ex. being informative) with what is easy to produce, that is what is more automatic and predictable. Futrell (2023)’s model can qualitatively explain production ordering preferences, the use of filled pauses, and good-enough word production.

1.5 Where does this leave us?

A core theoretical question is whether the observed patterns of reduction, convention, and partner-specificity are due to specific mechanisms that only operate in a limited scope of iterated-reference-game-like situations or whether these results are explainable by broader coverage theories of pragmatic language use.

The phenomena observed in iterated reference games do not appear under other circumstances where people use pragmatic reasoning, such as single-shot reference games. One option is that iterated reference game phenomena are driven by mechanisms that only apply to narrow circumstances including iterated reference games. For instance, theories of conversational coordination with common ground and a backchannel do not specify an exact scope, but are clearly not trying to account for all pragmatic referential language productions. In some cases such as theories specific to dyads, the implicit scope of a verbal theory is smaller than the range of circumstances where the phenomena is observed.

Alternatively, the same mechanisms that account for pragmatics broadly could also account for the phenomena in iterated reference games. Thus, iterated reference game phenomena like reduction would arise from how general mechanisms play out when targets have low codeability and reference is repeated to the same target with the same partner. Certain combinations of the parameters of a more general model (what the targets are, what the contextualized lexicon is, what the communication channel is) would lead to the emergence of the observed phenomena of partner-specific convention formation.

The iterated reference game literature has primarily focused on descriptive results and developing small-scope theories. I am more interested in what parameterizations and auxiliary hypotheses are needed to interface between broad-scope theories and the results of iterated reference games. I believe our default hypothesis should be that convention formation emerges from the same mechanisms that account for all the rest of language use. The overall goal is to figure out the mechanisms of language use, and iterated reference games are a useful model system for studying highly contextual language production and comprehension.

Having surveyed the phenomena and small-scope theories of iterated reference games, I now return to the three broad-coverage frameworks and synthesize how they

might apply and what the key open questions are.

1.5.1 Efficiency

It is easy judge an utterance to be inefficient if there is a shorter alternative that is both easier to produce and easier to comprehend. In the general case, however, efficiency is very hard to cache out in specific predictions. What’s efficient for an utterance in isolation may not be efficient when considered over an entire life of language use.

The formation of conversational pacts is informally characterized as efficient (Clark & Wilkes-Gibbs, 1986; Hawkins, Frank, et al., 2020). Here, “efficiency” is capturing that shorter expressions are *more* efficient, but a given shortened expression could still be mis-calibrated to a particular situation. It could be too short, leading to misunderstanding, or still longer than it needs to be for effective communication. Mis-calibration does not necessarily contradict claims about partner specificity of reduction, depending on auxiliary hypotheses. For example, communicators could intend to be optimally calibrated but run into production difficulties, masking their competence. Current claims of efficiency in conversation do not connect directly with the formal claims of optimal efficiency in the language evolution literature, and it is unknown whether speakers and listeners are calibrated or not.

One approach is to test the idea that speakers are reducing “as fast as possible” for listeners; that is, both that extra words in an utterance are adding utility (adding information, adding certainty), and that the context of earlier utterances (and their repetition) is necessary for understanding later utterances (i.e. reducing faster would negatively impact accuracy). This could be tested with models and information theory applied to utterances and by testing new matchers on manipulated samples or orders of utterances.

If speakers are not reducing “as fast as possible”, one should consider the auxiliary hypothesis that deviations are the result of production-side demands. If speakers have the hard task, then their utterances may be driven primarily by ease of planning and not information-theoretic optimality. As discussed above, pinpointing speaker-side demands is difficult, but could be approached with monologue studies, eye-tracking, and varying amounts of planning time prior to utterance initiation.

One challenge to studying naturalistic human behavior is that the stated task goal and participant’s utility function are not the same. People have exogenous norms about communication that influence how they converse with each other even in experimental settings.

Efficient language systems results from a process of cultural evolution where pressures apply at individual moments of language use. Efficiency does not require that every utterance be on the efficiency frontier, merely that there is enough pressure towards expressivity and ease overall that the system shifts in this direction. Thus, whether or not people are efficient at iterated reference games or not is relevant for explaining the behavior in reference games, but cannot invalidate ideas of efficiency because people’s behavior on this task is not shaping language use on a large scale.

1.5.2 Bayesian pragmatics

Bayesian pragmatics is a promising large-scale framework for explaining linguistic behavior in context and incorporating history and non-linguistic context into a model of utterance production and comprehension. The main challenge in applying Bayesian pragmatics is scaling up the models to handle the natural language produced in reference games.

An issue with Bayesian pragmatic models is how many free parameters are needed to fit the data, and whether one can actually achieve a parsimonious broad coverage model, or whether one just ends up with a number of related models each tuned to a specific dataset. The need to avoid overfitting to specific datasets is a broad issue in modeling where the goal is to fit a generative or cognitive model that could represent more than the one dataset.

One approach is to extend the existing RSA models to fit a wider range of experimental designs. Group communication, either in the form of one speaker - multiple listeners (S. O. Yoon & Brown-Schmidt, 2014, 2019a, 2019b) or network structures (Hawkins, Goodman, et al., 2020), poses a challenge as there is not yet a non-dyadic extension of RSA that can do partial pooling across listeners and jointly optimize for the responses of multiple interlocutors. It should be possible to design extensions to CHAI to accommodate different group structures.

Connecting the existing small-scale or limited domain models to the goal of predicting the patterns of full-scale natural language requires solving problems related to the open vocabulary. How can one set baseline literal meanings for any utterance someone might produce, and how should meanings update over repetitions? With toy models, one can do compositional semantics and update based on the compositionality. One could ground semantics in some neural model such as a VLM, but that doesn't provide a straightforward way of propagating updates down compositional units. One of the bottlenecks for getting a robust test of Bayesian pragmatic models is doing the NLP and ML engineering to assign and update meanings on an open space of alternatives.

CHAI, while not a full explanation for the patterns of real-world data, provides a framework that explains why reduction would be optimal. It remains to be seen whether RSA-style models can predict the slope of reduction and the content of words that stay versus drop, but attempts to build formal pragmatic models of these phenomena would push forward our understanding of pragmatics and communication. Even without a fully cached out model, the Bayesian pragmatic approach can inform how we think about experiments and results.

1.5.3 Information theory

Unlike the frameworks of efficiency and Bayesian pragmatics, information theory is more of a tool. A robust finding in iterated reference games is that the utterances get shorter over repetition, but what's actually going on with the language? In order to quantify the changes in iterated reference game language that occur over time, one needs metrics to characterize the changing language. For instance, one might theorize that words that are more informative are more likely to be retained and become the conventions. To test this, one would need a measure of "information" derived from computational models of language and stimuli.

Information theory could be used to formalize ideas of common ground by considering the joint information between the conversation history (as the common ground) and the current utterance. To measure this at scale would require models to approximate the information, but this may be possible with the current era of language and vision models.

For normal psycholinguistic texts, information theoretic predictors like surprisal or entropy reduction have substantial predictive value on processing difficulty as indexed by measures of reading time and neural signals (Smith & Levy, 2013; Wilcox et al., 2021). For a task like the iterated reference game, the measure of interest is task-relevant information that reduces the entropy over the target object. Surprisal and task-relevant information are not always the same, such as in gibberish which is high surprisal, but low information.

An iterated reference game could be thought of as a process where language is adapting to the specific task of optimally coding the target options. An optimal code would only need 3.5 bits to express each target when there are 12 options. Many more bits than this of natural language are used in early rounds – suggesting that, initially, standard natural language is not a good distributional fit for this task. Some of this language is conveying task-relevant information in a wordy way, while other parts may be not-task relevant at all (ex: describing features that do not vary across targets). Over time, people drop some of the non-relevant words, but they also develop a more efficient code for representing the task-relevant information. This formulation makes some testable predictions about the entropy of distributions over utterances at different points during the game.

1.5.4 Summary

In this introduction, I’ve highlighted some avenues for future work that is needed to test how well broad-coverage psycholinguistic approaches explain iterated reference game results. Some of these directions are completely outside the reference domain – better understandings of language planning and production generally, or natural language processing models that respect compositionality. Other directions involve computational and modeling directions as discussed above.

One other direction is to better characterize the **explanandum**: what exactly are the phenomena of interest? While there have been many studies of iterated reference games, some theoretically relevant aspects of the experimental space remain under-explored. The broad coverage approaches expect a mostly continuous relationship between circumstantial knobs and the outcome measures, which is easiest to assess with a more even sampling of experimental space. From a modeling perspective,

having more data and a wider range of data also reduces some concerns that we are over-fitting to a very narrow circumstance.

One necessity to substantial progress in the iterated reference game domain is a better empirical and descriptive understanding of the phenomena of interest. How do circumstantial knobs (group size, communication modality, stimuli, etc.) push around the extent to which we observe reduction, convention formation, and partner specificity? We don't have to finely map the entire landscape, but some sense of the geography of experimental space is necessary (Almaatouq et al., 2022).

More variable experimental circumstances are also needed to determine how closely related the different phenomena of interest are. In the narrow regions of experimental space that have been sampled so far, reduction, semantic alignment, and partner-specificity all pattern together, to the point where some metrics are used as proxies for each other. A greater range of experimental situations is needed to determine how well these measures separate versus correlate in different situations, which will in turn clarify whether these features should be linked mechanistically in theories.

1.6 Outline of the dissertation

In this dissertation, I expand the empirical foundation of iterated reference games in three directions, with the aim of addressing some theoretical questions and providing data for computational modeling in the future. I also apply some computational models to survey coarse-grained patterns of changes in semantic similarities.

In Chapter 2, I address the generalization patterns of the key phenomena in iterated reference games, by examining the trajectories of convention formation in games of varying small group sizes and varying communication channel capacities. This expands the empirical foundation to mimic more of the different circumstances in which communication occurs in the real world, and allows for a test of how entangled the different reduction phenomena really are.

In Chapter 3, I turn to the developmental timeline of the communication and coordination skills embedded in iterated reference tasks. I re-visit historical claims that children are too egocentric to produce successful ad-hoc referential expressions, which is fairly incompatible with more recent work showing nascent pragmatic language use

and communication skills in preschoolers. This different context again provides a test of how entangled different observed reduction phenomena really are and provides insights into how language use fits into broader cognitive and linguistic development.

Finally, in Chapter 4, I measure the level of opacity in the referential expressions produced in iterated reference games (using the corpus of reference games from Chapter 2). Taking inspiration from lexical uncertainty models, I use both naïve human matchers and computational proxies to determine how far from “baseline semantics” referential expressions are at different stages of convention formation. These patterns are a test of at least the strongest versions of efficiency and audience design claims and give us a better understanding of what sort of partner-specificity is at play.

In the last chapter, I conclude with a discussion of where this leaves us and what directions seem most promising from here.

Chapter 2

Interaction structure constrains the emergence of conventions in group communication

2.1 Introduction

Much of human social life revolves around communication in groups. At school, teachers address large classrooms of children (Cazden, 1988); at home, we chat with groups of friends and family members over dinner (Tannen, 2005); and at work, we attend meetings with colleagues and managers (Caplow, 1957; Zack, 1993). Such settings present considerable challenges that do not arise in the purely two-party (dyadic) settings typically studied in psychology (H. Branigan, 2006; Ginzburg & Fernandez, 2005; Traum, 2004). For example, producers need to account for the fact that different comprehenders in the group may have different mental states or levels of background understanding (Fox Tree & Clark, 2013; Horton & Gerrig, 2002, 2005; Weber & Camerer, 2003; S. O. Yoon & Brown-Schmidt, 2014, 2019a), while comprehenders must account for the fact that utterances are not necessarily tailored to them (Carletta et al., 1998; Cohn-Gordon et al., 2019; Fay et al., 2000; Metzing & Brennan, 2003; Rogers et al., 2013; Tolins & Fox Tree, 2015; S. O. Yoon & Brown-Schmidt, 2014). What enables producers and comprehenders to nevertheless overcome these challenges and navigate multi-party settings with relative ease?

One promising set of hypotheses centers on the group’s *interaction structure*, the set of constraints placed on the group’s shared communication channel. Many different aspects of interaction structure have been implicated in the effectiveness of dyadic communication, including the availability and quality of concurrent feedback (Krauss & Bricker, 1967; Krauss & Weinheimer, 1966; Kraut et al., 1982), the bandwidth of the communication modality (Dewhirst, 1971; Krauss & Glucksberg, 1977), and the group’s access to a shared workspace (Clark & Krych, 2004; Garrod et al., 2007). Yet larger groups introduce qualitatively different dimensions of interaction structure, leading to a large but often inconsistent body of findings even for these well-understood factors (Hiltz et al., 1986; Swaab et al., 2011). While communication is generally expected to deteriorate as groups get larger (MacMillan et al., 2004; Seaman & Basili, 1997), the structural “thickness” of the feedback channel may slow such deterioration (Ahern, 1994; Parisi & Brungart, 2005).

In this paper, we develop an experimental paradigm for evaluating the relative contribution of these factors: a *multi-party repeated reference game*. The ability to distinguish one particular entity from other possible entities, known as *reference*, is one of the most primitive and ubiquitous functions of communication. Reference games (Lewis, 1969; Wittgenstein, 1953) have been widely used to study dyadic communication under controlled conditions in the lab. They provide a clear metric of communicative effectiveness: how many words are required before a matcher successfully chooses a target image from a context of distractors? *Repeated* reference games, where the same target images appear multiple times in succession, were introduced to examine how interlocutors establish shared reference in the absence of conventional labels (Clark & Wilkes-Gibbs, 1986; Krauss & Weinheimer, 1964). At the beginning of the game, long and costly descriptions are typically required to succeed. A key finding, however, is that dyads become increasingly efficient over the course of interaction. Fewer words are required to achieve the same accuracy, but referring expressions also become more impenetrable to outsiders (Schober & Clark, 1989; Wilkes-Gibbs & Clark, 1992). The evolution of referring expressions over repetitions shows the characteristic dynamics of conventions: *stability*, or convergence on labels within a group, and *arbitrariness*, or divergence to different across groups, suggesting that dyads leverage their shared communication history to coordinate on

expectations about how to label the target images (Hawkins, Frank, et al., 2020).

In principle, repeated reference games provide a strong operationalization of communicative effectiveness for the problem of multi-party communication: describers must simultaneously achieve shared reference with multiple matchers. However, empirically studying multi-party communication raises a number of difficulties in practice. A much larger pool of participants must be recruited to achieve sufficient power at the relevant unit of analysis – the group – spanning a very high-dimensional space of possible parameter settings (Almaatouq et al., 2022). We address this problem by drawing on recent technical advances that have made it newly possible to achieve such samples using interactive web-based platforms (Almaatouq et al., 2021; Haber et al., 2019; Hawkins, Frank, et al., 2020). Repeated reference games in web-based platforms have previously replicated earlier results from face-to-face studies, although people produce fewer words in text modalities than oral modalities (Hawkins, Frank, et al., 2020). The text-based chat modalities arguably more closely resemble the interfaces used by modern teams who increasingly communicate through group text threads or popular platforms like Slack or Discord.

We leverage our platform to explore effects of group size and interaction channel thickness in a series of three experiments. While we find that small groups reliably converge on group-specific “shorthand” regardless of the interaction structure, larger groups require thicker channels – richer conversational feedback among members – to achieve the same degree of coherence. Thus, increasing group size alone does not impede communication; rather, larger groups may require stronger social and linguistic cues to establish common ground among all members. More broadly, our work suggests that studying communication in larger groups is necessary to unveil critical aspects of interaction structure that have not been evident in typical dyadic settings.

2.2 Results

We recruited 1319 participants through Prolific, an online crowd-sourcing platform. Participants were organized into 313 groups of size two to six for a communication game (Figure 2.2A). On each trial, everyone in the group was shown a gallery of 12




Figure 2.1: (A) Participants played a repeated reference game in groups of size 2 to 6. On each trial, a describer described the target image to the group of matchers. Each image appeared once per block for six blocks. (B) Experiments varied along 3 dimensions: Group size, group coherence, and matcher contributions. (C) Experiment 1 (pink) varied group size from 2 to 6 players while holding group coherence and matcher contributions constant. Experiment 2 (blue) held group size constant at 6 and manipulated the other dimensions. Experiment 3 (green) tested 4 corners of the space, crossing group size (2 vs. 6 players) with the thickness of interaction structure (high vs. low coherence and matcher contributions).

tangram images (Clark & Wilkes-Gibbs, 1986; Hawkins, Frank, et al., 2020; Ji et

al., 2022). One player was designated the *describer* and the others were designated the *matchers*. The describer was asked to use a chat box interface to describe a privately indicated *target* image. After all matchers guessed which of the 12 images was the target, they received task feedback and proceeded to the next trial. The game consisted of 72 trials structured into 6 repetition blocks, where each image appeared as the target exactly once per block.

We manipulated the interaction structure of this game across 11 distinct conditions in 3 distinct pre-registered experiments (Figure 2.2B). We systematically sampled points along four dimensions parameterizing different aspects of the interaction space. We manipulated *group size* (ranging from two to six), *role stability* (whether or not participants took turns in the describer role), richness of *task feedback* (whether or not matchers were able to see each other’s responses), and richness of the *matcher contributions* (whether matchers were able to freely respond through a chatbox or could only use emojis; Figure 2.2C). Other factors, such as the set of stimuli and background knowledge about one’s partners, were held constant across games.

2.2.1 Overview of experiments

Experiment 1 began by investigating how performance scaled with group size. Based on prior qualitative work, we predicted that larger groups face a more challenging coordination problem. We continuously varied the number of players from 2 to 6 while keeping other factors constant. For these conditions, the describer role rotated after each block, so that all players had at least one turn as describer. Matchers had access to an unrestricted chat box, but only received binary task feedback about whether their individual selection was correct without revealing others’ selections or the intended target.

Experiment 2 focused on the most challenging 6-player groups and explored the role of interaction structure. Each condition in Experiment 2 varied one aspect of the experiment relative to the Experiment 1 6-player baseline. We tried two variants that we expected to increase group coherence and improve performance, and a third variant we expected to interfere with the ability to establish mutual understanding and thus impede performance. In the first variant, we maintained the same describer throughout rather than a rotating describer, such that the same individual has the

opportunity to aggregate feedback across trials and track which matchers are struggling with which targets. In the second variant, we gave the group of matchers full feedback about what every other member of the group had selected, and we showed the intended target. In the third variant, we changed how matchers could make contributions to the group. In contrast to prior experiments, where matchers could contribute freely to the chat; here, we limited matchers to sending four discrete emojis (green check, thinking face, red x, and laughing-crying face) that could convey simple valence and level of comprehension, but not any referential content.

Experiment 3 crossed the extremes of group size from experiment 1 (2 vs. 6 people) with the extremes of group interactions from Experiment 2 (*thick* vs. *thin* interaction structure). In the *thick* condition, we maintained a consistent describer, gave all matchers full task feedback, and allowed them to freely use a chat box. In the *thin* condition, we forced the describer to rotate on each block, restricted feedback to their own binary correctness, and restricted matcher contributions to the four emojis. Note that the 2-player thick game most closely resembles the design of classic repeated reference games (Clark & Wilkes-Gibbs, 1986).

2.2.2 Smaller and higher-coherence groups are more accurate

Our first set of hypotheses focused on group performance: how accurately and efficiently groups were able to perform the referential task. We characterize group performance along two complementary metrics: (1) matcher accuracy and (2) describer efficiency. Matcher accuracy is given by the percent of matchers on each trial who successfully selected the target referent. Describer efficiency is given by the number of words produced by the describer to achieve that degree of matcher accuracy in the group. The degree to which describers are able to communicate more efficiently without negatively impacting matcher accuracy is indicative of convergence on a more effective shared communication protocol within the group.

We begin by examining matcher accuracy, the extent to which the intended target was reliably transmitted to all matchers. We constructed a series of 5 logistic mixed-effects regression models predicting accuracy as a function of condition and repetition block (separate models were run for experiment 1, each condition in experiment 2, and experiment 3). For this and other effects, there was substantial variation at the




Figure 2.2: Behavioral results across all three experiments. (A-C). Matcher accuracy at selecting the target image. (D-F). Number of words produced by the describer each trial. For all, small dots are per game, per block means, and smooth lines are predictions from model fixed effects with 95% credible intervals. Y-axes are truncated, and a few outliers points are not visible.

tangram and game levels, with some tangrams being markedly easier than others and some groups performing differently than others. This wide variation made it difficult to precisely estimate population-level main effects, leading to wide credible intervals. See Appendix Figure A11 for a visualization of the relative magnitudes of population effects and game and tangram level variations.

Across all conditions, we observed strong positive effects of repetition block, indicating improved performance over time (Figure 2.2A-C, Appendix Tables A4-A8). In Experiment 3, larger games began with lower initial accuracy ($\beta = -0.64$, 95% CrI = $[-1.05, -0.25]$) and improved more slowly ($\beta = -0.34$, 95% CrI = $[-0.43, -0.25]$) than smaller games, although group size differences were not reliable in Experiment 1 (Table A4), and these experiment 3 differences were not robust in a sensitivity analysis (Appendix Figure A4C and Table A58). Among large groups in Experiment 2, accuracy was higher in the thicker conditions than in the condition with thin interaction structure (Tables A5-A7), although effects of game thickness were not reliable in Experiment 3 (Table A8).

Because each experiment only explored a slice of the full parameter space, we also considered an exploratory analysis that pooled data across experiments, aiming to mitigate the loss in power from running entirely separate regression models. Specifically, we aggregated data from all experiments into a post-hoc mega-analytic model predicting accuracy as a function of repetition block, game thickness (thin v. not-thin) and game size. Overall, we found evidence that accuracy increased over time ($\beta = 0.46$, 95% CrI = $[0.4, 0.52]$) but the rate of increase was reduced for thin games ($\beta = -0.12$, 95% CrI = $[-0.21, -0.02]$) and larger games ($\beta = -0.07$, 95% CrI = $[-0.09, -0.05]$) compared to smaller or thicker games. That is, smaller groups and groups with higher coherence tended to be more accurate, though the magnitude and reliability of these effects varied across individual experiments.

2.2.3 Smaller and higher-coherence groups are more efficient

After establishing that groups were able to communicate accurately, we turned to the challenges faced by describers when deciding how much information to provide. Specifically, we predicted that larger and more heterogeneous groups may initially require more information, but that thicker interaction structure may similarly allow describers to communicate more effectively over time. We tested these predictions using linear mixed-effects models predicting the number of words a describer produced on each trial as a function of condition and block. These models counted all words the describer produced, including after matcher contributions (similar effects were

found in models predicting the length of describer’s utterances before any matcher contributions, see Tables A21-A24).

First, as predicted, describers in larger groups produced longer descriptions at the outset than describers in smaller groups (Figure 2.2D-F). This effect held for the continuous manipulation of group size for Experiment 1 ($\beta = 1.6$, 95% CrI = $[0.62, 2.6]$) as well as the 2-person versus 6-person manipulation in Experiment 3 ($\beta = 7.51$, 95% CrI = $[3.63, 11.3]$). Smaller groups also continued to use shorter descriptions than larger groups over the course of the game. In Experiment 1, the rate at which efficiency increased was similar across different size groups ($\beta = -0.09$, 95% CrI = $[-0.37, 0.18]$). In Experiment 3, larger groups reduced faster than smaller ones ($\beta = -1.22$, 95% CrI = $[-2.06, -0.29]$), but the faster reduction did not fully make up for the longer initial starting point, and was not robust to a sensitivity analysis (Figure A4F and Table A63).

While thin 6-person games showed a flatter reduction trajectory than thicker 6-person games in Experiment 2 (Tables A10-A12), there was no reliable effect of game thickness on reduction in Experiment 3 (Table A13).

The reduction patterns of description lengths is paralleled by how long matchers took to make selections; across conditions, matchers selected faster in later conditions (Figure A9), and the correlation between speed and description length was consistent across experiments (Figure A10).

Aggregating across experiments with a mega-analytic model, however, suggested that larger games were associated with steeper reduction ($\beta = -0.36$, 95% CrI = $[-0.51, -0.2]$) from a more verbose starting point ($\beta = 2.12$, 95% CrI = $[1.5, 2.75]$) than smaller games, and thin games had shallower reduction curves ($\beta = 0.79$, 95% CrI = $[0.04, 1.52]$) than thicker games. Overall, then, smaller games used shorter descriptions than larger games across various time points in the experiment, and thinner games reduced less than thicker games.

Table 2.1: Examples from 6-player groups in Experiment 3 of successful descriptions for the same image across repetitions. Describers are indicated with an asterisk. More example descriptions are in Appendix Tables A1 and A2.

6-person thick game			
<i>Rep 1: 4/5 correct</i>		6-person thin game	
A*	sitting down no legs showing	<i>Rep 1: 5/5 correct</i>	
C	no arms?	L*	man with no arms or legs
B	Bunny ears?	O	😞
A*	legs showing to one side no arms	P	😞
C	kinda like kneeling?	L*	slight protuberance on the right
A*	no feet	<i>Rep 2: 5/5 correct</i>	
A*	yes kneeling	N*	Can't see any arms. Imagine wrapped in a blanket completely.
<i>Rep 2: 5/5 correct</i>		N*	Armless and legless
A*	sitting down with no feet showing. legs to one side	N*	Burrito with a head
A*	no arms	M	😄
E	Swaddled up like a bb	O	😄
C	cute bb	P	😞
D	I think that's the one that looks like a ghost to me, like he's wearing a sheet. (?)	<i>Rep 3: 5/5 correct</i>	
<i>Rep 3: 5/5 correct</i>		O*	burrito
A*	sitting down no arms showing legs to one side	<i>Rep 4: 5/5 correct</i>	
E	the bb	Q*	burrito
B	Swaddled baby?	<i>Rep 5: 5/5 correct</i>	
F	Baby?	M*	burrito
A*	the bb	<i>Rep 6: 5/5 correct</i>	
<i>Rep 4: 5/5 correct</i>		P*	And our last one! Anyone else hungry after this? I quite fancy a burrito!
A*	sitting down with no arms legs to one side	L	😄
<i>Rep 5: 5/5 correct</i>		Q	✓
A*	sitting no arms and legs to one side. I will call them Kevin	N	😄😄😄
D	<3 Kevin	Q	😄✓
C	bb kevin	M	✓
<i>Rep 6: 5/5 correct</i>		Q	😄
A*	kevin the baby	N	✓✓✓
E	yess kevin		

2.2.4 Larger groups make greater use of matcher contributions

As a final measure of group performance, we examined the back-and-forth interactions between the describer and the group of matchers. Matchers use their chat contributions to actively provide feedback, ask questions, offer alternative descriptions, and seek clarification about the describer’s referring expressions. Example transcripts from successful games, one in the 6-thick condition and one in the 6-thin condition, are shown in Table 2.1. Additional examples are in the Appendix Tables A1 and A2. Overall, we found that larger groups displayed a higher proportion of trials where at least one matcher produced utterances (Figure A6A, $\beta = 0.79$, 95%CrI = [0.58, 0.98]), which declined across repetition blocks ($\beta = -0.8$, 95% CrI = [-0.97, -0.62]). On an individual level, a matcher in a larger group was more likely to make contributions than a matcher in a smaller group, although each contribution tended to be shorter (Figure A7, Tables A18, A20). The length of matcher interjections also decreased over time, especially for large groups (Figure A6D, $\beta = -0.41$, 95% CrI = [-0.72, -0.11]) consistent with the need for early matcher involvement in establishing referential conventions. Emoji use in Experiment 3 followed similar trends (Figure A8). Overall, describers in larger groups receive more total input from matchers, suggesting larger groups may require greater participation by matchers to reliably establish common ground.

2.2.5 Descriptions converge faster in groups with thicker channels

In the previous sections, we examined three metrics of communicative performance in groups of different sizes and interaction structures. We confirmed that groups in all conditions replicated the classic patterns of increasing accuracy and decreasing description length. We also found some initial evidence that larger groups may struggle to improve performance in the absence of thick communication channels. Here, we aim to better understand the mechanisms that allow describers to use shorter descriptions without sacrificing accuracy. In particular, we explore the hypothesis that




Figure 2.3: Example utterances describing the shown tangram figure produced by two 3-player games in Experiment 1. To measure convergence within a game (blue), we measured the cosine similarity between SBERT embeddings of descriptions and the embedding of the round 6 utterance (taken to be the convention). Higher cosine similarity indicates more similar meaning. To measure divergence between games (green), we measured the similarity between embeddings of utterances from the same round across games.

interaction structure and group size affect performance through a *convention formation* process (Clark & Wilkes-Gibbs, 1986). Under a recent model of convention formation (Hawkins, Franke, et al., 2023), groups are able to leverage their shared history to coordinate on stable expectations about how to refer to particular images. This model makes specific predictions about how interaction structure affects the ability to coordinate, in terms of the available feedback.

First, due to heterogeneity in the group – 6 individuals who may have diverging conceptualizations — a rational describer should provide a strictly more detailed initial description to hedge against multiple possible misunderstandings, as we previously observed. Second, all groups should display the characteristic dynamics of conventions: *stability*, or convergence within group, and *arbitrariness*, or divergence

to multiple equilibria across groups. Third, convergence should be faster when a single individual is consistently in the describer role and when matchers are able to freely respond in natural language, as describers are able to aggregate feedback about the effectiveness of their own utterances from block to block and also immediately correct specific misunderstandings within a given trial.

To assess the dynamics of describer descriptions, we examine the *semantic similarity* of descriptions within and across games. We quantified description similarity by concatenating describer messages together within a trial and embedding this description into a high-dimensional vector space using SBERT. SBERT is a BERT-based sentence embedder designed to map semantically similar sentences to embeddings that are nearby in embedding space. Semantically meaningful comparisons between sentences are made by taking pairwise cosine similarities between the embeddings (Reimers & Gurevych, 2019).

To measure stability, or convergence within groups, we compared utterances from blocks one through five to the final (block six) description for the same image from the same game. To measure arbitrariness, or divergence across groups depending on group-specific history, we compared utterances produced by different describers for the same image in the corresponding blocks. Figure 2.3 illustrates these two measures with example utterances and their within-game and between-game cosine similarities.

We modeled semantic convergence with a mixed effects linear regression model predicting the similarity between a block 1-5 utterance and the corresponding block 6 utterance as a function of the earlier block number and condition (Figure 2.4A-C; Tables A25-A29). All conditions showed some convergence toward a conventional “shorthand” for the picture, but the speed of convergence was affected both by group size and channel width. First, we found that smaller groups reached stable descriptions faster than larger games. In Experiment 1, initial similarity was invariant across group size ($\beta = -0.008$, 95% CrI = $[-0.021, 0.005]$), but smaller groups converged faster (Figure 2.4A, $\beta = -0.008$, 95% CrI = $[-0.011, -0.005]$). In Experiment 3, 6-person thick games started off further from their eventual convention than 2-person thick games ($\beta = -0.069$, 95% CrI = $[-0.113, -0.025]$) but closed the gap over time (Figure 2.4C, $\beta = 0.009$, 95% CrI = $[0.001, 0.017]$, this effect was not robust to sensitivity analysis, Figure A5C and Table A68). Second, thicker games tended to converge




Figure 2.4: Language similarity results measured with pairwise cosine similarity between embeddings of two utterances. (A-C). Convergence of descriptions within games as measured by similarity between an utterance from block 1-5 to the block 6 utterance in the same game for the same image. (D-F). Divergence of descriptions across games as measured by the similarity between two utterances produced for the same image by different groups in the same block. For all, small dots are per game, per block means, and smooth lines are predictions from model fixed effects with 95% credible intervals. Y-axes are truncated, and a few outliers points are not visible.

faster than thin games (Figure 2.4B-C). In Experiment 3, small thin games started off slightly further from their convention than small thick games, and this gap widened

over time ($\beta = -0.025$, 95% CrI = $[-0.033, -0.017]$). Finally, the combination of thin interaction structure and larger group hindered convergence more than either factor individually. Beyond the generally slower convergence in thin games, 6-person thin games showed substantially slower convergence even compared to 2-person thin games in Experiment 3 ($\beta = -0.035$, 95% CrI = $[-0.047, -0.025]$).

Pooling across experiments in a mega-analysis confirms this pattern. Thin games converge less than thick games overall ($\beta = -0.016$, 95% CrI = $[-0.025, -0.008]$), and *large* thin games are especially slow to converge ($\beta = -0.007$, 95% CrI = $[-0.01, -0.004]$). Across games, convergence towards the last utterance was driven by cumulative increasing similarity between pairs of utterances in adjacent blocks (Figure A12D-F, Tables A40-A44). In early rounds, descriptions could change substantially between rounds, but by later rounds, many descriptions had already reduced and solidified and varied little round to round. In summary, we found that stable descriptions emerged earlier if the group was smaller, or if the group had a thick interaction structure.

2.2.6 Games with thicker channels diverge from one another more quickly

While groups may initially overlap in their descriptions, including details of shapes or body parts, we predicted that their descriptions would become increasingly dissimilar as groups increasingly adapt to their own idiosyncratic shared history. To test this effect, we constructed a mixed-effects linear regression model predicting the cross-game similarity between a pair of utterances for the same image. A decrease in the similarity between different groups descriptions occurred in every condition, indicating increasing arbitrariness and group-specificity of descriptions (Figure 2.4D-F, Tables A30-A34). However, different game sizes and interaction structures revealed very different strengths of divergence.

First, smaller games used more group-specific language. In Experiment 1, smaller games diverged more quickly than larger games ($\beta = 0.001$, 95% CrI = $[0.001, 0.002]$). In Experiment 3, 2-person thick games started off more dissimilar than 6-person thick games, although 6-person games diverged faster and eventually approached

the dissimilarity levels of 2-person thick games (Table A34). Second, thicker interaction structure was associated with stronger group-specific divergence. In Experiment 3, 2-person thin games diverged more slowly than 2-person thick games ($\beta = 0.004$, 95% CrI = $[0.002, 0.005]$). As with the convergence patterns, large games with thin interaction structures had the flattest trajectories, as thinness and largeness compounded. In Experiment 3, 6-person thin games diverged even less than 2-player thin games (Figure 2.4F, $\beta = 0.017$, 95% CrI = $[0.015, 0.019]$), and in Experiment 2, 6-person thin games barely diverged at all (Figure 2.4E, $\beta = -0.004$, 95% CrI = $[-0.006, -0.001]$). A mega-analytic model confirms this pattern: thin games differentiate less between groups ($\beta = 0.005$, 95% CrI = $[0.004, 0.007]$) and large thin groups differentiate even less ($\beta = 0.004$, 95% CrI = $[0.004, 0.005]$).

As a complement to the embedding analysis, we also examined the frequency of a few classes of words in the descriptions. Literal geometric words (ex. square, triangle, etc) and words for body parts (leg, arm, etc) are common early in games, but decline over repetition in most conditions, to be replaced by more abstract descriptions that do not contain these classes of words (Figure A2). The 6-person thin condition, however, retains a higher level of literal geometric and body part words, along with high levels of positional words (above, left, below, etc) and posture words (kicking, standing, seated, etc), with a lower level of utterances that do not contain any of these classes of words.

2.3 Discussion

From classrooms to boardrooms, human communication often takes place in multi-party settings. However, experimental research rarely focuses on such settings, largely due to practical obstacles. In the current work, we asked how convention formation processes, typically studied in dyadic reference games, unfold in larger groups and under varying interaction structures. Across 3 online experiments and 11 experimental conditions, we varied multiple features of interaction structure including group size, modality of matcher contributions, and degree of group coherence. All conditions replicated classic dyadic phenomena: increasing accuracy and efficiency, semantic

convergence within games, and differentiation of descriptions between groups. However, we also found that the interaction structure substantially affects how rapidly groups develop partner-specific conventions. Small groups may be able to successfully form conventions under limited feedback, but larger groups require thicker interaction structure. Multi-player groups may therefore reveal key factors which are masked in purely dyadic settings.

Increasing efficiency, for example, has often been taken as an index of group-specific convention formation (Brennan & Clark, 1996; Clark & Wilkes-Gibbs, 1986; S. O. Yoon & Brown-Schmidt, 2014, 2019a). In our work, however, we observe distinct patterns for measures of raw utterance length compared to the dynamics of semantic content. In Experiment 3, thin 6-person games showed much less group-specific divergence despite comparable accuracy and efficiency. This gap raises the possibility that it is possible to become more efficient and accurate without negotiating a unified group-specific label. Instead, they may be relying more strongly on the group’s priors (Guilbeault et al., 2021). Thus, we encourage measures of semantic content (and not just performance) when evaluating convention formation. The transcripts for these games provide a rich dataset for exploring different ways language is used to form referential conventions.

The causal mechanisms driving group size effects remain unclear. There are many differences between a 2-person group and a 6-person group that could plausibly lead to different outcomes. For example, in a dyad, producers can tailor their utterances to the one matcher, but in large groups, producers must balance the competing needs of different comprehenders (Schober & Clark, 1989; Tolins & Fox Tree, 2015; S. O. Yoon & Brown-Schmidt, 2019a). These effects likely vary with the knowledge state and the communication channels available to comprehenders (Fox Tree & Clark, 2013; Horton & Gerrig, 2002, 2005). Further work digging into the language used and the interactions between participants might unearth plausible mechanisms for how differences in group size and interaction structure influence outcomes, pointing towards future experimental conditions.

Even within the boundaries of the repeated reference game paradigm, there is a high-dimensional space of possible experiments. We sampled only a few points along

a few salient dimensions. In our experiment 3, we grouped some factors together in order to run more games in each condition: a fully factorial design would have been too expensive to power adequately. We instantiated a “thin” channel by limiting matchers to 4 discrete utterances (emojis), but there are other possible restrictions that could be placed on the channel, such as rate-limited typing or explicit time pressure. Future work could explore other dimensions of the interaction structure, introducing pre-existing relationships or familiarity among group members, alternative incentives involving competition and power, or alternative referential targets involving more complex concepts.

A particularly important dimension shaping interaction is the modality of communication, including whether the participants use oral or written language, whether they are co-present in the same space, and whether they have visual access to each others’ faces and gestures. Distinct modalities carry distinct affordances and norms. In this work, we relied on a text-based chat modality without allowing co-presence or visual access.

We suspect that the general pattern of effects we see, in terms of group size and coherence, are likely to extend to other modalities. However, different modalities may allow for different strategies that may be more or less sensitive to group size, describer rotation, or different levels of matcher contributions. For instance, in face-to-face oral settings, it may be easier for describers to continuously talk until interrupted, or to monitor the comprehension of individual group members from their facial expressions.

In conclusion, narrowly focusing on the settings that are easy to study in the lab – dyads with rich communication channels – can lead to theories that mispredict how interactions play out in multi-party groups. By studying common ground and coordination across a wider range of interaction structures, we can develop a more nuanced understanding of the obstacles that stand in the way of successful communication and how groups can overcome them. This understanding can inform the design of policies and collaborative platforms that promote effective communication in various contexts, from small-scale conversations to large-scale civic discourse. As remote work and online communication become increasingly prevalent, it is increasingly crucial to understand how the structure of group communication environments shapes the effectiveness of human communication.

2.4 Materials and Methods

Our iterated reference task was implemented with Empirica (Almaatouq et al., 2021), a React-based web development framework for real-time multi-player tasks. Our experiments were designed sequentially and pre-registered individually.¹ We followed the pre-registered analysis plan for each experiment, although accuracy models were not explicitly specified until Experiment 3, and linguistic analyses were only verbally described starting with Experiment 2b. Results from some pre-registered models are omitted from the main text for brevity but are shown in the Appendix. Exploratory mega-analytic models pooling across the three experiments were not pre-registered.

All materials, data, and analysis code is available at <https://github.com/vboyce/multiparty-tangrams>.

2.4.1 Participants

This research was covered by the Stanford IRB under protocol 20009 “Online investigations of language learning”. Participants were recruited using the Prolific platform. All participants self-reported as fluent native English speakers on Prolific’s demographic prescreen. Experiment 1 took place between May and July 2021, Experiment 2 between March and August 2022, and Experiment 3 in October 2022. Each participant took part in only one experiment and was blocked from participating in subsequent experiments. As games with more participants tended to run longer, we paid participants different rates based on group size, with the goal of a consistent \$10 hourly rate. Participants were paid \$7 for 2-player games, \$8.50 for 3-player games, \$10 for 4-player games, and \$11 for 5- and 6-player games. When one player occupied the describer role for the entirety of a 6-player game, they were rewarded an additional \$2 bonus. Across all games, participants could earn up to \$2.88 in performance bonuses.

A total of 1319 people participated across the 3 experiments. We recruited enough participants for 20 games in each condition in experiments 1 and 2 and 40 games per

¹Experiment 1: <https://osf.io/cn9f4> for the 2-4 player groups, and <https://osf.io/rpz67> for the 5-6 player data run later. Experiment 2: same describer at <https://osf.io/f9xyd>, full feedback at <https://osf.io/j5zbn>, and thin at <https://osf.io/k5f4t>. Experiment 3: <https://osf.io/untzy>

condition in experiment 3. However, due to attrition in filling the games initially and due to participants dropping out of the games, we ended up with fewer games in some conditions. For logistical reasons of matching participants into real-time games, we had to recruit participants in fairly large batches, and so did not have precise control to add new games to replace games that did not fill or had participants drop out early. A breakdown of number of games and participants in each condition is shown in Table A3 along with further discussion of recruitment logistics.

2.4.2 Materials

The same 12 tangram images, drawn from Hawkins, Frank, et al. (2020) and Clark & Wilkes-Gibbs (1986), were used every block. These images were displayed in a 4×3 grid with the order randomized across participants to disincentivize spatial descriptions such as “top left,” as the image might be in a different place on the describer’s and matchers’ screens. To reduce cognitive load from visual search, the locations were fixed for each participant across trials.

2.4.3 Procedure

The experimental procedure was very similar across the three experiments. We first describe the procedure used in Experiment 1 and then describe the differences in later experiments.

Experiment 1 Participants were directed from Prolific to our custom web application, where they were presented with a consent form and a series of instruction pages explaining the protocol. After finishing the instructions, they needed to pass a quiz to proceed. They were then directed to a “waiting room” lobby. Once the lobby filled to the required number of players, the game began. One lobby was filled before another was started; if a participant was waiting for 5 minutes, that lobby timed out, and the participant was paid without completing the experiment. Due to technical constraints with assigning participants to lobbies and games, only games of a single experimental condition could be active at a time. Thus, different conditions were run on different days or times of day.

One of the participants was randomly selected to begin in the role of describer, and the other participants were assigned to the role of matchers. On each trial, the describer saw a fixed array of tangrams with one tangram (privately) highlighted as the *target*. They were given a chat interface to communicate the target to the matchers, who were asked to determine which of the 12 images was the referential target. All participants were free to use the chat box to communicate at any time, but matchers could only make a selection after the describer had sent a message. Once a matcher clicked, they could not change their selection. There was no signal to the describer or other matchers about who had already made a selection. We recorded what all participants said in the chat, as well as who selected which image and how long they took to make their selections.

Once all matchers had made a selection (or a 3-minute timer ran out), participants were given feedback and proceeded to the next trial. Matchers only received *binary* feedback about whether they had chosen correctly or not; that is, matchers who made an incorrect choice were not shown the correct answer (see Figure A1 for example feedback). The describer saw which tangram each matcher selected, but matchers did not see one another's selections. Matchers got 4 points for each correct answer; the describer got points equal to the average of the matchers' points. These points were translated into performance bonuses at the end of the experiment (1 point = 1 cent bonus). After the describer had described each of the 12 images as targets, in a randomized sequence, the process repeated with the same set of targets, for a total of 6 such repetition blocks (72 trials).

The same person was the describer for an entire block, but participants rotated roles between blocks. Thus, over the course of the 6 blocks, participants were describers 3 times in 2-player games, twice in 3-player games, once or twice in 4 and 5-player games, and once in 6-player games. Rotating the describer was chosen in this first experiment to keep participants more equally engaged (the describer role is more work), and to provide a more robust test of our hypotheses regarding efficiency and convention formation. After the game finished, participants were given a survey asking for optional demographic information and feedback on their experience with the game.

If a participant disconnected from the experiment, the game would stop.

Experiment 2 Experiment 2 consisted of three different variations on Experiment 1, all conducted in 6-player games. Each of these conditions differed from the Experiment 1 baseline in exactly one way. In the *same describer* condition, one person was designated the describer for the entire game, rather than having the describer role rotate. In the *full feedback* condition, all participants were shown what all others had selected as well as the identity of the correct target. This condition was similar to previous dyadic work, such as Hawkins, Frank, et al. (2020), where the correct answer was indicated during feedback. In the *thin* condition, we altered the chatbox interface for matchers. Instead of a textbox, matchers had 4 buttons, each of which sent a different emoji to the chat. Matchers were given suggested meanings for the 4 emojis during the instruction phase. They could send as many emojis as desired; for instance, they might initially indicate confusion, and later indicate understanding. In addition, for the thin condition, we added notifications that appeared in the chat box marking the time when each player had made a selection.

Experiment 3 The thin channel condition in Experiment 3 was the same as the thin condition in Experiment 2. The thick condition combined the two coherency-enhancing variations from Experiment 2: the same participant remained in the describer role throughout, and full feedback was given about the correct answer and what all other players had selected. Across both conditions in Experiment 3, notifications were sent to the chat to indicate when a participant had made a selection. For experiment 3, game lobbies worked slightly differently, and 5 minutes after the first participant had joined the lobby, the game started if there were at least two participants. Correspondingly, in experiment 3, games did not stop if a player disconnected, instead if there were at least two players still active, the game continued, swapping a player into the role of describer if necessary to continue the game.

2.4.4 Data pre-processing and exclusions

Participants could use the chat box freely, which meant that the chat transcript contained some non-referential language. The first author skimmed the chat transcripts, tagging utterances that did not refer to the current tangram. These were primarily pleasantries (“Hello”), meta-commentary about how well the task was going, and bare

confirmations or denials (“ok”, “got it”, “yes”, “no”). We excluded these utterances from our analyses. Note that chat lines sometimes included non-referential words in addition to words referring to the tangrams (“ok, so it looks like a zombie”, “yes, the one with legs”); these lines were retained intact.

In Experiments 1 and 2, games did not start if there were not enough participants and ended if any participant disconnected. In Experiment 3, games started after a waiting period even if they were not entirely full and continued even in the event that a participant disconnected (with describer role reassigned if necessary), unless the game dropped below 2 players. The distribution of player counts in games that were initially recruited to be 6 player games is shown in Figure A3. The realities of online recruitment and disconnection meant that the number of games varied between conditions. We excluded incomplete blocks from analyses, but included complete blocks from partial games (See Table A3). When skimming transcripts to tag non-referential utterances, we noticed that one game in the 6-player thick condition had a describer who did not give any sort of coherent descriptions, even with substantial matcher prompting. We excluded this game from analyses.

2.4.5 Modelling strategy

We fit all regression models in brms (Bürkner, 2018) with weakly regularizing priors. We were unable to fit the full pre-registered mixed effects structure in a reasonable amount of time for some models, so we included the maximal hierarchical effects that were tractable. All model results and priors and formulae are reported in the Appendix. Models of accuracy used by-group random intercepts only, models of word count used full mixed effect structure, and models of S-BERT similarities used by-group and by-target random intercepts as applicable (see Figure A11). Models of matcher accuracy were logistic models with normal(0,1) priors for betas and sd. Models of describer efficiency were run as linear models with an intercept prior of normal(12,20), a beta prior of normal(0,10), an sd prior of normal(0,5) and a random-effect correlation prior of lkj(1). For all of the models of SBERT similarity, we used linear models with the priors normal(.5,.2) for the intercept, normal(0,.1) for betas, and normal(0,.05) for sd. As an additional post-hoc analysis, we ran mega-analytic models combining data across all experiments. For these models, we grouped the 3

thin-ish conditions (2c, and the two thin conditions of experiment 3) as one level, and coded the rest of the conditions as thick-ish. Game size was coded as a continuous measure (2 through 6). The priors for the mega-analytic models were the same as for the per-experiment models described above.

As a sensitivity analysis, we re-ran the primary models on the subset of the data from games that a) completed all 72 trials and b) had the full complement of players the entire time (relevant to 6-player experiment 3 games where games could start or continue with fewer players). Discrepancies are mentioned in the results, and these analyses are depicted in Figures A4 and A5 and Tables A54-A73. We also needed to decide how to handle dropout in Experiment 3, as some of the 6-player games did not retain all 6 players for the entire game. Our decision was to follow an intent-to-treat analysis and treat data as missing completely at random. Note that this choice underestimates differences between 2-player and (genuine) 6-player games by labeling some smaller groups as 6-player groups. We do not know exactly what leads some participants to drop out, but it is possible that some factors may be random (ex. connection issues) and others may be correlated with performance (ex. frustration because group is struggling).

We do not know whether groups that start and continue at the full size differ from games where some participants drop out. This is potentially an issue across all experiments; in experiments 1 and 2, groups stopped playing if anyone dropped out, and in experiment 3 they kept playing as a smaller group. The number of games in each condition and rates of dropoff are shown in Table A3 and Figure A3.

Chapter 3

Preschoolers can form conventional pacts with each other to communicate about novel referents

3.1 Introduction

Learning a language requires learning not only the content of that language, but also how to use the language to communicate. One case study for language use is referential communication, the ability to describe a target so an interlocutor can pick it out from a set of possibilities. Adults show sensitivity to both the visual context and their audience during referential communication, calibrating the description they provide to their beliefs about the interlocutor’s knowledge state.

Iterated reference games provide an important paradigm for studying referential communication. In these games, one player repeatedly describes a set of abstract shapes to a partner so they can identify the target images (Boyce et al., 2024; Clark & Wilkes-Gibbs, 1986; Hawkins, Frank, et al., 2020; Krauss & Weinheimer, 1964). Over repetition, features of the initial descriptions are conventionalized as each pair comes to agree on a shared understanding of how to label each image. Success at this task requires mastery of a number of linguistic and communicative skills, including producing adequate initial descriptions, monitoring for comprehension, asking for clarification, and appropriately using the shared conversation history to inform later

referring expressions. Studying how children play iterated reference games can provide insight into the developmental trajectory of the ability to produce referential expressions in order to achieve joint understanding.

One influential early study suggested that 4-5-year-old preschoolers struggle with child-child referential communication (Glucksberg et al., 1966). In their paradigm, one child was given a set of 6 blocks in a specific order. Their task was to describe the image on each block so their partner could pick out their corresponding block. As children described and selected blocks, they stacked them on pegs. While 4-5-year-old children succeeded on practice trials with familiar shapes and visual access to each other's blocks, children failed on critical trials where the blocks had abstract drawings and there was no visual access. Even after multiple rounds with the same images, children were not able to correctly order the blocks. Glucksberg et al. (1966) attributed children's communicative failures to their production of ego-centric descriptions that did not account for the other child's perspective.

Similar experiments with older children indicated a gradual improvement through adolescence both for initial accuracy and for the increase in accuracy across repetitions. Still, even the 9th grade sample was noticeably worse than the adult college student sample (Glucksberg & Krauss, 1967; Krauss & Glucksberg, 1969). Given that even teenagers had difficulties with the task, the complex stacking paradigm and the large number of potential targets may have posed task demands that prevented accurate measurement of children's abilities.

More recently, though no evidence has directly contradicted Glucksberg et al. (1966), a number of other studies have revealed early emerging skills in the preschool years that support the use of referential communication. Preschool-aged children are faster to select a target when it is referred to in a consistent way (Graham et al., 2013; Matthews et al., 2010) and 6-year-olds are more likely to use consistent referential expressions when their partner is consistent (Köymen et al., 2014), suggesting that young children are sensitive to the norm of consistent descriptions in cooperative communication. Further, preschool-aged children adapt the informativeness of their referential expressions based on the visual content available to their interlocutor (MATTHEWS et al., 2006; Nadig & Sedivy, 2002; Nilsen & Graham, 2009). By age 5, children integrate information about other's perspectives into their comprehension




Figure 3.1: Experimental setup and procedure. Panel A shows the experimental setup. Panel B shows the 4 possible targets; names for targets are for cross-reference with later figures only. Panel C shows the procedure for Experiment 1; within critical blocks targets were ordered randomly. Panel D shows the procedure for Experiment 2.

of utterances (San Juan et al., 2015). Overall, by late preschool, children can reason about others’ perspectives to communicate more effectively.

While children are sensitive to others’ perspectives, they still struggle to appropriately tailor the specificity of their utterances to the visual context, often resulting in under-informative utterances (Leung et al., 2024; Matthews et al., 2012). When children must describe one of two very similar images, 4 and 5-year-olds sometimes neglect to mention the relevant features, although they do better when playing in an interactive game than when not (Grigoroglou & Papafragou, 2019). By 5 years old, children are sensitive to over- and under-informative utterances, asking for clarification on some under-informative utterances and taking longer to make selections on over-informative utterances (Morisseau et al., 2013).

A common thread among many developmental studies is that children perform better when the cognitive load of the task is reduced, i.e. when there are fewer possible referents (Abbot-Smith et al., 2015; San Juan et al., 2015). These findings paired with evidence of emerging communicative skills in young children suggest that tailored referential expression production is possible in children, but it is likely to be masked

when the task demands are too high.

Since Glucksberg et al. (1966), few studies have revisited the question of whether children can successfully communicate in an iterated reference game. In one study, 8-10-year-olds exhibited adult-like patterns of increasing accuracy, increasing speed, and shorter descriptions across repetitions, but children’s accuracy was still far below adult performance and highly variable between dyads (H. P. Branigan et al., 2016). 4-6-year-olds successfully used conventionalized gestures with their partners in an iterated reference game where children could only use gestures to communicate (Bohn & Frank, 2019). 4-8-year-old children succeeded at playing an iterated reference game with a parent using a simple, child-friendly tablet-based task (Leung et al., 2024). In this task, even 4-year-olds had an initial accuracy above 80%, which rose to above 90% in later repetitions (Leung et al., 2024). Together, these findings provide evidence of young children’s ability to communicate about novel referents under the right conditions.

Given the task demands in Glucksberg et al. (1966) and work showing that children can succeed in a less demanding paradigm, here we revisit the question of child-child referential communication. In the present study, we re-examine young children’s ability to establish effective referring expressions with each other in an iterated reference game using a simplified tablet-based paradigm. Across an initial study and a replication including a total of 51 pairs of 4-5-year-old children, we found that preschool-aged children were successful in an iterated reference game, suggesting that children’s capacity to construct effective referring expressions in novel contexts emerges earlier than once claimed.

3.2 Experiment 1

3.2.1 Methods

Our goal for Experiment 1 was to test young children’s ability to coordinate to produce descriptions of abstract shapes that their partner could understand. Young children can be very sensitive to task demands and cognitive load (Carruthers, 2013; Keen, 2003; Turan-Küçük & Kibbe, 2024), so we adapted the experimental framework from Leung et al. (2024), and further simplified it by reducing the total pool of targets

and the number of trials. This experiment was pre-registered at [anonymized link](#).

Participants

4 and 5-year-old children were recruited from a university preschool during the school day. Children played with another child from the same class. Experiment 1 was conducted between June and August 2023. Pairs of children were included in analyses if they completed at least 8 of the 12 critical trials. We had 19 games that completed 12 critical trials, and 1 game that completed 11 critical trials. Of the 40 children, 21 were girls, and the median age was 57 months, with a range of 48-70 months.

Materials

For the target stimuli, we used four of the ten tangram images from Leung et al. (2024), chosen based on visual dissimilarity (Figure 3.1B). We coded the matching game using Empirica (Almaatouq et al., 2021), hosted it on a server, and then accessed the game on tablets that were locked in a kiosk mode so children could not navigate away from the game.

Procedure

Once a pair of children agreed to play the game, a research assistant took them to a quiet testing room. Children were introduced to a stuffed animal “Smurfy” who wanted to play a matching game. Children sat across a table from each other, each with a tablet in front of them (Figure 3.1A). On each trial, one child was the “teller” and saw a black box around one of two images on their screen and was asked to “tell Smurfy what they see” in the black box. The “guesser” saw the same two images in a randomized order and tried to select the described image to help Smurfy make a match. When the guesser selected an image, both children received feedback in form of a smiley or frowny face and an excited or disappointed sound. After each trial, children switched roles. Children passed Smurfy back and forth to keep track of whether they were the “guesser” or “teller” on a given trial.

Children completed two warm-up trials with black and white images of familiar shapes, followed by 3 blocks of the 4 target images (Figure 3.1C). Targets were

randomly paired with another of the critical images as the foil.

The experimenters running the game did not volunteer descriptions, but they did scaffold the interaction, prompting children to describe the images, and sometimes repeating children’s statements (especially when utterances were inaudible or the child did not respond immediately; this aspect of the procedure was modified in Experiment 2). The entire interaction was video-recorded.

Data processing

Children’s selections and the time to selection were recorded from the experiment software. Children’s descriptions were automatically transcribed from the video using Whisper (Radford et al., 2022) for the first pass and then hand-corrected by experimenters. Transcripts were hand-annotated for when each trial started, who said each line, and what referential descriptions were used. We excluded trials where the “teller” did not produce a description, or where all description was unintelligible and impossible to transcribe. After exclusions, we had 231 trials remaining.

Statistical analyses were run in brms (Bürkner, 2018) with weakly informative priors. We report estimates and 95% credible intervals. The experimental set-up, analysis code, de-identified transcripts, and performance data for both experiments is available at anonymized repo.

3.2.2 Results

Accuracy and speed

Our primary measure of interest was whether children could accurately communicate the intended target. To test for changes in accuracy over time, we fit a Bayesian mixed effects logistic regression predicting accuracy.¹ Children’s accuracy was above chance (Odds Ratio: 3, (95%CrI = [1.14, 8.09])), and their accuracy slightly increased over the game (OR of one trial later: 1.17, (95% CrI = [1.03, 1.39]), Figure 3.2). This level of accuracy is generally in-line with accuracies from 4-year-olds playing with their parents in Leung et al. (2024) and indicates that children can understand and succeed at the task.

¹ $\text{correct.num} \sim \text{trial.num} + (\text{trial.num}|\text{game}) + (1|\text{target})$




Figure 3.2: Children’s accuracy at selecting the correct target over time. Error bars are bootstrapped 95% CIs with a linear trend line overlaid.

As another measure of children’s performance, we looked at how long children spent on each trial. We ran a Bayesian mixed effects linear regression predicting the time to selection in seconds.² The first critical trial averaged 27.48, (95% CrI = [20.48, 34.53]) seconds, and children got faster over time (-1.22, (95%CrI = [-1.9, -0.53]) seconds / trial). Children were able to achieve the same accuracy in less time, suggesting that they were becoming more efficient at completing the task.

Description length

In iterated reference games with adults, description lengths usually shorten over repeated references (Boyce et al., 2024; Clark & Wilkes-Gibbs, 1986; Hawkins, Frank, et al., 2020). We were curious if children’s descriptions would display the same trend, so we ran a Bayesian mixed effects linear regression predicting the number of words in the description the “teller” produced.³ On the first critical trial, descriptions averaged 3.66, (95% CrI = [2.56, 4.74]) words, and description length was relatively stable over

² $\text{time.sec} \sim \text{trial.num} + (\text{trial.num}|\text{game}) + (1|\text{target})$

³ $\text{words} \sim \text{trial.num} + (\text{trial.num}|\text{game}) + (1|\text{target})$




Figure 3.3: Length of description produced by the teller each trial. Grey dots are individual data points, colored dots are per trial means with bootstrapped 95% CIs.

time (change of 0.12, (95% CrI = $[-0.04, 0.28]$) words per trial, Figure 3.3). Thus, children’s increasing speed was not from shorter utterances, but instead from some combination of improved task understanding, faster utterance planning, and faster decisions of what to select.

Some examples to illustrate the variety of effective descriptions children employed are shown in Table 3.1 (these specific examples are from Experiment 2, but both experiments had similar distributions of descriptions).

Convergence

While description length is often used as a proxy for measuring convention formation, it does not capture semantic overlap between utterances. Boyce et al. (2024) introduced a more sensitive measure of semantic convergence that compares the content of utterances using word embeddings to trace how similarities within and across games change over time. With only 3 blocks of descriptions, we do not expect semantic similarity for descriptions of a given target to show any meaningful change over time. However, we can test for a more coarse measure of sensitivity to partner: whether




Figure 3.4: Semantic similarity between pairs of descriptions from different sources. Dots are means and lines are bootstrapped 95% CIs.

children’s utterances are more like their partner’s than children in other games’. If children are fully ego-centric (as suggested by Glucksberg et al., 1966), their choices of descriptions would be independent from their partners.

Following the methods of Boyce et al. (2024), we embedded each description in a semantic vector space using S-BERT (Reimers & Gurevych, 2019), and then used the cosine between embeddings as a measure of semantic similarity.

We compared the semantic similarities between descriptions of the same target based on who produced the description (Figure 3.4). We used a Bayesian mixed effects linear regression to predict similarity.⁴ Utterances were more similar if they came from the same partnership (0.222, (95% CrI = [0.184, 0.26])). Utterances were slightly more similar still if they came from the same person (0.135, (95% CrI = [0.077, 0.193])), which is expected since children are likely to be fairly consistent with themselves. However, children used descriptions that were much more similar to their partner’s than to other children’s (Figure 3.4), indicating sensitivity to their partner’s expressions.

⁴ $\text{sim} \sim \text{same_game} + \text{same_speaker} + (1|\text{target})$

Table 3.1: Example descriptions children successfully used to identify different target images in Experiment 2.

• person • a person holding a sandwich • a people carrying a box of dirt • a monster • someone holding a plate and giving it to a restaurant and has watermelon	
• vampire • hopping • a person flying • a person • a kite • a triangle with a head on it with feet • somebody skydiving, not in the airplane	
• racecar • airplane • alligator • a person fell down • a boat • a person that's in a race car that has one triangle and two triangles	
• person • a person walking • a person looking down • a people, but it doesn't have any arms	

3.2.3 Discussion

In Experiment 1, we adapted the paradigm of Leung et al. (2024) for pairs of children, taking an already simple set-up and making it shorter. Our goal was to see if young children were at all able to provide adequate descriptions, so children received a lot of scaffolding around the experimental interaction. Sometimes, this scaffolding included experimenters echoing children's descriptions, which could potentially influence children's responses. In Experiment 2, we repeated the same paradigm, with a tighter experimental script and a larger sample size.

3.3 Experiment 2

3.3.1 Methods

As Experiment 2 was very similar to Experiment 1, we focus on the changes made compared to Experiment 1. Experiment 2 was pre-registered at [anonymized link](#).

The biggest change between the experiments was increasing the number of repetitions of target stimuli from 3 to 4 (from 12 to 16 trials). The greater number of trials in Experiment 2 made it possible to look for changes over time that could be indicative of convergence to shared descriptions within a game and divergence between games.

Participants

Experiment 2 was run between March and August of 2024, at the same preschool as Experiment 1. No children participated in both experiments. 30 pairs of children completed all 16 critical trials, and 1 pair of children completed 10 critical trials. Our target age range was 4 and 5-year-olds, but one older 3-year-old was unintentionally included. Of the 62 children, 30 were girls, and the children had a median age of 56 months, and a range of 45-69 months.

Materials

The same 4 critical images were used as in Experiment 1. In response to some children struggling with the abrupt switch from familiar to non-nameable shapes, we introduced more practice trials for Experiment 2. We used a total of 4 practice trials to provide a gradient from familiar shapes to less recognizable, blockier shapes (Figure 3.1D).

Procedure

The procedure was much the same as Experiment 1 (Figure 3.1D). We added an initial “bubble popping” exercise to give children practice tapping the tablet appropriately (this was an issue for some children in Experiment 1). The experimental script was fully written out and memorized by experimenters so children all received the same

instructions. We wrote up contingency statements that the experimenter could use to prompt children who were not giving descriptions or making selections. Experimenters helped with game mechanics such as whose turn it was to tell and who should press the screen, but avoided contributing or repeating any content about the images or the descriptions.

Data processing

Data were processed in the same way as Experiment 1. After excluding trials where children did not give a description or where the experimenter echoed a child's description, we had 466 trials total.

3.3.2 Results

Accuracy and speed

In Experiment 2, children's accuracy was above chance (Odds Ratio: 5.95, (95%CrI = [3.07, 11.89])) and relatively stable over time (OR of one trial later: 1.01, (95% CrI = [0.94, 1.09]), Figure 3.2). The first critical trial averaged 22.06, (95%CrI = [15.86, 28.58]) seconds, and children got faster over time (-0.7, (95% CrI = [-0.99, -0.41]) seconds / trial). Children were initially faster in Experiment 2 than Experiment 1, possibly due to the increased number of practice trials and pre-training on how to press the screens. Taken together, we find more evidence that children can successfully communicate with each other about these abstract shapes, and do so with increasing efficiency.

Description length

The average length of descriptions on the first trial was 3.44, (95% CrI = [2.23, 4.75]) words and description length was relatively stable over time (change of 0.02, (95%CrI = [-0.05, 0.09]) words / trial, Figure 3.3). This finding is comparable to Experiment 1, again finding that children produce short utterances without much change in length over time.

Convergence

As a coarse measure of partner-sensitivity, we repeated the semantic analysis from Experiment 1. Utterances were more similar if they came from the same partnership (0.27, (95% CrI = [0.243, 0.297])) and were slightly more similar if they came from the same person (0.097, (95% CrI = [0.059, 0.132])).

As Experiment 2 had 4 blocks, we examined whether descriptions were converging semantically toward the final description. We compared the utterances from the first three blocks to the descriptions in the last block using a Bayesian mixed effects linear regression predicting similarity.⁵ Over the first three blocks, descriptions became increasingly similar to the last block description (0.042, (95% CrI = [0.007, 0.078])). Descriptions were more similar if they came from the same child, which is expected as a sign of internal consistency (0.067, (95% CrI = [0.006, 0.127])). Although over time descriptions did get more similar to the last block utterance, the semantic distance between adjacent block utterances was relatively constant (0.009, (95% CrI = [-0.026, 0.044])).

Partnerships often diverged from one another as groups focused on distinct aspects of the image. We tested whether descriptions in different games diverged over time using a Bayesian mixed effects linear regression.⁶ As the games progressed, descriptions to the same target from different games became slightly less similar (-0.013, (95% CrI = [-0.018, -0.008])), which indicates that games are converging to different conventions. These patterns of increasing similarity within games and increasing divergence between games qualitatively match the patterns found for adults (Boyce et al., 2024).

3.4 Joint analysis

As the two experiments were similar to one another, we re-ran models pooling the data across the two experiments, using experiment number as a random effect. Pooling the two experiments, children’s accuracy was above chance (OR: 4.64, (95% CrI =

⁵ $\text{sim} \sim \text{earlier_block.num} + \text{same_speaker} + (1|\text{game1}) + (1|\text{target})$

⁶ $\text{sim} \sim \text{block.num} + (1|\text{target})$




Figure 3.5: Semantic similarity between descriptions from earlier blocks (1-3) and the last block in Experiment 2. Heavy dots are means with bootstrapped 95% CIs; light dots are individual values.

[1.42, 12.16])) and accuracy numerically increased over the course of the game, although the credible interval included 0 (OR of one trial later: 1.04, (95% CrI = [0.98, 1.12])). Descriptions produced by tellers averaged 3.62, (95% CrI = [0.09, 7.38]) words on the first trial, and description length was relatively stable over time (change of 0.04, (95% CrI = [-0.02, 0.11]) words per trial). Utterances were more similar if they came from the same partnership (increase in cosine similarity: 0.257, (95% CrI = [0.234, 0.279])) and were slightly more similar still if they came from the same person within the partnership (0.109, (95% CrI = [0.078, 0.14])).

We might expect that whether a description was successful influences whether the same description, or a variant of it, is employed in future rounds. Intuitively, successful descriptions can be copied and built upon, while unsuccessful descriptions should be replaced by a fresh attempt. To test whether accuracy is predictive of similarity to future descriptions, we ran a post-hoc Bayesian linear model predicting similarity to the next block description in terms of accuracy.⁷ Descriptions that elicited a correct response were more similar to the next block description (0.145,

⁷ $\text{sim} \sim \text{earlier_block.num} \times \text{correct} + \text{same_speaker} \times \text{correct} + (1|\text{game1}) + (1|\text{target}) + (1|\text{expt})$

(95% CrI = [0.055, 0.236])) with no substantial interaction with block number or whether both descriptions came from the same teller. This pattern of results is consistent with the expectation that children are more likely to stick to their own description or repeat the other child's description if it was previously successful.

3.5 General discussion

Prominent early studies claimed that young children cannot overcome their egocentrism to coordinate with each other in reference games (Glucksberg et al., 1966). However, more recent developmental work has found that young children show emerging communicative and pragmatic sensitivity, especially when cognitive demands are low. Here, we revisited the question of preschoolers' performance in child-child reference games, using a scaffolded paradigm to reduce extraneous task demands.

Across 2 experiments and 51 pairs of 4 and 5-year-old children, we tested how well children could produce referential expressions that allowed their partner to find a matching abstract shape. Children varied substantially in what sorts of descriptions they produced, but overall accuracy was high (85%), indicating that children were generally able to produce adequate descriptions. Additionally, children's utterances showed signs of converging toward conceptual pacts. While this task is substantially scaled down relative to measures used for adult competence, it does suggest that the relevant communication skills are present at least in rudimentary form by the end of the preschool years.

Unlike adults, children did not display an increase in accuracy or a shortening of referential expressions over the course of the game. Still, our findings show that descriptions became increasingly similar to descriptions in the last block and that successful utterances were more similar to future utterances, suggesting that children are adapting their descriptions as the game unfolds. These null findings are likely a result of initially high accuracy in the first block and initially short utterances that leave little room for reduction.

It is unclear to what extent the uniformly short descriptions we observed are a product of the simplified task or children's behavioral differences from adults. In this case, the low number of options and relatively easy-to-describe shapes may have

obviated the need for long initial descriptions. Indeed, adult controls in Leung et al. (2024) used shorter initial descriptions than adults in studies with larger arrays of harder to distinguish images (Boyce et al., 2024; Hawkins, Frank, et al., 2020). However, young children may also struggle to produce longer descriptions, and young children may be more willing to take guesses when adults would seek additional clarification. Especially in light of other work suggesting that conceptual pact formation and reduction in utterance length sometimes decouple in adults (Boyce et al., 2024), further empirical work on the factors driving verbosity in reference games is warranted.

The generalizability of our results is limited by the target population, the target images, and the task structure. We sampled a convenience population of children at a university nursery school. The set of tangram images may be easier to distinguish and have higher codability than other target images used in adult reference games. We specifically targeted children’s abilities to construct referring expressions that can be jointly understood, so children were provided scaffolding around taking turns and talking to their partner. Thus, children’s performance should be taken as a proof-of-concept about ability, rather than a claim about how generally children spontaneously demonstrate these abilities.

In the broader picture of language acquisition, there is debate over the timing of the emergence of communicative and pragmatic abilities relative to the acquisition of grammar and meaning. On one side, children seem to learn literal semantics far before they display an understanding of some pragmatic implicatures (Huang & Snedeker, 2009; Noveck, 2001); on the other, sensitivity to communicative intent is an early emerging skill that develops in parallel with linguistic knowledge and may bootstrap language learning (Bates, 1974; Bohn et al., 2019; Tomasello, 2008). Our current findings are most consistent with a gradual development of children’s communicative and linguistic skills, where the skills emerge early and then are refined over time, as children’s cognitive capacities increase. At 4-5 years old, children are already able to establish novel referential conventions with one another as part of their broader ability to communicate and coordinate.

Chapter 4

Idiosyncratic but not opaque: Linguistic conventions formed in reference games are interpretable by naïve humans and vision–language models

4.1 Introduction

When a teen says about someone that “he’s got rizz,” what does this mean? The idea that teen slang is arbitrary and opaque to outsiders (i.e., older generations) is enough of a cultural touchstone that late night comedy shows have segments about it. Teens are far from the only ones to have in-group naming conventions; many communities form stable linguistic conventions including professional jargon, regionalisms, and of course slang.

Iterated reference games are commonly used as a model system for studying the formation of linguistic conventions. In these games, a describer tells their partner how to sort or match a series of abstract images (e.g., Clark & Wilkes-Gibbs, 1986; Hawkins, Frank, et al., 2020). Over repeated rounds of referring to the same targets, pairs develop conventionalized nicknames for the target images. These nicknames

are often partner-specific, in that different pairs develop different nicknames for the same targets. When describing the targets to a new person, describers return to more elaborated descriptions, indicating an expectation that prior conventions are not appropriate descriptions to use for new matchers (Hawkins et al., 2021; Wilkes-Gibbs & Clark, 1992; S. O. Yoon & Brown-Schmidt, 2019a). This partner specificity is generally interpreted to be efficient, implying that conventions are not understandable to those outside of the convention formation process.

For one-shot reference games, the choice of referring expression can be described as a process of recursive inference, in which speakers reason about how listeners will interpret their utterance and vice versa. This process is described in Bayesian pragmatics models such as the Rational Speech Acts model (RSA) (Frank & Goodman, 2012; Goodman & Frank, 2016). The Continual Hierarchical Adaptation through Inference model (CHAI) builds on RSA by adding rules for how agents update their belief distributions after each interaction, to account for the dynamics of repeated interaction in iterated reference games (Hawkins, Franke, et al., 2023). In models like CHAI, both speaker and listener agents are cooperatively reasoning about each other, but related models can be used when the interaction is one-sided, such as when a naïve listener is overhearing a conversation. A key factor permitting variation in referring expressions is lexical uncertainty; RSA models incorporating lexical uncertainty (Bergen et al., 2016; Potts et al., 2015) treat listeners as Bayesian agents who jointly infer the meaning of a speaker’s utterance and their model of the speaker’s lexicon. Models incorporating lexical uncertainty have been used to model children’s word learning (Bohn et al., 2022) as well as person- and group-level differences (Hawkins, Goodman, et al., 2020; Schuster & Degen, 2020). Usually models assume that the person-to-person variation in lexica is small, so the agent is learning a general meaning and a hierarchical set of tweaks to account for slightly different word extensions based on speaker and circumstance, rather than learning arbitrary meanings for each person (Hawkins, Franke, et al., 2023; Schuster & Degen, 2020; but see Misyak et al., 2016).

Conventions in iterated reference games are formed between people without a salient shared group identity prior to the interaction, but people do not use the conventions with new partners. How opaque are the temporary linguistic conventions

created in reference games: are they opaque like “rizz” or interpretable like “round-about”? From a theoretical perspective, measuring this opacity can help inform the development of models like CHAI that aim to account for lexical change in conversation. We attempt to answer this question here.

We consider opacity of references to be a measure of the semantic distance between the referring expression and the target. Transparent expressions (those with low opacity) have signifiers and referents that are semantically close, such that anyone with the same general lexicon can identify the appropriate referent given the signifier. In contrast, expressions that are opaque have signifiers and referents that are semantically distant in the lexicon, such that the relations are arbitrary and inaccessible without additional clues such as the partner-specific conversation history. We note that properties of utterances such as groundedness or compositionality and properties of the target such as iconicity and nameability may load on opacity but are not the same as opacity.

One option for measuring the transparency of referring expressions is to use vision–language models to operationalize a shared semantic space for both language and images. Computational methods have enabled the embedding of various stimuli (including images and text) into high-dimensional feature spaces; these embeddings have properties which suggest that they are reasonable approximations of humans’ semantic spaces, including similarity in representational geometries (e.g., Grand et al., 2022; Muttenthaler & Hebart, 2021). Indeed, embeddings from neural network models have been used as a form of semantics in a range of reference game scenarios (e.g., Gul & Artzi, 2024; Ji et al., 2022; Kang et al., 2020; Le et al., 2022; Ohmer et al., 2022). In particular, such embeddings can be treated as the default context- and speaker-independent lexicon, since they are not updated to account for convention formation within an iterated reference game.

A second option for measuring the transparency of referring expressions is to measure how often naïve humans, who were not part of the group who formed the convention, can correctly associate the target referent with the referring expression. Prior work with naïve matchers has been limited and has focused on the role of conversational history. In this work, naïve matchers tend to do better the more their observation history resembles that of the original conversation—when hearing

descriptions in order instead of in reverse order (Murfitt & McAllister, 2001), when listening to the entire game instead of starting in the third round (Schober & Clark, 1989), and when seeing yoked trials from a single game rather than trials sampled across 10 games (Hawkins, Sano, et al., 2023). In these studies, naïve matchers had worse accuracy than in-game matchers, but their performance was still far above chance, suggesting that the convention–target relationship is not purely arbitrary. In fact, even when pairs of participants try to obfuscate their meaning, overhearers can still do quite well at identifying the target referents (Clark & Schaefer, 1987). Nonetheless, receiving more context from an interaction—and in particular having that context be in order—is beneficial to matchers. Except for the shuffled condition of Hawkins, Sano, et al. (2023), these studies do not address how opaque descriptions are when they are presented without any prior context. Such an approach would be important to have truly naïve matchers that lack even the context of trial order, providing a clearer understanding of when these expressions are opaque and when they are merely idiosyncratic.

In the current work, we measure the opacity of the referring expressions created in iterated reference games. We use conversations from Boyce et al. (2024), who created a large corpus of iterated reference games that were played online using a chatbox for communication. This corpus is made up of 6-round iterated reference games using the same 12 target images. Games varied in how large the describer–matcher groups were (2–6 participants) and how “thick” the communication channels were (for example, if matchers could send back messages or just emoji). The varied conditions within a consistent framework allowed us to test how the opacity of referring expressions varies depending on the conditions the referring expressions came from.

We use both human experiments and models to assess when and why expressions are opaque or understandable to outside observers. We first present a computational approach using a vision–language model to measure the semantic similarities between referring expressions and their targets, and we validate our model against naïve human matchers (Experiment 1). We then use both naïve human matchers and the model to compare the opacity of referring expressions across different game conditions and time points (Experiment 2). Finally, we address the role of conversation history by comparing naïve matcher performance on game transcripts presented in order versus



Figure 4.1: Experimental setup. Naïve matchers read transcripts from trials in reference games from Boyce et al. (2024) and selected which image they thought was being described. Matchers received bonus payments for correct selections.

out of order (Experiment 3).

4.2 Task setup

4.2.1 Materials

We drew our referring expressions from Boyce et al. (2024), excluding utterances that were marked as not containing referential content. For our naïve matcher experiments, we sampled different subsets of this corpus. Within the subsets, we excluded transcripts that contained swear words or crude or sexual language. For the computational model, we used the entire corpus, and pre-processed the text by concatenating all the referential messages sent by the describer for a given trial.

4.2.2 Experimental procedure

We recruited English-speaking participants from Prolific. On each trial, participants saw the full transcript from that trial, containing all the chat messages marked by whether they were from the speaker or a listener. Participants selected the image they thought was the target from the tableau of 12 (Figure 4.1). Participants received feedback on whether they were right or wrong on each trial. Except when the

specific viewing order was part of the experimental manipulation, we randomized the order of trials, subject to the constraint that the same target could not repeat on adjacent trials. The task was implemented in jsPsych (de Leeuw et al., 2023). We paid participants \$10 an hour plus a bonus of 5 cents per correct response. All our experimental code is at this anonymized repo.

4.2.3 Computational models

We used the Contrastive Language-Image Pretraining model (CLIP; `clip-vit-large-patch14`) as a comprehender model for our domain (Radford et al., 2021). CLIP is a vision-language model that uses a text transformer and a vision transformer to embed text and images into the same space, trained to maximize the similarity between representations of images and their English captions. It is a natural choice for reference games, as the model is trained to estimate the correspondence between images and phrases in natural language. We ran CLIP for the concatenated describer utterances and all 12 tangram shapes. For each utterance, we computed probabilities for each tangram shape using logit scores from CLIP. The simplest way to do this is simply taking the softmax of the logits. However, tangram shapes are outside of the training distribution for the model, perhaps explaining why it favored some images over others regardless of the content of the text.

To improve the performance of base CLIP, we trained a set of readout models to assign probabilities to images using CLIP’s logits as features. Models were trained to maximize task performance (i.e., to assign high probability to the target tangram given the concatenated describer utterance). We compared four types of models: random forest, logistic regression, multi-layer perceptron (MLP), and gradient-boosted tree. Classifiers were implemented in the `scikit-learn` and `XGBoost` libraries (Chen & Guestrin, 2016; Pedregosa et al., 2011). Each readout model was evaluated using 10-fold cross-validation, where the model was trained on 90% of the data and evaluated on the remaining 10%. Table 4.1 shows the cross-validated accuracy of different readout models, as well as the performance of CLIP with no readout. The MLP with two hidden layers of size 100 performed the best on held-out data; in subsequent analyses, we use the MLP trained on all the data.

Table 4.1: Cross-validated accuracies for classifiers. Standard deviations in accuracy across the 10 folds are shown in parentheses. Best performance within each model class is underlined, and best overall performance is bolded.

Classifier	Accuracy
Random baseline	0.08
CLIP without readout	0.31
Logistic regression	
No penalty	<u>0.50 (0.01)</u>
L2 penalty	0.50 (0.01)
Random forest	
10 estimators	0.46 (0.02)
50 estimators	0.51 (0.02)
100 estimators	0.52 (0.02)
500 estimators	<u>0.52 (0.02)</u>
Gradient-boosted tree	
10 estimators	0.48 (0.02)
100 estimators	<u>0.51 (0.02)</u>
Multi-layer perceptron	
1 $\times$ 32-dim hidden layer	0.50 (0.01)
1 $\times$ 100-dim hidden layer	0.52 (0.01)
1 $\times$ 512-dim hidden layer	0.53 (0.02)
1 $\times$ 1028-dim hidden layer	0.53 (0.02)
2 $\times$ 32-dim hidden layers	0.51 (0.02)
2 $\times$ 100-dim hidden layers	<u>0.55 (0.02)</u>

4.3 Experiment 1

Our CLIP-MLP computational model was optimized for task accuracy. To validate whether this objective also results in human-like response patterns, we conducted a calibration experiment to determine if model-assigned target probabilities were aligned with the probabilities that naïve human matchers would choose particular target images across a range of utterance-target pairs.




Figure 4.2: Correlation between human accuracy and CLIP-MLP probability of target in Experiment 1. Small points are individual descriptions, colored by decile of CLIP-MLP probability, large points and error bars are the bootstrapped mean and 95% CI across descriptions for each decile.

4.3.1 Methods

We first obtained target probabilities from our CLIP-MLP model for all utterances from Boyce et al. (2024). We then used stratified sampling to select 217 trials by dividing model-predicted probabilities into deciles and choosing approximately 22 utterances per decile, spanning the 12 different possible target images. We recruited 61 participants who each saw 64 trials randomly sampled from the 217 tested trials. On average, each trial was seen by 18 participants. This experiment was pre-registered at this anonymized link.

4.3.2 Results and discussion

We obtained human accuracies on each trial by dividing the number of participants who selected the target by the total number of participants who saw the trial (Figure 4.2). There was a modest but significant positive correlation between model-predicted probabilities and human accuracies ($r = 0.33$ [0.21, 0.45]). This result suggests that




Figure 4.3: Accuracies for naïve human matchers and the CLIP-MLP model for Experiments 2a and 2b, grouped by the source of the referential description. Facets are the communication thickness of the original game and x-axis is when in the game the transcript came from. Point estimates and 95% CrI are predictions from the fixed effects of logistic and beta regressions. Bootstrapped mean accuracy from the original matchers is included as a ceiling, and random chance as a baseline.

model predictions were calibrated to human response patterns, albeit not perfectly. It is possible to use these calibration results to tune model predictions to better approximate human responses; we leave this approach for future work. Nonetheless, the observed positive correlation suggests that our computational model carries some signal about human accuracies, validating its use in subsequent experiments as a computational comparison.

4.4 Experiment 2

As a starting point for examining the opacity of referring expressions, we focused on referring expressions from the first and last rounds of reference games. Based on the idea that conventions form across repeated communication, later-round utterances should be more opaque. To test this hypothesis, we ran a recognition experiment




Figure 4.4: Accuracies for naïve human matchers and the CLIP-MLP model for Experiments 2a and 2b, split out by target image. Point estimates and 95% CI are predictions from the fixed effects and by-tangram random effects of logistic and beta regressions, bootstrapped across conditions. Bootstrapped mean accuracy from the original matchers is included as a ceiling, and random chance as a baseline.

including descriptions from games of different sizes and communication thicknesses. Based on the patterns of cross-game similarity in Boyce et al. (2024), we expected that smaller and thicker games, whose descriptions diverged fastest, would have more idiosyncratic and opaque conventions than larger groups with thinner communication channels.

4.4.1 Methods

Experiment 2a

To establish a baseline of how well naïve matchers could understand descriptions without context, we ran a 2×2 within-subjects experiment, drawing target transcripts from 2- and 6-player games from Experiment 1 of Boyce et al. (2024) and from the first and last blocks of these games. These games had medium-thick communication channels, where matchers could send text messages to the chat, the describer

role rotated each round, and matchers received limited feedback. We recruited 60 participants who each saw 60 trials (15 in each of the 4 conditions). Overall, participants saw 774 transcripts from 40 games. This experiment was pre-registered at this anonymized link.

Experiment 2b

After observing limited condition differences in Experiment 2a, we ran a follow-up experiment on descriptions from Experiment 3 of Boyce et al. (2024), where the communication channel thicknesses were more extreme. Here, we used a $2 \times 2 \times 2$ within-subjects design, drawing our transcripts from the first and last rounds of thick and thin, 2- and 6- person games. In the “thick” condition, matchers could send text messages to the chat, one person was the describer for the whole game, and matchers received feedback on everyone’s selections. In contrast, in the “thin” condition, matchers could only communicate by sending 4 emoji, the describer role rotated, and matchers received limited feedback. As the emoji did not have referential content, we did not include them in the transcripts shown to naïve matchers. For experiment 2b, we recruited 60 participants who each saw 64 trials (8 in each of the 8 conditions). Overall, participants saw 2392 transcripts from 163 games. This experiment was pre-registered at this anonymized link.

4.4.2 Results

Experiment 2a

For Experiment 2a, we ran a Bayesian mixed-effects logistic model of naïve matcher accuracy in `brms` (Bürkner, 2018).¹ For this and other models, we report the model estimate and 95% credible interval. Overall, naïve matchers were right 62% of the time, far above the $1/12 = 8.3\%$ expected by random chance ($OR = 1.93$, (95% CrI = [1.05, 3.62])). There were no large effects of condition (Figure 4.3 middle panel). Participants tended to be less accurate at descriptions from the last round (OR of last round = 0.77, (95% CrI = [0.53, 1.1])). There was no clear effect of original

¹`correct ~ group_size × round + trial_order + (group_size × round | correct_tangram) + (group_size × round + trial_order | workerid)`

group size (OR of 6-player game = 1.15, (95% CrI = [0.89, 1.47])), but there was an interaction between round and group size (OR = 1.49, (95% CrI = [1.06, 2.1])). Later transcripts from larger games were easier to understand, but earlier transcripts from smaller games were easier to understand. Much of the variation in accuracy was driven by the target image, which accounted for more variation than participant differences (standard deviation of image distribution = 0.98, (95% CrI = [0.63, 1.51]); SD of participant distribution = 0.64, (95% CrI = [0.42, 0.88])). Some images were much easier to identify as the target than others (Figure 4.4).

Experiment 2b

For Experiment 2b, we ran a similar Bayesian mixed-effects logistic model.² Naïve matchers were above chance (OR = 1.81, (95% CrI = [1.06, 3.08])), Figure 4.3). As in Experiment 2a, there were not substantial effects of condition. Last-round descriptions had slightly lower accuracy (OR of last round = 0.64, (95% CrI = [0.47, 0.85])), but there was an interaction with thickness, where for thin games, last round descriptions were less opaque (OR = 1.55, (95% CrI = [1.02, 2.33])). Again there was strong variation based on target image (0.81, (95% CrI = [0.51, 1.28])), which exceeded by-participant variation (0.62, (95% CrI = [0.43, 0.83])).

Additional predictors

We considered the accuracy of the in-game matchers from Boyce et al. (2024) and the length of the description as post-hoc predictors. In both experiments, in-game accuracy was predictive of naïve matcher accuracy (Expt 2a OR = 3.33, (95% CrI = [2.45, 4.53]), Expt 2b OR = 2.39, (95% CrI = [1.88, 3.03])). The log number of words in the description was not predictive in Experiment 2a (OR = 1.05, (95% CrI = [0.94, 1.17])), but longer descriptions were slightly beneficial in Experiment 2b (OR = 1.1, (95% CrI = [1.01, 1.2])).

The pattern of which conditions became more opaque in later rounds resembled the pattern of which conditions produced descriptions that diverged the most

²correct ~ group_size × thickness × round + trial_order + (group_size × thickness × round | correct_tangram) + (group_size × thickness × round + trial_order | workerid)

in semantic space in Boyce et al. (2024). As a post-hoc test of whether opacity might be related to semantic divergence from other descriptions, we used the mean semantic similarity between an utterance and other utterances in the same condition as an additional predictor of accuracy.³ Similarity to other utterances was strongly predictive of increased accuracy in both experiments (Expt 2a: OR = 12.49, (95% CrI = [4.7, 33.64]), Expt 2b: OR = 14.75, (95% CrI = [5.93, 36.87])) and was more predictive for the last round descriptions (Expt 2a: OR = 3.49, (95% CrI = [1.09, 10.92]), Expt 2b: OR = 4.78, (95% CrI = [1.61, 14.25])). While exploratory, this analysis suggests that referring expressions that are further from shared semantic priors (i.e., more idiosyncratic) are harder for naïve matchers to understand.

4.4.3 Model results

As a computational comparison, we used the probability the CLIP-MLP model assigned to the correct target as our dependent measure and fit a Bayesian mixed-effects beta regression on the descriptions from Experiment 2.⁴ The CLIP-MLP model was far above chance, but had lower accuracy than the human participants (OR = 0.6, (95% CrI = [0.45, 0.82])). The strongest predictor of accuracy was later round (OR = 1.32, (95% CrI = [0.94, 1.83])), but even this was uncertain. There was substantial by-target image variation (SD = 0.46, (95% CrI = [0.27, 0.76])).

In additional models, we checked the effect of in-game matcher accuracy, length of the description, and semantic divergence. CLIP-MLP had higher accuracy when in-game matcher accuracy was higher (OR = 1.52, (95% CrI = [1.35, 1.71])), and when descriptions were shorter (OR for log words = 0.85, (95% CrI = [0.82, 0.9])). The model may perform poorly on long descriptions because they are further from the model’s training distribution of image captions. A description’s semantic similarity to other descriptions was predictive of higher accuracy (OR = 11.85, (95% CrI = [6.51, 21.06])), especially for last round utterances (OR = 3.46, (95% CrI = [1.65, 7.36])), in line with the human results.

³Semantic similarity was operationalized as cosine similarity between S-BERT embeddings (Reimers & Gurevych, 2019), the measure of semantic distance used in Boyce et al. (2024).

⁴ $\text{correct} \sim \text{group_size} \times \text{thickness} \times \text{round} + (\text{group_size} \times \text{thickness} \times \text{round} | \text{correct_togram})$

4.4.4 Discussion

Overall, naïve human matchers were fairly accurate overall, but less accurate than matchers in the original game, consistent with prior work. The computational model was less accurate, but still far above chance. The largest source of variability in accuracy was from target images, and whether earlier or later utterances were more opaque varied by game condition. The level of semantic divergence from other expressions was strongly predictive of the opacity of the expression. While this analysis was not prespecified, it still provides some suggestion that descriptions that were closer to shared semantic priors were also more interpretable.

4.5 Experiment 3

The experience of naïve matchers in Experiment 2 differed from in-game matchers in several ways; any of these could explain differences in accuracy. In-game matchers received descriptions from a consistent group, in the order they were created, and were the intended audience of the descriptions. In Experiment 3, we focused on the role of context and group-specific interaction history to tease apart some of these differences. One question of interest was how much seeing the entire the conversation history in order would increase the interpretability of later round descriptions.

Additionally, we were interested in how quickly naïve matchers process the linguistic descriptions and how this might relate to the predictability and informativity of the descriptions. Word predictability (often measured as surprisal) is a predictor of language processing times. The iterated reference game transcripts present an interesting test case because expressions are likely to be globally uncommon (and thus unpredictable), but may be licensed and unsurprising given the immediate visual context and conversation history. Thus, another question of interest is how rapidly comprehenders are able to adapt to the changing context.

4.5.1 Methods

We compared naïve matchers in yoked and shuffled conditions. In the yoked condition, naïve matchers saw all the descriptions from a single game in the order they originally




Figure 4.5: Accuracies for Experiment 3. Error bars are bootstrapped 95% CIs.

occurred. In the shuffled condition, naïve matchers saw all the descriptions from a single game in a randomized order.⁵

Because some descriptions are already fairly comprehensible in isolation, we focused on games that showed strong group-specificity. We hand-picked 10 games from Boyce et al. (2024) on the basis of high in-game matcher accuracy, strong patterns of descriptions shortening over repetition, and the use of idiosyncratic or non-modal referring expressions. Thus, these games showed the hallmarks of strong conventionalization to terms that were more likely to be opaque to outsiders.

We recruited 196 participants (99 in the yoked condition and 97 in shuffled) who each saw all 72 trials of one of the 10 games. This experiment was pre-registered at this anonymized link. Participants read the transcripts in a modified self-paced reading procedure where they uncovered the text word-by-word (revealed words stayed visible); only after uncovering the entire transcript could participants select an image.

⁵For clarity, we note this is the same yoking but a different shuffling than that used in Hawkins, Sano, et al. (2023).

4.5.2 Results and discussion

Accuracy

One primary question of interest was how much seeing the conversation history unfold in order would help participants interpret descriptions, especially those from later rounds.

We compared accuracy across the yoked and shuffled conditions with a Bayesian mixed-effects logistic regression.⁶ The descriptions were more transparent when they were presented in a yoked order (OR = 2.2, (95% CrI = [1.63, 3]), Figure 4.5). In the shuffled condition, there was no main effect of round number (OR for one round later = 0.99, (95% CrI = [0.95, 1.02])), but there was a marginal interaction where the benefit of the yoked condition decreased for later rounds (OR for one round later = 0.94, (95% CrI = [0.89, 1])). This was offset by matchers in both conditions improving at the task over time (OR for one trial later in matcher viewing order = 1.02, (95% CrI = [1.02, 1.02])). In the yoked condition round and trial number were aligned, so an improvement over time could be either from matcher practice or from descriptions being easier to understand. In the shuffled condition, matcher practice effects did not correlate with position in the original game.

Comparing to the performance of in-game matchers, we separated out the benefits of seeing the descriptions in order versus being a participant in the group.⁷ There was a benefit to seeing the items in order (OR = 2.24, (95% CrI = [1.63, 3.04])) and a larger benefit to being a participant during the game (OR = 4.35, (95% CrI = [2.77, 6.89])). The benefit of seeing the items in order waned in later blocks (OR = 0.94, (95% CrI = [0.89, 1])), but the benefit of being in the game did not (OR = 1.06, (95% CrI = [0.95, 1.18])). In all cases, there was a baseline improvement over trials (OR = 1.02, (95% CrI = [1.02, 1.02])). As a caveat, we note that in-game matchers and naïve matchers may have varied from each other in terms of effort and time spent on the task, and thus the comparison should be interpreted cautiously.

The accuracy of the CLIP-MLP model was worse than the shuffled human results,

⁶ $\text{correct} \sim \text{orig_repNum} \times \text{condition} + \text{matcher_trialNum} + (1|\text{gameId}) + (1|\text{correct_tangram}) + (1|\text{workerid})$

⁷ $\text{correct} \sim \text{orig_repNum} \times \text{order} + \text{orig_repNum} \times \text{setting} + \text{matcher_trialNum} + (1|\text{gameId}) + (1|\text{correct_tangram}) + (1|\text{workerid})$




Figure 4.6: Schematic of informativity predictors used in the incremental reading time analysis.

and did not change across rounds (OR for one round later = 1.02, (95% CrI = [0.97, 1.07])). The larger difference between naïve human and CLIP-MLP accuracies in Experiment 3 than Experiment 2 suggests that the shuffled ordering still provides useful context that helps matchers understand the conventions. This history was not available to the CLIP-MLP model which saw every description as a one-shot task.

Incremental reading times

Our second question was whether comprehenders would rapidly adapt to the task, which we addressed by considering different potential predictors of word-by-word reading times that varied on whether they took in local context. As a baseline, we included word length and log frequency. We also considered three measures of a word’s informativity: **1**) surprisal from a base large language model (LLM, LLama 3.1 8B) (Grattafiori et al., 2024), **2**) surprisal from a vision language model (VLM, LLama 3.2 11B) (Grattafiori et al., 2024), conditioned on grid of 12 possible targets, and **3**) the word-wise change (operationalized as KL divergence) in predicted target distribution from the CLIP-MLP model discussed above (Figure 3). This last measure represents the task-relevant information content of a word based on how much it shifts the model’s distribution over possible targets. We residualized the measures of




Figure 4.7: Estimates and 95% credible intervals for coefficients in the incremental processing model. Word-level predictors are shown in panel A, colored by spillover position. Condition effects are shown in panel B.

informativity and normalized all word-level predictors, and then modeled RTs as a function of these predictors for the current and 3 prior words, in addition to condition-level predictors (Figure 4.6).⁸

Model results are shown in Figure 4.7. As expected, shorter and more frequent words were faster to read. Unexpectedly, current word LLM surprisal was not predictive of reading time, and both VLM surprisal and KL divergence were predictive of shorter reading times, in contrast to the typical positive relationship between surprisal and RTs. LLM surprisal, VLM surprisal, and KL all measure different sources of information with increasing task-specificity, and may have different impulse response functions, leading to different patterns of lag.

Overall, participants tended to read faster in the yoked condition than the shuffled condition, indicative of context having a facilitative role on reading time. Participants also tended to get faster over the course of the task, which could be either adaptation

⁸The model formula was $RT \sim [word-length \times frequency + LLM + VLM + KL] + original-block \times condition + order-viewed + (1 | game) + (1 | target) + (1 | participant)$. Bracketed predictors were included for current word and 3 prior words to account for spillover.

to the task, or an increasing comprehension of the descriptions (which would be supported by their increasing accuracy over time). Participants tended to read more slowly per word for descriptions coming from later original blocks; however, this measure is confounded both by viewing order being correlated with block for the yoked condition, and by reduced representation of short descriptions in the reading time data, due to the need for spillover predictors.

The reading time paradigm we used is noisy and prone to spillover effects. For user interface reasons, we did not require a button press after the last word to enable target selection, so we do not have RTs for the final words. Additionally, our analytic approach meant that we only considered words where we had predictors for the three previous words (to account for spillover); however, this does mean that we do not have usable data for utterances that were 4 words or fewer. We do think we were getting some signal as the RTs were correlated with word length and word frequency, but unfortunately, we are unable to interpret the measures of information because of the noisy paradigm.

4.6 General Discussion

Real-world conventions vary in whether they are opaque to outsiders (“rizz”) or interpretable even to those who don’t produce them (“roundabout”). Convention formation in the real world is difficult to study, so iterated reference games are a method for operationalizing convention formation for experimental study. In reference games, conventions are partner-specific: different groups’ evolving conventions follow different paths through semantic space. Despite the number of studies on iterated reference games, few studies have examined whether the descriptions are interpretable by outsiders.

Across multiple experiments, we found that naïve human matchers were far above chance at identifying the targets, and our computational model was also above chance. For both humans and models, more variation was explained by the target image than the round or game condition the descriptions came from, suggesting that conventionalization was not the primary driver of how difficult an expression was to interpret.

Even for games selected for strong conventionalization, naïve matchers had high accuracy overall, although this accuracy was further increased if they saw the conversation history in order. Exploratory analyses also suggested that more idiosyncratic descriptions were more opaque. Our findings are consistent with a lexical uncertainty approach, where expressions that are closer to overall priors are easier to understand, and groups that have fewer people and thicker channels are more able to break away from these priors and have conventions that drift farther apart in semantic space.

This work suggests that conventions formed within a small group may still be fairly comprehensible to those outside the group, as the semantics of these conventionalized expressions may still be relatively close to the semantics of the image. Hence, such conventions may not be completely arbitrary, and are instead strongly related to the image and the shared general lexicon, although it is possible that greater arbitrariness and opacity may be obtained with more extended interaction.

While this first attempt to combine incremental sentence processing methods with descriptions from iterated reference games was not very informative, it points out that this genre of methods could be used to address questions of how both visual and conversational context influence word predictability and in turn the speed at which words are integrated in online sentence processing. Further work could refine the measures using more sensitive techniques such as eye-tracking, analyses that better account for time course, and measures of information that take prior informativity into account. Beyond reading time measures, one could also consider indexing informativity more directly with a gated paradigm that allows for periodic guesses of the target to provide human measures of incremental task-informativity.

Limiting the generality of our findings, our experimental and computational results were only on a specific set of iterated reference game transcripts and images. Some images may lend themselves to more transparent descriptions because they are more iconic, with a narrower prior over different ways they could be conceptualized, or they may be further from competitors within this pool of images. Future work sampling across larger sets of images (such as Ji et al., 2022) could probe image-level factors.

Nonetheless, this work has demonstrated the utility of adopting a broader perspective on convention comprehension. In particular, the use of computational modelling allowed for a means to estimate semantics under lexical uncertainty without requiring symbolic semantic representations. Future work could capitalize on this approach to better understand semantic dynamics underlying convention formation, as well as provide further quantitative investigations of pragmatics models like RSA and CHAI. These directions will help us to better understand the nature of reference and conventions, and how humans navigate the complex and ever-evolving landscape of communication.

Conclusion

This dissertation started with questions of robustness and generalizability of the phenomena associated with convention formation in iterated reference games, with a goal of understanding whether these phenomena were “special”, necessitating their own theoretical explanations, or whether they could be accounted for as non-central examples within broader-coverage theories. Our approach was to run iterated reference game experiments over a wider set of parameters, to probe how well empirical evidence lined up either with the special verbal theories or a wider terrain of smoother variability.

Synthesizing across experiments, one observation is that the central findings of the iterated reference game literature are quite robust. People very frequently and naturally create shorthand names for targets as a way to effectively communicate about them, and which nicknames are used vary widely, with some popular concepts shared across many groups, but also a long tail of idiosyncratic, yet effective, terms. Initial descriptions are usually long, evolving over time into shorter expressions. With “thick” communication channels, similar to those used in much of the literature, but mediated by text instead of speech, the trends are very consistent with prior findings from the in-person dyadic literature. People are just good at communicating and good at coordinating with each other.

The interesting part is that some of the classic patterns occur even under more limiting circumstances, even beyond what seems to be predicted from the informal theories. For instance, at least for pairs, people do just fine without much in the way of backchannel. Even for larger groups of 6 people without backchannel and with limited knowledge of performance, accuracy increases and utterances shorten. The descriptions produced under these circumstances are so effective that even naïve

matchers get them right more than 50% of the time on 12-way choice. In one view, this is robust evidence for our ability to produce and interpret language effectively. Another view is that common ground and audience design can't be so important if people can succeed without very much of it.

We find robust evidence of the path-dependency of what referring expressions are used. Many different descriptions can be effective, and dyads find many of them. Children display similar patterns in miniature – pairs can use many different conceptualizations effectively. However, most of them can be understood reasonably well by others, and many of them are reasonably transparent to models of language statistics. Thus, while the formation of referring expressions may be idiosyncratic and depend on the partner, the interpretation of the expression is less dependent. It is as if people start out pluripotent in the set of associations they will accept and the referring expression process mostly is of picking one association, with rarer instances of creating new associations beyond the shared pool.

Finally, a wider range of circumstances complicates the use of reduction as a proxy for partner-specific convention formation. In larger groups with constrained communication channels, the degree of group-specificity of referring expressions seems to uncouple from the length of the utterances. In a different uncoupling, children do show some partner-specificity but do not show a decrease in expression length.

Overall, we do find quantitative patterns where more communication and more context generally leads to better performance and faster conventionalization. It's easy to come up with a theory that predicts that directionally. The hard part is determining *how much* more supportive some circumstances should be than others.

So, what should we think about human communicative capacities? Some of the ideas that people are really attuned to conversation with a partner and say just the right things for their partner to understand probably aren't true, at least in the strong form. Which is not at all to say that people aren't great at communicating – we find instead that the communication abilities are wide and flexible. This work opens up new questions: how do we explain the path-dependence of people's utterances, if it isn't driven by necessity and efficiency? How much of the path-dependency and referring expression length are driven by production constraints? How do we model the generation process that occasionally leads to very idiosyncratic descriptions and

more often results in less idiosyncratic descriptions? How do these patterns relate to iconicity and group dynamics?

While the work presented in this dissertation pushes forward our empirical understanding of adaptive language use in iterated reference games, there are a set of limitations that should be noted. Some of these limitations apply more generally to the literature as a whole, and thus point to opportunities for future advancement.

First, it must be noted that the generalizability of the work is limited by the generalizability of the stimuli, the participant populations, and the paradigm. We used a limited set of 12 vaguely humanoid tangrams in Chapters 2 and 4 and a different set of 4 vaguely humanoid tangrams in Chapter 3. The literature as whole has more variability, so we have no reason to believe that our results are contingent on the stimuli, but features of stimuli are very likely to influence the results via their levels of iconicity, affordances for describing strategies (pose, geometry, body parts), and levels of codeability and discriminability. We are limited in our ability to assess variation due to stimuli because of the small number of stimuli that we sampled.

Our studies were all conducted in English using convenience populations of adults recruited from Prolific and children recruited from a university preschool. From the literature we have some evidence that the similar patterns of key convention phenomena do occur in French and Dutch as well. To the extent that the paradigm is accessing some general communicative and linguistic abilities, we should expect the general patterns to generalize across demographic factors, languages, and cultures. However, different languages may have different affordances for how to identify unknown referents and how to shorten and develop nicknames. Typological variation in language might lead different parts of speech to be more common in reduced utterances, for instance.

As we transition from documenting that humans can use language to communicate under these circumstances and become increasingly concise to trying to model how humans do this, features of the language itself may become more relevant. One blind spot in the literature is with signed languages, which have some different affordances related to iconicity in that modified meanings can be expressed either with additional signs or with modifications to the head signs (Ferrara et al., 2025).

From a social perspective, we might suspect that social distances, social hierarchies, and social norms could all constrain the interactions and patterns of language use in reference games. We put people in an artificial game environment with specific communication goals, but participants are interacting with people and bring their social knowledge with them. Thus, there are unmeasurable expectations about how conversation works and unmeasured goals for social connection that could be influencing all of our measures. One could imagine that people with closer connections or different social mores might show some different patterns and quantitative variation. We might be somewhat reassured by general robustness of the underlying patterns to modalities with different social norms (in-person, by phone, and by web-chat) and the fact that some studies don't find differences between friends and strangers. However, while suggestive of some generalizability, we have not tested how social factors are influencing game play. Shared society and culture might be particularly relevant to our naïve matcher results since guessing the target may depend on shared conceptual interpretations and priors over what the images could represent.

Another concern with this work is that we do not know the causal mechanisms behind the dependent variables that are measured. Thus, we also do not know what moderators are relevant and what dimensions we expect results to generalize over. Current statistical models used to assess reduction (or disfluency rate, or speed, etc) have the issue that they are fit to a specific dataset and can't be compared across datasets and situations, so it's difficult to assess their quantitative generality (Yarkoni & Westfall, 2017). Developing formal theory, even at a computational level, is difficult as many dynamics and equations could suggest the same directional outcome (van Dongen et al., 2025). It is particularly complicated here where the measurable outcomes are interrelated and are not the actual target of interest. To think about a generating model, we would need to determine what latent constructs are core invariants and which variables are affected by implementational details.

A final consideration is how ecologically valid the experimental paradigm is. To what extent are iterated reference games a useful model of natural human communication that captures something relevant about convention formation? There are various dis-analogies between normal conversation and reference games. In real life, many interactions are more than just about reference and have a more open set of

possible referents. This may be an acceptable simplification, and in other work, we do find suggestions that referential reduction and alignment can occur in coordination games (Boyce & Frank, 2023; Mankewitz et al., 2021). In real life, many of our interactions are with friends or other people who we interact with regularly, unlike in reference games where participants are generally strangers. Thus, the amount of shared history and the timeline of accumulating shared history is quite dissimilar between friendships and reference games. It is possible that the differences in type and distribution of shared context could lead to different types of convention formation. I do not think iterated reference games are a complete model of all linguistic convention formation, but I do think it is a sufficiently good analogy, especially for more temporary conventions between people who do not have years-long history together. Iterated reference games certainly can tell us something about how people use language to do a task that has an atypical distribution of communicative needs.

In ongoing work, I am using a couple of new approaches to partially address some of the limitations discussed here. The first method is to aggregate and harmonize data from multiple existing datasets to get some generalizability and answer questions that cannot be answered in a satisfactory way with only one dataset. The other is to use tools from natural language processing to explain semantic and lexical changes in referring expressions over the process of conventionalization.

One question that currently cannot be adequately answered by the data presented here is what is the functional form predicting the number of words used as a function of repetition. We could try to fit that model, but we would not be learning anything generalizable since we only have 6 time points and the entire dataset shares the same set of targets and some other paradigm properties. It’s not even clear that repetition is the right predictor versus something more like how much to reduce based on what informational profile the last utterance had. However, by using a larger set of data with more variability in experiment length, number of options, paradigm details, and stimuli, we might be able to more meaningfully test a set of different functional form relations and have some confidence that we were not just overfitting to irrelevant design choices.

Other questions best addressed across multiple datasets include the robustness and generalizability of the key phenomena. Different papers report somewhat differing sets

of outcomes and operationalize outcomes differently even when the same construct is of interest. Thus, we might know that the directionality of effects is consistent across set-ups, but we don't know how much variability there is, either based on moderators, or just in general.

One approach to pooling effects and studying moderators is meta-analysis, which combines like effect size measures across different exact experimental implementations. However, given the different measures of interest and different ways of measuring change over time, for iterated reference games, the numerical results are not easy to compare. Instead, we are working on harmonizing raw data across different datasets, in order to do mega-analytic statistical measures where the individual data points all go into the same regressions, with data source as a hierarchical random effect.

In addition to studying potential moderators of interest, pooling data across datasets also allows us to check for robustness to factors that we have posited as unlikely to be crucial. For instance, while the conditions reported in Chapter 2 all shared a text-based chat interface and the same set of target stimuli, other experiments from other researchers have used face-to-face oral paradigms and line drawing stimuli. Thus, applying a mega-analytic approach across multiple datasets allows for estimating the effect from these moderators. In addition to getting robust estimates and measures of generalizability, this could also potentially inform future experimental design, both by pointing out gaps in the explored experimental space, but also identifying what dimensions of experimental variation are most important. For instance, one could imagine that type of stimuli has more interacting effects on number of words whereas written versus spoken might just cause an overall decrease in verbosity without any interactions. One might then decide to use the more convenient modality but focus on varying experimental stimuli more systematically.

The other future direction we are pursuing is applying NLP tools to try to get a better grasp on the linguistic patterns in iterated reference games. We believe that the changes observed are probably operating more on semantics and information of the words than on metrics like length. While we don't have the models we would need for a generative model of iterated reference game utterances, we do have the tools to more closely look at semantic changes using embeddings. For instance, we

might be able to fine tune a model on utterance–selection pairs and use it as a proxy for matchers to then do attribution of information to words within the utterance. Determining what the information profile of utterances is and whether informativity is a strong predictor of what does or doesn’t get reduced is a way we can use the tools now available to answer scientific questions about the process of reduction. Since this process involves model training, we need as much language from reference games as we can get. The more data and data diversity we have, the better and more robust a model we can tune.

There are many possible approaches to understanding linguistic communication. In the dissertation work, we have built on many previous studies of language from a variety of sub-disciplines. We’ve drawn from the study of conversation, from sentence processing, and from computational linguistics. Beyond language research, we’ve also drawn on cognitive science and developmental psychology. Advances in computational methods and experimental methods have enabled this work, and further theoretical and technical connections will lead to further insights. Understanding language deeply requires explaining patterns of language on different time scales and levels of abstraction. I look forward to continuing to explore language use in its duality as a communicative modality and a mental process.

Appendix A

Feedback

Figure A1 shows the type of feedback shown to participants in different conditions.

Semantic content of descriptions

As an exploratory, qualitative analysis of the describer’s language, we performed dictionary searches to classify whether the describers utterances on each trial contained the following types of words.

Geometric/literal words: squar*, triangle, triangular, diamond, shape, trapez*, angle, degree, parallel*, rhomb*, box, cube, line, white, black

Words for body parts: face, head, heads, back, shoulder, shoulders, arm, arms, leg, legs, foot, feet, body, knee, knees, toe, toes, hand, hands, body, butt, heel, heels, ear, ears, nose, neck, chest, hair

Positional words: right, left, above, below, under, over, top, bottom, behind, side, beneath

Posture words: kick*, crouch*, squat*, kneel*, knelt, stood, stand*, sit*, sat, lying, walk*, facing, fall*, looking, lean*, seat*, laying

Qualitatively, descriptions that contain none of these words are usually shorter, more abstract, more conventionalized descriptions. Figure A2 shows the fraction of trials where the speaker used these words. Generally, the more concrete word categories become less frequent over time, although this effect seems weaker in larger, thinner games.




Figure A1: Feedback shown to participants. In low-coherence games, matchers saw feedback as shown in A and B. In all games, describers saw feedback like in C; icons were associated with players and indicated who selected which. In high feedback games, matchers saw grids like C, but with text as in A and B.




Figure A2: Fraction of trials where the describer used certain types of words. See text for word lists. Faint dots are per tangram, larger dots and error bars are means and bootstrapped 95% CIs.

Additional game transcripts

Table A1: Additional examples from 6-player thick games for the same image across repetitions. Describers are indicated with an asterisk.

6-person thick game	
<i>Rep 1: 5/5 correct</i>	
Q*	oh no. okay. this one is similar to the one i called minnesota, so like pretty simple, with just a distinct "head" diamond and then a body beneath
Q*	the body has sort of sloping shoulders
P	BEAN SHAPED BODY?
O	with a small triangle on the bottom right?
Q*	yeah i think so! bean is a good description lol
Q*	bean body
<i>Rep 2: 5/5 correct</i>	
Q*	bean body!!
<i>Rep 3: 5/5 correct</i>	
Q*	bean body!
<i>Rep 4: 5/5 correct</i>	
Q*	bean body!
<i>Rep 5: 5/5 correct</i>	
Q*	you guys are doin' great. :) bean body for this one
<i>Rep 6: 5/5 correct</i>	
Q*	bean body for this one



6-person thick game	
<i>Rep 1: 2/5 correct</i>	
V*	This shape hasn't got much sticking out. Small triangle pointing right, tilted square on top
Y	Is the triangle on the right at the bottom?
V*	Long tall shape, not far off symmetrical.
V*	Yes triangle bottom right
W	Are their two triangles on it?
X	does it look like a rabbit?
V*	Not the rabbit
X	does it look like someone on their back kicking a square soccer ball
V*	It's like one block like a gemstone
V*	but square on top
V*	Not on back
V*	It's the one with the least sharp corners I think
V*	One long gem shape with small triangle bottom right, tilted square on top
<i>Rep 2: 4/5 correct</i>	
V*	Long tall shape. Little triangle sticking out bottom right. Diamond on top
<i>Rep 3: 3/5 correct</i>	
V*	Long tall hexagon, triangle bottom right, diamond on top
<i>Rep 4: 4/5 correct</i>	
V*	Long hexagon
V*	little triangle bottom right
V*	Diamond top
<i>Rep 5: 4/5 correct</i>	
V*	Long hexagon.
V*	Diamond on top
V*	little triangle bottom right
<i>Rep 6: 4/5 correct</i>	
V*	Long hexagon
V*	Triangle bottom right

Table A2: Additional examples from 6-player thin games for the same image across repetitions. Target image is same as Table A2. Describers are indicated with an asterisk.

6-person thin game	
<i>Rep 1: 2/5 correct</i>	
J*	Looks like a skinnier candle but with a shadow at the bottom
K	😞
H	✅
J*	The top square looks like the flame
J*	It has a triangle leading to that flame
I	❌
<i>Rep 2: 3/5 correct</i>	
K	😞
I*	the simplest image in the whole, no much issues
J	😞
L	😞
I*	1 square as the head
L	😞
J	😞
H	😞
K	😞
I*	candle with more missing chunks
<i>Rep 3: 1/5 correct</i>	
L*	looks like a candle with more chunks missing
J	😞
I	😞
L*	The bottom part faces the right and forms a triangle
L*	The top is a square
I	✅
J	❌
<i>Rep 4: 3/5 correct</i>	
H*	Candle one with shadow
H*	The straight one not the one with bits missing
L	😞
I	😞
L	😞
H*	Lower case j but the bottom swings to the right
I	✅✅✅✅
<i>Rep 5: 0/5 correct</i>	
K*	Candle with chunks missing on the right
I	✅✅✅
<i>Rep 6: 4/5 correct</i>	
G*	candle shape with a chunk of the bottom left missing and added to the bottom right instead
K	😞😞
H	✅
I	✅

6-person thin game	
<i>Rep 1: 2/5 correct</i>	
A*	Square on top
A*	At the bottom of the base part of it sticks out to the right
A*	The base has 8 faces (8 sides of the shape)
C	✅
A*	The square balances on a triangular point of the base
A*	It is a thick base shape, the bottom curves right
D	✅
<i>Rep 2: 5/5 correct</i>	
B*	tombstone and its shadow with a rotated square on top
D	✅
F	✅
C	😞
<i>Rep 3: 5/5 correct</i>	
D*	tombstone casting a shadow (sorry i've pinched that from before)
C	😞
B	😞
E	😞
C	😞
E	😞😞
D*	square on top sitting on top of a triangle which is on top of a square, with 5 side shape at the bottom
<i>Rep 4: 4/5 correct</i>	
C*	looks mummified
C*	person with no limbs
D	✅
B	😞
C*	square atop, chunky with triangle bottom right
<i>Rep 5: 5/5 correct</i>	
F*	Gravestone casting shadow. Square on top
D	✅
<i>Rep 6: 5/5 correct</i>	
E*	tombstone
D	✅

Recruitment and number of games

The number of games in each condition varied due to the realities of online recruitment for real-time games. We had control over how many participants we recruited on

Prolific. Because of the real-time nature of the games, we relied on the fact that most Prolific workers who accepted the task would rapidly open it and go to our experiment site, where they would read the instructions and be assigned into games. Due to technical details about how games filled, we only ran one condition of games at any one time.

In experiments 1 and 2, games would only start if the game was full. Thus, if for example, we recruited 28 people, and 27 people went to the experiment when we were running a 4-person game, three people would be unable to make a complete game, and we also wouldn't have groupmates available for a 28th person who opened the experiment late.

We ran games in large batches to try to minimize the fraction of people who wouldn't be matched into a game. If we got many fewer (complete) games than anticipated, we could recruit another batch of participants, taking the expected attrition rate into account. However, the logistics of needing participants to all open the study at roughly the same time to match into games meant that it was infeasible to recruit fewer than 20 or so participants at a time, so we could not precisely achieve the targeted number of games. Thus, the number of games varied somewhat across conditions.

In addition to not having control over exactly how many games started, some games also ended prematurely when participants disconnected. These partial games were more common in larger groups. We speculate it could be due either to just having more people who might have connection issues or need to leave the game, or more frustration causing people to quit the game (sometimes very early in the game).

Both of these factors led to a low yield rate in terms of completed games per participants recruited, especially in experiment 2 where all the games were 6-player. To try to maximize the data collected per money spent on recruitment, we made adjustments for experiment 3. In experiment 3, we allowed games to them to start and continue with fewer than 6 people. Thus, while many of the "6-player" games did not have 6 players the entire game, we at least have data from the entire course of the game.

In experiment 3, the 6* player games did not all have 6 players, both because games continued as participants dropped out and because if there weren't enough

Table A3: The number of games in each experiment and condition. Complete games finished all 6 blocks; partial games ended early due to disconnections, but contributed at least one complete block of data. 6* indicates that some games started with fewer than 6 players or continued with fewer than 6 players after participants disconnected.

Experiment	Players	Complete	Partial	Total Participants
1: baseline	2	15	4	38
1: baseline	3	18	2	60
1: baseline	4	19	2	84
1: baseline	5	17	3	100
1: baseline	6	12	6	108
2: same describer	6	15	3	108
2: full feedback	6	13	4	102
2: thin	6	10	6	96
3: thin	2	35	3	76
3: thin	6*	44	0	235
3: thick	2	39	3	84
3: thick	6*	38	2	222

players after 5 minutes of waiting, the game would start with whoever was there. All analyses use “intent to treat” and call these 6 player games.

In Figure A3, the number of games goes up in some cases because only complete blocks (where the describer said something every trial) are analysed. If there was initial confusion and a describer missed a trial, that block was excluded.




Figure A3: Number of players during 6 thin and 6 thick games in experiment 3. Blocks that were incomplete were excluded, so if a describer said nothing during a trial, that block was excluded.

Sensitivity analyses

As (post-hoc) sensitivity analyses, we re-ran our main analyses on the subset of games that completed all 6 blocks with a full set of players. This is a check on whether any of the results are driven by a) partial games that were included in early blocks but not later blocks or b) experiment 3 games that started without a full set of players or had players drop out.

Versions of main paper Figures 2 and 4 with just these complete games are shown here in Figures A4 and A5. The full model output for the sensitivity analyses is in Tables A54-A73.

The only qualitative differences in the results between these models and those on the full dataset are in experiment 3. Note that large experiment 3 games were the ones most likely to be excluded, both for incompleteness and for not having the full 6 players throughout.

In the accuracy model, the block:gameSize interaction is not robust, and there is a marginal effect of block:gameSize:thin instead. The main effect of gameSize is also weaker. In the reduction model, the block:gameSize interaction is not robust. In the convergence model, the earlier:gameSize interaction is not robust.




Figure A4: Equivalent of Main text Figure 2, with only full complete games. Behavioral results across all three experiments. (A-C). Matcher accuracy at selecting the target image. (D-F). Number of words produced by the describer each trial. For all, small dots are per game, per block means, and smooth lines are predictions from model fixed effects with 95% credible intervals. Y-axes are truncated, and a few outlier points are not visible.




Figure A5: Equivalent of Main text Figure 4 with only full complete games. Language similarity results measured with pairwise cosine similarity between embeddings of two utterances. (A-C). Convergence of descriptions within games as measured by similarity between an utterance from block 1-5 to the block 6 utterance in the same game for the same image. (D-F). Divergence of descriptions across games as measured by the similarity between two utterances produced for the same image by different groups in the same block. For all, small dots are per game, per block means, and smooth lines are predictions from model fixed effects with 95% credible intervals. Y-axes are truncated, and a few outliers points are not visible.

Matcher contributions

We can think about matcher utterances either at the level of one matcher or at the level of a group. In larger groups, the total number of matchers who could contribute is higher, so matcher utterances might be more common. On a per-matcher basis, having more group-mates could potentially increase the amount a given matcher contributes (more people to potentially have lateral conversations with) or it could decrease it (others are already asking the questions you were going to).

We analyzed matcher utterances both ways – on a game level and on an individual level. We broke it down into two steps: 1) are there matcher contributions on a given trial and 2), for trials with matcher contributions, how long are the contributions.

Matchers' contributions declined over the course of the game. The use of emoji in the thin games is not directly comparable to matcher language use in thick games, since some emoji usage (such as the green checkmark) is most likely equivalent to non-referential matcher language ("got it" etc.) that was excluded. The higher rate of emoji use versus referential language thus could be due to its non-equivalence, a lower level of accuracy in thin games, or matchers having a lower threshold for sending emojis compared to writing out clarifications.

We note a deviation from the pre-registration here in the analysis of the emojis. In the pre-registration we said we would "analyse the distribution of emoji's produced as a function of block and its relation to accuracy and speaker utterance length." We did not do this beyond the visualization shown here.




Figure A6: Matcher contributions at the group level. A-C: Fraction of trials where any matcher said anything that was referential. Dots are per game averages. Smooths are binomial fit lines. D-F: On trials where at least one matcher contributed, the number of words of referential language produced by matchers. Dots are per game averages. Smooths are quadratic fit lines. Y-axis is truncated, and a few outliers points are not visible.




Figure A7: Matcher contributions at the individual level. A-C: Fraction of trials where an individual matcher said anything that was referential. Dots are per game averages. Smooths are binomial fit lines. D-F: On trials where a matcher contributed, the number of words of referential language that matcher produced. Dots are per game averages. Smooths are quadratic fit lines. Y-axis is truncated, and a few outliers points are not visible.




Figure A8: Fraction of trials on which at least one matcher produced the labelled emoji. Fraction of trials when any emoji was produced are shown in black. Dots are per condition, per block estimates with 95% bootstrapped CIs. Smooths are binomial fit lines.

Reaction times

As the number of words produced by the speaker each trial decreases over repetition, so does the time each trial takes. In Figure A9, the time to selection (per-matcher) is shown in panels A-C, and the time per trial (equivalent to how long the slowest matcher took) is shown in panels B-D.

Matchers sped up over time across conditions, although thin large groups may not have speed up as much. Larger groups appear to be a bit slower at the level of individual matchers, which compounds to overall slower games as there are more matchers to wait for.

We also confirm in Figure A10 that there is a strong relationship between the number of words from the describer and the length of a trial.




Figure A9: Behavioral results across all three experiments. (A-C). Individual matcher time to selection. (D-F). Time to end of trial for each group. For all, small dots are per game, per block means, and smooth lines are linear fits. Y-axes are truncated, and a few outlier points are not visible.




Figure A10: Relationship between number of words produced by describer and trial length. Across conditions, a similar pattern holds. X-axes are truncated, and a few outliers points are not visible.

Group level variation in models

One problem with precisely estimating the fixed effects of interest is that there is substantial variation at the game and tangram levels leading to wide confidence intervals on some of the fixed effects.

Models were fit with the maximum mixed effect structure for which the model would run in a reasonable amount of time.

As shown in Figure A11, the standard deviations for the grouping levels are substantial compared to the non-Intercept coefficient estimates. This means that the between-game or between-tangram variation is large compared to the effects of interest. This variation causes wide estimates as there is model uncertainty in how to apportion observed effects between group-level idiosyncrasies and population-level fixed effects.

This variability is diminished in the mega-analytic models which were fit with less extensive mixed effects (to aid in model convergence) and with more data.




Figure A11: Comparison between the population-level coefficients (point estimate and 95% CI) and the group-level standard deviations (point estimates and 95% CIs) for the heirarchical effects.

Additional measures of convergence

The main text included the graph for convergence comparing utterances from blocks 1-5 to the utterance from block 6. Here we show two other measures of semantic shifts for descriptions for the same tangram in the same game: similarity to the first utterance and similarity to the next utterance.

Similarity to the first utterance is not very informative (but we pre-registered it). Similarity to the next utterance is what actually drives the convergence phenomena: pairs of utterances from adjacent blocks become closer together over time.




Figure A12: Additional measures of convergence and divergence. A-C is the similarity between utterances on a given block to the first block utterance for the same image, in the same game. Dots are per-game averages, smooths are quadratic. D-F is the similarity between utterances on a given block to the corresponding utterances in the next block. Dots are per-game averages, smooths are quadratic.

Distinctiveness of tangrams

An additional measure of convergence/divergence patterns is how different tangrams get described in the same game – as nicknames evolve, different tangrams get more different descriptions.




Figure A13: Divergence in descriptions of different tangrams. Cosine similarity between the descriptions of two different tangrams in the same block and group are shown. Dots are per-game averages, smooths are quadratic.

Summaries of model outputs

The following sections contain model outputs. All models were run using BRMS. We report the priors and pre-registration status for each group of models. Tables provide the individual model formulae and the point estimates and 95% credible intervals for the fixed effects.

Note that for all models, block was 0 indexed, so intercepts are what happened during the first block.

Accuracy models

Accuracy models were all run as logistic models with $\text{normal}(0,1)$ priors for both betas and sd. This model was not explicitly included in the experiment 1 and 2 pre-registrations; it was included with more ambitious mixed effects (which did not run in a timely manner) in the experiment 3 pre-registration.

Table A4: Experiment 1 logistic model of matcher accuracy: $\text{correct.num} \sim \text{block} \times \text{numPlayers} + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	2.10	[1.57, 2.65]
block	0.44	[0.31, 0.58]
block:numPlayers	-0.02	[-0.05, 0.01]
numPlayers	-0.07	[-0.2, 0.05]

Table A5: Experiment 2: 6 same describer logistic model of matcher accuracy: $\text{correct.num} \sim \text{block} + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	1.78	[1.4, 2.19]
block	0.45	[0.39, 0.52]

Table A6: Experiment 2: 6 full feedback logistic model of matcher accuracy: $\text{correct.num} \sim \text{block} + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	1.35	[0.59, 2.06]
block	0.47	[0.39, 0.54]

Table A7: Experiment 2: 6 thin logistic model of matcher accuracy: $\text{correct.num} \sim \text{block} + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.88	[0.64, 1.12]
block	0.23	[0.19, 0.28]

Table A8: Experiment 3 logistic model of matcher accuracy: $\text{correct.num} \sim \text{block} \times \text{gameSize} \times \text{channel} + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	1.69	[1.39, 1.99]
block	0.41	[0.32, 0.5]
block:channelthin	-0.07	[-0.18, 0.04]
block:gameSize6	-0.34	[-0.43, -0.25]
block:gameSize6:channelthin	0.07	[-0.05, 0.19]
channelthin	-0.36	[-0.78, 0.05]
gameSize6	-0.64	[-1.05, -0.25]
gameSize6:channelthin	0.31	[-0.22, 0.87]

Reduction models

Primary reduction model

Reduction models were run as linear models with an intercept prior of $\text{normal}(12,20)$, a beta prior of $\text{normal}(0,10)$, an sd prior of $\text{normal}(0,5)$ and a correlation prior of $\text{lkj}(1)$. This model was pre-registered for each experiment and run with the mixed effects structure as pre-specified.

Table A9: Experiment 1: $\text{words} \sim \text{block} \times \text{numPlayers} + (\text{block}|\text{tangram}) + (1|\text{playerId}) + (1|\text{tangram:gameId}) + (\text{block}|\text{gameId})$

Term	Est.	95% CrI
Intercept	16.87	[11.63, 21.89]
block	-3.36	[-4.56, -2.18]
block:numPlayers	-0.09	[-0.37, 0.18]
numPlayers	1.60	[0.62, 2.6]

Table A10: Experiment 2: 6 same describer: words $\sim$ block + (block|tangram) + (1|tangram:gameId) + (block|gameId)

Term	Est.	95% CrI
Intercept	29.65	[24.82, 34.49]
block	-5.31	[-6.35, -4.3]

Table A11: Experiment 2: 6 full feedback: words $\sim$ block + (block|tangram) + (1|tangram:gameId) + (block|gameId)

Term	Est.	95% CrI
Intercept	25.79	[20.97, 30.29]
block	-4.64	[-5.81, -3.53]

Table A12: Experiment 2: 6 thin: words $\sim$ block + (block|tangram) + (1|tangram:gameId) + (block|gameId)

Term	Est.	95% CrI
Intercept	20.3	[17.37, 23.53]
block	-2.1	[-3.37, -1.12]

Table A13: Experiment 3: words $\sim$ block $\times$ channel $\times$ game-Size + (block $\times$ channel $\times$ gameSize|tangram) + (1|tangram:gameId) + (block|gameId)

Term	Est.	95% CrI
Intercept	14.74	[11.68, 17.72]
block	-2.24	[-2.92, -1.57]
block:channelthin	0.29	[-0.56, 1.23]
block:channelthin:gameSize6	0.64	[-0.59, 1.81]
block:gameSize6	-1.22	[-2.06, -0.29]
channelthin	0.80	[-2.85, 4.26]
channelthin:gameSize6	-2.21	[-7.16, 3.08]
gameSize6	7.51	[3.63, 11.3]

Extra reduction model

For experiment 1, we also pre-specified a model about whether the describer's correctness on the prior block (when they were a matcher) had an effect on how many words of description they produced. Priors were the same as for primary reduction model.

Table A14: Experiment 1: words $\sim$ block $\times$ numPlayers + block $\times$ was-INcorrect + (block|tangram) + (1|playerId) + (1|tangram:gameId) + (block|gameId)

Term	Est.	95% CrI
Intercept	12.16	[6.48, 18.07]
block	-2.17	[-3.39, -1]
block:numPlayers	-0.22	[-0.5, 0.06]
block:wasINcorrect	0.24	[-0.24, 0.72]
numPlayers	2.09	[0.88, 3.3]
wasINcorrect	3.07	[1.67, 4.45]

Log reduction model

As an exploratory check, we reran some of the primary reduction models using the log number of words as the DV; otherwise model specifications and priors were the same.

Table A15: Experiment 1 log reduction: $\log\text{words} \sim \text{block} \times \text{numPlayers} + (\text{block}|\text{tangram}) + (1|\text{playerId}) + (1|\text{tangram:gameId}) + (\text{block}|\text{gameId})$

Term	Est.	95% CrI
Intercept	2.75	[2.48, 3.01]
block	-0.38	[-0.47, -0.29]
block:numPlayers	0.01	[-0.01, 0.03]
numPlayers	0.07	[0.02, 0.12]

Table A16: Experiment 3 log reduction: $\log\text{words} \sim \text{block} \times \text{channel} \times \text{gameSize} + (\text{block} \times \text{channel} \times \text{gameSize}|\text{tangram}) + (1|\text{tangram:gameId}) + (\text{block}|\text{gameId})$

Term	Est.	95% CrI
Intercept	2.50	[2.32, 2.69]
block	-0.27	[-0.32, -0.22]
block:channelthin	0.06	[0, 0.14]
block:channelthin:gameSize6	0.03	[-0.06, 0.13]
block:gameSize6	-0.02	[-0.09, 0.04]
channelthin	0.06	[-0.17, 0.29]
channelthin:gameSize6	-0.13	[-0.44, 0.18]
gameSize6	0.45	[0.23, 0.67]

Matcher reduction models

These models were not pre-registered.

For the model of how often any matchers made contributions, the priors were $\text{normal}(0,1)$ for both beta and sd.

For the model of how much was said on trials when matchers talked, the priors were the same as for the primary (describer) reduction model.

We ran these models both at the group level (did any matcher make a contribution, how many words of matcher contributions were there) and at the individual level (for a given matcher, did they make a contribution, and if so, how many words did they contribute).

Table A17: Experiment 1: group-level: $\text{is.words} \sim \text{block} \times \text{numPlayers} + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	-2.67	[-3.54, -1.79]
block	-0.80	[-0.97, -0.62]
block:numPlayers	0.03	[-0.01, 0.07]
numPlayers	0.79	[0.58, 0.98]

Table A18: Experiment 1: individual-level: $\text{is.words} \sim \text{block} \times \text{numPlayers} + (1|\text{playerId})$

Term	Est.	95% CrI
Intercept	-1.65	[-2.24, -1.05]
block	-0.82	[-0.96, -0.68]
block:numPlayers	0.05	[0.02, 0.08]
numPlayers	0.20	[0.07, 0.33]

Table A19: Experiment 1: group-level: $\text{words} \sim \text{block} \times \text{numPlayers} + (\text{block}|\text{gameId})$

Term	Est.	95% CrI
Intercept	4.72	[0.09, 9.44]
block	-0.17	[-1.53, 1.3]
block:numPlayers	-0.41	[-0.72, -0.11]
numPlayers	2.07	[1, 3.12]

Table A20: Experiment 1: individual-level: $\text{words} \sim \text{block} \times \text{numPlayers} + (\text{block}|\text{playerId})$

Term	Est.	95% CrI
Intercept	9.85	[8.3, 11.4]
block	-1.45	[-2.05, -0.86]
block:numPlayers	0.14	[0.02, 0.27]
numPlayers	-0.52	[-0.84, -0.21]

Initial utterance reduction model

These models were not pre-registered. They looked at describer reduction only on words that were produced prior to the first matcher message each trial. These models were only run on experimental conditions where matchers could contribute textual responses.

Reduction models were run as linear models with the same priors as the primary reduction model.

Table A21: Experiment 1: words $\sim$ block $\times$ numPlayers + (block|tangram) + (1|tangram:gameId) + (block|gameId)

Term	Est.	95% CrI
Intercept	18.66	[14.58, 22.71]
block	-3.56	[-4.54, -2.55]
block:numPlayers	0.27	[0.03, 0.5]
numPlayers	-0.33	[-1.14, 0.53]

Table A22: Experiment 2: 6 same describer: words $\sim$ block + (block|tangram) + (1|tangram:gameId) + (block|gameId)

Term	Est.	95% CrI
Intercept	18.06	[14.76, 21.44]
block	-2.49	[-3.19, -1.79]

Table A23: Experiment 2: 6 full feedback: words $\sim$ block + (block|tangram) + (1|tangram:gameId) + (block|gameId)

Term	Est.	95% CrI
Intercept	16.69	[13.41, 20.02]
block	-2.49	[-3.34, -1.62]

Table A24: Experiment 3: words $\sim$ block $\times$ gameSize + (block $\times$ gameSize|tangram) + (1|tangram:gameId) + (block|gameId)

Term	Est.	95% CrI
Intercept	13.88	[11.62, 16.2]
block	-2.08	[-2.66, -1.47]
block:gameSize6	-0.65	[-1.43, 0.14]
gameSize6	5.13	[2, 7.95]

Linguistic content models

We ran a number of models predicting the cosine similarity between pairs of S-BERT embeddings of utterances. For all of these models, we used linear models with the priors $\text{normal}(.5, .2)$ for intercept, $\text{normal}(0, .1)$ for beta, and $\text{normal}(0, .05)$ for sd.

These models were verbally described (but not formally specified) in the pre-registrations for experiment 2 in the full feedback and thin conditions and for experiment 3, for looking at divergence between games, convergence within games (compared to first block, next block, and last block utterances), and divergence between tangrams within games.

Convergence within games: comparison to last round

This is the primary convergence metric presented in the main paper.

Table A25: Experiment 1: $\text{sim} \sim \text{earlier} \times \text{condition} + (1|\text{tangram}) + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.517	[0.458, 0.573]
condition	-0.008	[-0.021, 0.005]
earlier	0.089	[0.076, 0.102]
earlier:condition	-0.008	[-0.011, -0.005]

Table A26: Experiment 2: 6 same describer: $\text{sim} \sim \text{earlier} + (1|\text{tangram}) + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.499	[0.444, 0.556]
earlier	0.086	[0.078, 0.094]

Table A27: Experiment 2: 6 full feedback: $\text{sim} \sim \text{earlier} + (1|\text{tangram}) + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.438	[0.389, 0.487]
earlier	0.062	[0.051, 0.072]

Table A28: Experiment 2: 6 thin: $\text{sim} \sim \text{earlier} + (1|\text{tangram}) + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.498	[0.453, 0.54]
earlier	0.023	[0.013, 0.033]

Table A29: Experiment 3: $\text{sim} \sim \text{earlier} \times \text{channel} \times \text{gameSize} + (1|\text{tangram}) + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.581	[0.542, 0.62]
channelthin	-0.034	[-0.08, 0.011]
channelthin:gameSize6	0.039	[-0.021, 0.097]
earlier	0.080	[0.074, 0.086]
earlier:channelthin	-0.025	[-0.033, -0.017]
earlier:channelthin:gameSize6	-0.035	[-0.047, -0.025]
earlier:gameSize6	0.009	[0.001, 0.017]
gameSize6	-0.069	[-0.113, -0.025]

Divergence across games

This is the divergence metric presented in the paper.

Table A30: Experiment 1: $\text{sim} \sim \text{block} \times \text{condition} + (1|\text{tangram})$

Term	Est.	95% CrI
Intercept	0.468	[0.429, 0.507]
block	-0.035	[-0.038, -0.032]
block:condition	0.001	[0.001, 0.002]
condition	0.002	[0, 0.004]

Table A31: Experiment 2: 6 same describer: $\text{sim} \sim \text{block} + (1|\text{tangram})$

Term	Est.	95% CrI
Intercept	0.484	[0.442, 0.526]
block	-0.041	[-0.043, -0.039]

Table A32: Experiment 2: 6 full feedback: $\text{sim} \sim \text{block} + (1|\text{tangram})$

Term	Est.	95% CrI
Intercept	0.502	[0.46, 0.546]
block	-0.038	[-0.04, -0.035]

Table A33: Experiment 2: 6 thin: $\text{sim} \sim \text{block} + (1|\text{tangram})$

Term	Est.	95% CrI
Intercept	0.434	[0.406, 0.465]
block	-0.004	[-0.006, -0.001]

Table A34: Experiment 3: $\text{sim} \sim \text{block} \times \text{channel} \times \text{gameSize} + (1|\text{togram})$

Term	Est.	95% CrI
Intercept	0.411	[0.368, 0.453]
block	-0.024	[-0.025, -0.023]
block:channelthin	0.004	[0.002, 0.005]
block:channelthin:gameSize6	0.017	[0.015, 0.019]
block:gameSize6	-0.008	[-0.01, -0.007]
channelthin	0.014	[0.01, 0.018]
channelthin:gameSize6	-0.030	[-0.035, -0.024]
gameSize6	0.051	[0.047, 0.055]

Divergence across tangrams

This is an additional metric comparing the similarities between descriptions for different tangrams within a game. It measures how distinct the descriptions for different tangram images are.

Table A35: Experiment 1: $\text{sim} \sim \text{block} \times \text{condition} + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.429	[0.382, 0.473]
block	-0.043	[-0.046, -0.039]
block:condition	0.000	[-0.001, 0.001]
condition	0.003	[-0.008, 0.014]

Table A36: Experiment 2: 6 same describer: $\text{sim} \sim \text{block} + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.416	[0.389, 0.443]
block	-0.046	[-0.048, -0.044]

Table A37: Experiment 2: 6 full feedback: $\text{sim} \sim \text{block} + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.459	[0.422, 0.496]
block	-0.047	[-0.049, -0.044]

Table A38: Experiment 2: 6 thin: $\text{sim} \sim \text{block} + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.432	[0.393, 0.471]
block	-0.025	[-0.028, -0.022]

Table A39: Experiment 3: $\text{sim} \sim \text{block} \times \text{channel} \times \text{gameSize} + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.378	[0.352, 0.404]
block	-0.027	[-0.029, -0.025]
block:channelthin	-0.001	[-0.003, 0.002]
block:channelthin:gameSize6	0.011	[0.008, 0.015]
block:gameSize6	-0.010	[-0.013, -0.008]
channelthin	0.038	[-0.001, 0.082]
channelthin:gameSize6	-0.053	[-0.115, 0]
gameSize6	0.073	[0.035, 0.113]

Convergence to next

We also looked at how similar an utterance was to the next block utterance for the same image in the same group: this can be thought of as the derivative of the to-last comparison. (Although cosine similarities are not actually additive in the same way integrals are).

Table A40: Experiment 1: $\text{sim} \sim \text{earlier} \times \text{condition} + (1|\text{tangram}) + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.591	[0.541, 0.641]
condition	-0.004	[-0.014, 0.006]
earlier	0.063	[0.051, 0.075]
earlier:condition	-0.008	[-0.011, -0.006]

Table A41: Experiment 2: 6 same describer: $\text{sim} \sim \text{earlier} + (1|\text{tangram}) + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.660	[0.619, 0.702]
earlier	0.043	[0.037, 0.05]

Table A42: Experiment 2: 6 full feedback: $\text{sim} \sim \text{earlier} + (1|\text{tangram}) + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.605	[0.569, 0.643]
earlier	0.015	[0.006, 0.024]

Table A43: Experiment 2: 6 thin: $\text{sim} \sim \text{earlier} + (1|\text{tangram}) + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.533	[0.49, 0.578]
earlier	0.010	[0, 0.019]

Table A44: Experiment 3: $\text{sim} \sim \text{earlier} \times \text{channel} \times \text{gameSize} + (1|\text{tangram}) + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.714	[0.682, 0.746]
channelthin	-0.124	[-0.159, -0.088]
channelthin:gameSize6	0.000	[-0.051, 0.049]
earlier	0.046	[0.041, 0.052]
earlier:channelthin	-0.010	[-0.018, -0.002]
earlier:channelthin:gameSize6	-0.018	[-0.029, -0.007]
earlier:gameSize6	-0.003	[-0.011, 0.004]
gameSize6	-0.034	[-0.069, 0.003]

Divergence from first

We also looked at how similar an utterance was to the first block utterance for the same image. This is not very informative because first round utterances tend to be pretty noisy with lots of hedges and filler words.

Table A45: Experiment 1: $\text{sim} \sim \text{later} \times \text{condition} + (1|\text{tangram}) + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.647	[0.591, 0.705]
condition	-0.010	[-0.022, 0.003]
later	-0.030	[-0.041, -0.019]
later:condition	0.001	[-0.002, 0.004]

Table A46: Experiment 2: 6 same describer: $\text{sim} \sim \text{later} + (1|\text{tangram}) + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.680	[0.628, 0.728]
later	-0.042	[-0.049, -0.035]

Table A47: Experiment 2: 6 full feedback: $\text{sim} \sim \text{later} + (1|\text{tangram}) + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.644	[0.584, 0.706]
later	-0.044	[-0.052, -0.037]

Table A48: Experiment 2: 6 thin: $\text{sim} \sim \text{later} + (1|\text{tangram}) + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.537	[0.49, 0.584]
later	-0.014	[-0.023, -0.004]

Table A49: Experiment 3: $\text{sim} \sim \text{later} \times \text{channel} \times \text{gameSize} + (1|\text{tango}) + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.721	[0.681, 0.76]
channelthin	-0.076	[-0.123, -0.026]
channelthin:gameSize6	-0.062	[-0.127, 0.001]
gameSize6	-0.017	[-0.062, 0.03]
later	-0.034	[-0.039, -0.028]
later:channelthin	0.011	[0.003, 0.019]
later:channelthin:gameSize6	0.021	[0.01, 0.032]
later:gameSize6	-0.011	[-0.019, -0.004]

Exploratory Mega-analytic models

As an exploratory measure, we combined data across all experiments and re-ran the core set of models (same structure and priors as for the individual models). We binarized game thickness: all games that were not thin were counted as thick. The intercept condition is a 2-player game, in the thick condition, for the first block. Thinner means the game was in the thin condition; larger is per additional player, and block is per later block. Group size was coded based on assigned condition, not based on actual number of players at the time.

Table A50: Mega-analytic on accuracy: $\text{correct.num} \sim \text{block} \times \text{thinner} \times \text{larger} + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	1.83	[1.6, 2.07]
block	0.46	[0.4, 0.52]
block:larger	-0.07	[-0.09, -0.05]
block:thinner	-0.12	[-0.21, -0.02]
block:thinner:larger	0.01	[-0.02, 0.04]
larger	-0.07	[-0.15, 0]
thinner	-0.50	[-0.89, -0.08]
thinner:larger	-0.02	[-0.14, 0.11]

Table A51: Mega-analytic of reduction: $\text{words} \sim \text{block} \times \text{thinner} \times \text{larger} + (\text{block}|\text{gameId})$

Term	Est.	95% CrI
Intercept	17.38	[15.59, 19.21]
block	-2.80	[-3.24, -2.36]
block:larger	-0.36	[-0.51, -0.2]
block:thinner	0.79	[0.04, 1.52]
block:thinner:larger	0.26	[0, 0.51]
larger	2.12	[1.5, 2.75]
thinner	-1.59	[-4.61, 1.55]
thinner:larger	-0.90	[-1.92, 0.11]

Table A52: Mega-analytic on divergence between groups: $\text{sim} \sim \text{block} \times \text{thinner} \times \text{larger}$

Term	Est.	95% CrI
Intercept	0.428	[0.426, 0.431]
block	-0.026	[-0.026, -0.025]
block:larger	-0.002	[-0.002, -0.002]
block:thinner	0.005	[0.004, 0.007]
block:thinner:larger	0.004	[0.004, 0.005]
larger	0.012	[0.011, 0.013]
thinner	-0.003	[-0.007, 0.001]
thinner:larger	-0.007	[-0.008, -0.006]

Table A53: Mega-analytic on convergence to last: $\text{sim} \sim \text{earlier} \times \text{larger} \times \text{thinner}$

Term	Est.	95% CrI
Intercept	0.546	[0.534, 0.558]
earlier	0.072	[0.067, 0.077]
earlier:larger	0.000	[-0.002, 0.002]
earlier:larger:thinner	-0.007	[-0.01, -0.004]
earlier:thinner	-0.016	[-0.025, -0.008]
larger	-0.016	[-0.021, -0.012]
larger:thinner	0.009	[0.002, 0.016]
thinner	0.001	[-0.021, 0.021]

Sensitivity analysis models

These are the same specification as the primary models, just on the subset of the data from games that were full and complete.

Accuracy

Table A54: Experiment 1 logistic model of matcher accuracy: $\text{correct.num} \sim \text{block} \times \text{numPlayers} + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	2.15	[1.5, 2.78]
block	0.43	[0.29, 0.57]
block:numPlayers	-0.02	[-0.05, 0.01]
numPlayers	-0.08	[-0.22, 0.07]

Table A55: Experiment 2: 6 same describer logistic model of matcher accuracy: $\text{correct.num} \sim \text{block} + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	1.81	[1.38, 2.27]
block	0.45	[0.37, 0.52]

Table A56: Experiment 2: 6 full feedback logistic model of matcher accuracy: $\text{correct.num} \sim \text{block} + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	1.72	[1.41, 2.05]
block	0.46	[0.39, 0.54]

Table A57: Experiment 2: 6 thin logistic model of matcher accuracy: $\text{correct.num} \sim \text{block} + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.86	[0.43, 1.29]
block	0.26	[0.21, 0.31]

Table A58: Experiment 3 logistic model of matcher accuracy: $\text{correct.num} \sim \text{block} \times \text{gameSize} \times \text{channel} + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	1.67	[1.4, 1.96]
block	0.41	[0.32, 0.5]
block:channelthin	-0.06	[-0.18, 0.06]
block:gameSize6	-0.02	[-0.12, 0.09]
block:gameSize6:channelthin	-0.14	[-0.27, 0]
channelthin	-0.36	[-0.77, 0.05]
gameSize6	-0.43	[-0.9, 0.03]
gameSize6:channelthin	0.42	[-0.21, 1.04]

Reduction

Table A59: Experiment 1: $\text{words} \sim \text{block} \times \text{numPlayers} + (\text{block}|\text{tangram}) + (1|\text{playerId}) + (1|\text{tangram:gameId}) + (\text{block}|\text{gameId})$

Term	Est.	95% CrI
Intercept	17.22	[11.42, 22.74]
block	-3.35	[-4.64, -2.05]
block:numPlayers	-0.09	[-0.39, 0.21]
numPlayers	1.55	[0.43, 2.69]

Table A60: Experiment 2: 6 same describer: $\text{words} \sim \text{block} + (\text{block}|\text{tangram}) + (1|\text{tangram:gameId}) + (\text{block}|\text{gameId})$

Term	Est.	95% CrI
Intercept	29.45	[23.78, 35.14]
block	-5.23	[-6.36, -4.05]

Table A61: Experiment 2: 6 full feedback: words $\sim$ block + (block|tangram) + (1|tangram:gameId) + (block|gameId)

Term	Est.	95% CrI
Intercept	25.17	[20.16, 30.14]
block	-4.40	[-5.5, -3.28]

Table A62: Experiment 2: 6 thin: words $\sim$ block + (block|tangram) + (1|tangram:gameId) + (block|gameId)

Term	Est.	95% CrI
Intercept	20.35	[16.21, 24.45]
block	-1.70	[-3.03, -0.38]

Table A63: Experiment 3: words $\sim$ block $\times$ channel $\times$ game-Size + (block $\times$ channel $\times$ gameSize|tangram) + (1|tangram:gameId) + (block|gameId)

Term	Est.	95% CrI
Intercept	14.93	[12.54, 17.43]
block	-2.21	[-2.77, -1.67]
block:channelthin	0.26	[-0.48, 1.05]
block:channelthin:gameSize6	0.07	[-1.28, 1.45]
block:gameSize6	-0.35	[-1.32, 0.64]
channelthin	0.52	[-2.53, 3.52]
channelthin:gameSize6	0.58	[-4.8, 6.19]
gameSize6	3.99	[0.09, 8.14]

Convergence within games

Table A64: Experiment 1: $\text{sim} \sim \text{earlier} \times \text{condition} + (1|\text{tangram}) + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.517	[0.458, 0.575]
condition	-0.008	[-0.02, 0.005]
earlier	0.090	[0.077, 0.103]
earlier:condition	-0.008	[-0.011, -0.005]

Table A65: Experiment 2: 6 same describer: $\text{sim} \sim \text{earlier} + (1|\text{tangram}) + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.501	[0.443, 0.556]
earlier	0.087	[0.079, 0.095]

Table A66: Experiment 2: 6 full feedback: $\text{sim} \sim \text{earlier} + (1|\text{tangram}) + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.438	[0.388, 0.486]
earlier	0.062	[0.052, 0.072]

Table A67: Experiment 2: 6 thin: $\text{sim} \sim \text{earlier} + (1|\text{tangram}) + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.499	[0.448, 0.55]
earlier	0.022	[0.011, 0.032]

Table A68: Experiment 3: $\text{sim} \sim \text{earlier} \times \text{channel} \times \text{gameSize} + (1|\text{tangram}) + (1|\text{gameId})$

Term	Est.	95% CrI
Intercept	0.583	[0.549, 0.618]
channelthin	-0.037	[-0.084, 0.008]
channelthin:gameSize6	0.010	[-0.069, 0.09]
earlier	0.080	[0.075, 0.086]
earlier:channelthin	-0.023	[-0.032, -0.015]
earlier:channelthin:gameSize6	-0.031	[-0.046, -0.015]
earlier:gameSize6	0.006	[-0.005, 0.017]
gameSize6	-0.058	[-0.116, -0.003]

Divergence across games

Table A69: Experiment 1: $\text{sim} \sim \text{block} \times \text{condition} + (1|\text{tangram})$

Term	Est.	95% CrI
Intercept	0.468	[0.426, 0.51]
block	-0.034	[-0.037, -0.03]
block:condition	0.001	[0, 0.002]
condition	0.002	[0, 0.005]

Table A70: Experiment 2: 6 same describer: $\text{sim} \sim \text{block} + (1|\text{tangram})$

Term	Est.	95% CrI
Intercept	0.469	[0.429, 0.516]
block	-0.038	[-0.041, -0.036]

Table A71: Experiment 2: 6 full feedback: $\text{sim} \sim \text{block} + (1|\text{tangram})$

Term	Est.	95% CrI
Intercept	0.509	[0.461, 0.554]
block	-0.040	[-0.043, -0.037]

Table A72: Experiment 2: 6 thin: $\text{sim} \sim \text{block} + (1|\text{tangram})$

Term	Est.	95% CrI
Intercept	0.441	[0.401, 0.478]
block	-0.007	[-0.011, -0.004]

Table A73: Experiment 3: $\text{sim} \sim \text{block} \times \text{channel} \times \text{gameSize} + (1|\text{tangram})$

Term	Est.	95% CrI
Intercept	0.416	[0.368, 0.458]
block	-0.025	[-0.026, -0.024]
block:channelthin	0.005	[0.003, 0.006]
block:channelthin:gameSize6	0.023	[0.019, 0.027]
block:gameSize6	-0.007	[-0.01, -0.004]
channelthin	0.011	[0.006, 0.015]
channelthin:gameSize6	-0.064	[-0.075, -0.052]
gameSize6	0.057	[0.048, 0.066]

References

- Abbot-Smith, K., Nurmsoo, E., Croll, R., Ferguson, H., & Forrester, M. (2015). How children aged 2;6 tailor verbal expressions to interlocutor informational needs. *Journal of Child Language*, 43(6), 1277–1291. <https://doi.org/10.1017/S0305000915000616>
- Ahern, T. C. (1994). The Effect of Interface on the Structure of Interaction in Computer-Mediated Small-Group Discussion. *Journal of Educational Computing Research*, 11(3), 235–250. <https://doi.org/10.2190/5P8J-10M6-GP9K-6REP>
- Almaatouq, A., Becker, J., Houghton, J. P., Paton, N., Watts, D. J., & Whiting, M. E. (2021). Empirica: A virtual lab for high-throughput macro-level experiments. *Behavior Research Methods*, 53(5), 2158–2171. <https://doi.org/10.3758/s13428-020-01535-9>
- Almaatouq, A., Griffiths, T. L., Suchow, J. W., Whiting, M. E., Evans, J., & Watts, D. J. (2022). Beyond playing 20 questions with nature: Integrative experiment design in the social and behavioral sciences. *Behavioral and Brain Sciences*, 47, e33. <https://doi.org/10.1017/S0140525X22002874>
- Bangerter, A., Mayor, E., & Knutsen, D. (2020). Lexical entrainment without conceptual pacts? Revisiting the matching task. *Journal of Memory and Language*, 114, 104129. <https://doi.org/10.1016/j.jml.2020.104129>
- Bates, E. (1974). Acquisition of pragmatic competence. *Journal of Child Language*, 1(2), 277–281. <https://doi.org/10.1017/S0305000900000702>
- Bergen, L., Levy, R., & Goodman, N. (2016). Pragmatic reasoning through semantic inference. *Semantics and Pragmatics*, 9(20), 1–91. <https://doi.org/10.3765/sp.9.20>
- Bohn, M., & Frank, M. C. (2019). The Pervasive Role of Pragmatics in Early Language. *Annual Review of Developmental Psychology*, 1(1), 223–249. <https://doi.org/10.1146/annurev-devpsy-010119-00001>

- [//doi.org/10.1146/annurev-devpsych-121318-085037](https://doi.org/10.1146/annurev-devpsych-121318-085037)
- Bohn, M., Kachel, G., & Tomasello, M. (2019). Young children spontaneously recreate core properties of language in a new modality. *Proceedings of the National Academy of Sciences*, 116(51), 26072–26077. <https://doi.org/10.1073/pnas.1904871116>
- Bohn, M., Schmidt, L. S., Schulze, C., Frank, M. C., & Tessler, M. H. (2022). Modeling Individual Differences in Children’s Information Integration During Pragmatic Word Learning. *Open Mind*, 6, 311–326. https://doi.org/10.1162/opmi_a_00069
- Bovet, V., Knutsen, D., & Fossard, M. (2023). Direct and indirect linguistic measures of common ground in dialogue studies involving a matching task: A systematic review. *Psychonomic Bulletin & Review*, 31(1), 122–136. <https://doi.org/10.3758/s13423-023-02359-2>
- Boyce, V., & Frank, M. C. (2023). Communicative reduction in referring expressions within a multi-player negotiation game. *Proceedings of the Annual Meeting of the Cognitive Science Society*.
- Boyce, V., Hawkins, R. D., Goodman, N. D., & Frank, M. C. (2024). Interaction structure constrains the emergence of conventions in group communication. *Proceedings of the National Academy of Sciences*, 121(28), e2403888121. <https://doi.org/10.1073/pnas.2403888121>
- Branigan, H. (2006). Perspectives on Multi-party Dialogue. *Research on Language and Computation*, 4(2-3), 153–177. <https://doi.org/10.1007/s11168-006-9002-2>
- Branigan, H. P., Bell, J., & McLean, J. F. (2016). Do You Know What I Know? The Impact of Participant Role in Children’s Referential Communication. *Frontiers in Psychology*, 7. <https://doi.org/10.3389/fpsyg.2016.00213>
- Branigan, H. P., Catchpole, C. M., & Pickering, M. J. (2011). What makes dialogues easy to understand? *Language and Cognitive Processes*, 26(10), 1667–1686. <https://doi.org/10.1080/01690965.2010.524765>
- Brennan, S. E., & Clark, H. H. (1996). Conceptual pacts and lexical choice in conversation. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 22(6), 1482–1493. <https://doi.org/10.1037/0278-7393.22.6.1482>

- Brown-Schmidt, S. (2012). Beyond common and privileged: Gradient representations of common ground in real-time language use. *Language and Cognitive Processes*, 27(1), 62–89. <https://doi.org/10.1080/01690965.2010.543363>
- Bürki, A. (2018). Variation in the speech signal as a window into the cognitive architecture of language production. *Psychonomic Bulletin & Review*, 25(6), 1973–2004. <https://doi.org/10.3758/s13423-017-1423-4>
- Bürkner, P.-C. (2018). Advanced Bayesian Multilevel Modeling with the R Package brms. *The R Journal*, 10(1), 395. <https://doi.org/10.32614/RJ-2018-017>
- Caplow, T. (1957). Organizational Size. *Administrative Science Quarterly*, 1(4), 484. <https://doi.org/10.2307/2390870>
- Carletta, J., Garrod, S., & Fraser-Krauss, H. (1998). Placement of Authority and Communication Patterns in Workplace Groups. *Small Group Research*, 29(5), 531–559. <https://doi.org/10.1177/1046496498295001>
- Carruthers, P. (2013). Mindreading in Infancy. *Mind & Language*, 28(2), 141–172. <https://doi.org/10.1111/mila.12014>
- Cazden, C. B. (1988). *Classroom discourse: The language of teaching and learning*. ERIC.
- Champagne-Lavau, M., Fossard, M., Martel, G., Chapdelaine, C., Blouin, G., Rodriguez, J.-P., & Stip, E. (2009). Do patients with schizophrenia attribute mental states in a referential communication task? *Cognitive Neuropsychiatry*, 14(3), 217–239. <https://doi.org/10.1080/13546800903004114>
- Chantraine, Y., Joannette, Y., & Ska, B. (1998). Conversational abilities in patients with right hemisphere damage. *Journal of Neurolinguistics*, 11(1-2), 21–32. [https://doi.org/10.1016/S0911-6044\(98\)00003-7](https://doi.org/10.1016/S0911-6044(98)00003-7)
- Chen, T., & Guestrin, C. (2016). Xgboost: A scalable tree boosting system. *Proceedings of the 22nd Acm Sigkdd International Conference on Knowledge Discovery and Data Mining*, 785–794. <https://doi.org/10.1145/2939672.2939785>
- Clark, H. H., & Krych, M. A. (2004). Speaking while monitoring addressees for understanding. *Journal of Memory and Language*, 50(1), 62–81. <https://doi.org/10.1016/j.jml.2003.08.004>
- Clark, H. H., & Schaefer, E. F. (1987). Concealing one's meaning from overhearers. *Journal of Memory and Language*, 26(2), 209–225. <https://doi.org/10.1016/>

0749-596X(87)90124-0

- Clark, H. H., & Wilkes-Gibbs, D. (1986). Referring as a collaborative process. *Cognition*, 22(1), 1–39. [https://doi.org/10.1016/0010-0277\(86\)90010-7](https://doi.org/10.1016/0010-0277(86)90010-7)
- Cohn-Gordon, R., Goodman, N. D., & Potts, C. (2018). An Incremental Iterated Response Model of Pragmatics. *arXiv:1810.00367 [Cs]*. <https://doi.org/10.48550/arxiv.1810.00367>
- Cohn-Gordon, R., Levy, R., & Bergen, L. (2019). *The pragmatics of multiparty communication*.
- Dahan, D. (2024). When hedging helps, rather than impedes, communication: Collaboration in the referential communication task. *Discourse Processes*, 62(2), 89–111. <https://doi.org/10.1080/0163853X.2024.2437738>
- Dahan, D. (2025). Coordinating reference in conversation: The choice between linguistic conventions and linguistic precedents. *Journal of Memory and Language*, 143, 104639. <https://doi.org/10.1016/j.jml.2025.104639>
- de Leeuw, J. R., Gilbert, R. A., & Luchterhandt, B. (2023). jsPsych: Enabling an Open-Source Collaborative Ecosystem of Behavioral Experiments. *Journal of Open Source Software*, 8(85), 5351. <https://doi.org/10.21105/joss.05351>
- de Ruiter, J. P., Bangerter, A., & Dings, P. (2012). The Interplay Between Gesture and Speech in the Production of Referring Expressions: Investigating the Tradeoff Hypothesis. *Topics in Cognitive Science*, 4(2), 232–248. <https://doi.org/10.1111/j.1756-8765.2012.01183.x>
- Degen, J., Hawkins, R. D., Graf, C., Kreiss, E., & Goodman, N. D. (2020). When redundancy is useful: A Bayesian approach to “overinformative” referring expressions. *Psychological Review*, 127(4), 591–621. <https://doi.org/10.1037/rev0000186>
- Degen, J., Kursat, L., & Leigh, D. (2021). Seeing is believing: Testing an explicit linking assumption for visual world eye-tracking in psycholinguistics. *Proceedings of the Annual Meeting of the Cognitive Science Society*, 43.
- Dewhirst, H. D. (1971). Influence of Perceived Information-Sharing Norms on Communication Channel Utilization. *Academy of Management Journal*, 14(3), 305–315. <https://doi.org/10.2307/255075>
- Doyle, G., & Frank, M. C. (2016). Investigating the Sources of Linguistic Alignment

- in Conversation. *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 526–536. <https://doi.org/10.18653/v1/P16-1050>
- Doyle, G., Yurovsky, D., & Frank, M. C. (2016). A Robust Framework for Estimating Linguistic Alignment in Twitter Conversations. *Proceedings of the 25th International Conference on World Wide Web*, 637–648. <https://doi.org/10.1145/2872427.2883091>
- Duff, M. C., Gupta, R., Hengst, J. A., Tranel, D., & Cohen, N. J. (2011). The Use of Definite References Signals Declarative Memory. *Psychological Science*, 22(5), 666–673. <https://doi.org/10.1177/0956797611404897>
- Duff, M. C., Hengst, J., Tranel, D., & Cohen, N. J. (2005). Development of shared information in communication despite hippocampal amnesia. *Nature Neuroscience*, 9(1), 140–146. <https://doi.org/10.1038/nn1601>
- Duff, M. C., Warren, D. E., Gupta, R., Vidal, J. P. B., Tranel, D., & Cohen, N. J. (2011). Teasing apart tangrams: Testing hippocampal pattern separation with a collaborative referencing paradigm. *Hippocampus*, 22(5), 1087–1091. <https://doi.org/10.1002/hipo.20967>
- Eliav, R., Ji, A., Artzi, Y., & Hawkins, R. D. (2023). *Semantic uncertainty guides the extension of conventions to new referents* (No. arXiv:2305.06539). arXiv. <https://doi.org/10.48550/arXiv.2305.06539>
- Erk, K. (2022). The Probabilistic Turn in Semantics and Pragmatics. *Annual Review of Linguistics*, 8(1), 101–121. <https://doi.org/10.1146/annurev-linguistics-031120-015515>
- Fay, N., Garrod, S., & Carletta, J. (2000). Group Discussion as Interactive Dialogue or as Serial Monologue: The Influence of Group Size. *Psychological Science*, 11(6), 481–486. <https://doi.org/10.1111/1467-9280.00292>
- Fay, N., Garrod, S., Roberts, L., & Swoboda, N. (2010). The Interactive Evolution of Human Communication Systems. *Cognitive Science*, 34(3), 351–386. <https://doi.org/10.1111/j.1551-6709.2009.01090.x>
- Ferrara, C., Lu, J. C., & Goldin-Meadow, S. (2025). Playing With Language in the Manual Modality: Which Motions Do Signers Gradiently Modify? *Cognitive Science*, 49(4), e70051. <https://doi.org/10.1111/cogs.70051>

- Ferreira, F. (2004). Production-comprehension asymmetries. *Behavioral and Brain Sciences*, 27(02), 196–196. <https://doi.org/10.1017/S0140525X04280050>
- Ferreira, V. S. (2019). A Mechanistic Framework for Explaining Audience Design in Language Production. *Annual Review of Psychology*, 70(1), 29–51. <https://doi.org/10.1146/annurev-psych-122216-011653>
- Fox Tree, J. E. (1999). Listening in on monologues and dialogues. *Discourse Processes*, 27(1), 35–53. <https://doi.org/10.1080/01638539909545049>
- Fox Tree, J. E., & Clark, N. B. (2013). Communicative Effectiveness of Written Versus Spoken Feedback. *Discourse Processes*, 50(5), 339–359. <https://doi.org/10.1080/0163853X.2013.797241>
- Frank, M. C., & Goodman, N. D. (2012). Predicting Pragmatic Reasoning in Language Games. *Science*, 336(6084), 998–998. <https://doi.org/10.1126/science.1218633>
- Franke, M., & Degen, J. (2016). Reasoning in Reference Games: Individual- vs. Population-Level Probabilistic Modeling. *PLOS ONE*, 11(5), e0154854. <https://doi.org/10.1371/journal.pone.0154854>
- Futrell, R. (2023). Information-theoretic principles in incremental language production. *Proceedings of the National Academy of Sciences*, 120(39), e2220593120. <https://doi.org/10.1073/pnas.2220593120>
- Futrell, R., & Hahn, M. (2022). Information Theory as a Bridge Between Language Function and Language Form. *Frontiers in Communication*, 7. <https://doi.org/10.3389/fcomm.2022.657725>
- Futrell, R., & Levy, R. (2017). Noisy-context surprisal as a human sentence processing cost model. In M. Lapata, P. Blunsom, & A. Koller (Eds.), *Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics: Volume 1, Long Papers* (pp. 688–698). Association for Computational Linguistics.
- Gandolfi, G., Pickering, M. J., & Garrod, S. (2022). Mechanisms of alignment: Shared control, social cognition and metacognition. *Philosophical Transactions of the Royal Society B: Biological Sciences*, 378(1870), 20210362. <https://doi.org/10.1098/rstb.2021.0362>
- Garrison, A. C. S., Yoon, S. O., Brown-Schmidt, S., Ariss, T., & Fairbairn, C. E.

- (2024). Alcohol and common ground: The effects of intoxication on linguistic markers of shared understanding during social exchange. *Psychology of Addictive Behaviors*, 38(1), 79–91. <https://doi.org/10.1037/adb0000922>
- Garrod, S., Fay, N., Lee, J., Oberlander, J., & MacLeod, T. (2007). Foundations of Representation: Where Might Graphical Symbol Systems Come From? *Cognitive Science*, 31(6), 961–987. <https://doi.org/10.1080/03640210701703659>
- Gershman, S. J. (2018). How to never be wrong. *Psychonomic Bulletin & Review*, 26(1), 13–28. <https://doi.org/10.3758/s13423-018-1488-8>
- Ghaleb, E., Rasenberg, M., Pouw, W., Holler, J., & Fernandez, R. (2024). *Analysing Cross-Speaker Convergence in Face-to-Face Dialogue through the Lens of Automatically Detected Shared Linguistic Constructions*. <https://doi.org/10.48550/arXiv.2405.08546>
- Gibson, E., Futrell, R., Piantadosi, S. P., Dautriche, I., Mahowald, K., Bergen, L., & Levy, R. (2019). How Efficiency Shapes Human Language. *Trends in Cognitive Sciences*, 23(5), 389–407. <https://doi.org/10.1016/j.tics.2019.02.003>
- Ginzburg, J., & Fernandez, R. (2005). Action at a distance: The difference between dialogue and multilogue. *Proceedings of DIALOR*, 5, 9.
- Glucksberg, S., & Krauss, R. M. (1967). What do people say after they have learned how to talk? Studies of the development of referential communication. *Merrill-Palmer Quarterly of Behavior and Development*, 13(4), 309–316. <https://www.jstor.org/stable/23082551>
- Glucksberg, S., Krauss, R., & Weisburg, R. (1966). *Referential Communication in Nursery School Children: Method and Some Preliminary Findings*.
- Goodman, N. D., & Frank, M. C. (2016). Pragmatic Language Interpretation as Probabilistic Inference. *Trends in Cognitive Sciences*, 20(11), 818–829. <https://doi.org/10.1016/j.tics.2016.08.005>
- Goodman, N. D., & Stuhlmüller, A. (2013). Knowledge and Implicature: Modeling Language Understanding as Social Cognition. *Topics in Cognitive Science*, 5(1), 173–184. <https://doi.org/10.1111/tops.12007>
- Graham, S. A., Sedivy, J., & Khu, M. (2013). That’s not what you said earlier: Preschoolers expect partners to be referentially consistent. *Journal of Child Language*, 41(1), 34–50. <https://doi.org/10.1017/S0305000912000530>

- Grand, G., Blank, I. A., Pereira, F., & Fedorenko, E. (2022). Semantic projection recovers rich human knowledge of multiple object features from word embeddings. *Nature Human Behaviour*, 6(7), 975–987. <https://doi.org/10.1038/s41562-022-01316-8>
- Grattafiori, A., Dubey, A., Jauhri, A., Pandey, A., Kadian, A., Al-Dahle, A., Letman, A., Mathur, A., Schelten, A., Vaughan, A., Yang, A., Fan, A., Goyal, A., Hartshorn, A., Yang, A., Mitra, A., Sravankumar, A., Korenev, A., Hinsvark, A., ... Ma, Z. (2024). *The Llama 3 Herd of Models* (No. arXiv:2407.21783). arXiv. <https://doi.org/10.48550/arXiv.2407.21783>
- Grice, H. P. (1989). Logic and conversation. In *Studies in the way of Words*. Academic Pr.
- Grigoroglou, M., & Papafragou, A. (2019). Interactive contexts increase informativeness in children’s referential communication. *Developmental Psychology*, 55(5), 951–966. <https://doi.org/10.1037/dev0000693>
- Guilbeault, D., Baronchelli, A., & Centola, D. (2021). Experimental evidence for scale-induced category convergence across populations. *Nature Communications*, 12(1), 327. <https://doi.org/10.1038/s41467-020-20037-y>
- Gul, M. O., & Artzi, Y. (2024). *CoGen: Learning from Feedback with Coupled Comprehension and Generation* (No. arXiv:2408.15992). arXiv. <https://doi.org/10.48550/arXiv.2408.15992>
- Gupta, R., Tranel, D., & Duff, M. C. (2012). Ventromedial prefrontal cortex damage does not impair the development and use of common ground in social interaction: Implications for cognitive theory of mind. *Neuropsychologia*, 50(1), 145–152. <https://doi.org/10.1016/j.neuropsychologia.2011.11.012>
- Gwilliams, L., & Davis, M. H. (2022). Extracting Language Content from Speech Sounds: The Information Theoretic Approach. In L. L. Holt, J. E. Peelle, A. B. Coffin, A. N. Popper, & R. R. Fay (Eds.), *Speech Perception* (pp. 113–139). Springer International Publishing. https://doi.org/10.1007/978-3-030-81542-4_5
- Gwilliams, L., Marantz, A., Poeppel, D., & King, J.-R. (2023). Top-down information shapes lexical processing when listening to continuous speech. *Language, Cognition and Neuroscience*, 39(8), 1045–1058. <https://doi.org/10.1080/23273798>

- 2023.2171072
- Haber, J., Baumgärtner, T., Takmaz, E., Gelderloos, L., Bruni, E., & Fernández, R. (2019). The PhotoBook Dataset: Building Common Ground through Visually-Grounded Dialogue. *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, 1895–1910. <https://doi.org/10.18653/v1/P19-1184>
- Hale, J. (2001). A probabilistic earley parser as a psycholinguistic model. *Second Meeting of the North American Chapter of the Association for Computational Linguistics on Language Technologies 2001 - NAACL '01*, 1–8. <https://doi.org/10.3115/1073336.1073357>
- Hale, J. (2016). Information-theoretical Complexity Metrics. *Language and Linguistics Compass*, 10(9), 397–412. <https://doi.org/10.1111/lnc3.12196>
- Hanna, J. E., Tanenhaus, M. K., & Trueswell, J. C. (2003). The effects of common ground and perspective on domains of referential interpretation. *Journal of Memory and Language*, 49(1), 43–61. [https://doi.org/10.1016/S0749-596X\(03\)00022-6](https://doi.org/10.1016/S0749-596X(03)00022-6)
- Hawkins, J. (1995). *A Performance Theory of Order and Constituency*. Cambridge University Press.
- Hawkins, R. D., Frank, M. C., & Goodman, N. D. (2020). Characterizing the Dynamics of Learning in Repeated Reference Games. *Cognitive Science*, 44(6). <https://doi.org/10.1111/cogs.12845>
- Hawkins, R. D., Franke, M., Frank, M. C., Goldberg, A. E., Smith, K., Griffiths, T. L., & Goodman, N. D. (2023). From partners to populations: A hierarchical Bayesian account of coordination and convention. *Psychological Review*, 130(4), 977–1016. <https://doi.org/10.1037/rev0000348>
- Hawkins, R. D., Goodman, N. D., Goldberg, A. E., & Griffiths, T. L. (2020). Generalizing meanings from partners to populations: Hierarchical inference supports convention formation on networks. *arXiv:2002.01510 [Cs]*. <https://doi.org/10.48550/arxiv.2002.01510>
- Hawkins, R. D., Liu, I., Goldberg, A. E., & Griffiths, T. G. (2021). Respect the code: Speakers expect novel conventions to generalize within but not across social group boundaries. *Proceedings of the Annual Meeting of the Cognitive Science Society*,

43.

- Hawkins, R. D., Sano, M., Goodman, N. D., & Fan, J. E. (2023). Visual resemblance and interaction history jointly constrain pictorial meaning. *Nature Communications*, 14(1). <https://doi.org/10.1038/s41467-023-37737-w>
- Hengst, J. A. (2003). Collaborative Referencing Between Individuals With Aphasia and Routine Communication Partners. *Journal of Speech, Language, and Hearing Research*, 46(4), 831–848. [https://doi.org/10.1044/1092-4388\(2003/065\)](https://doi.org/10.1044/1092-4388(2003/065))
- Hiltz, S. R., Johnson, K., & Turoff, M. (1986). Experiments in Group Decision Making Communication Process and Outcome in Face-to-Face Versus Computerized Conferences. *Human Communication Research*, 13(2), 225–252. <https://doi.org/10.1111/j.1468-2958.1986.tb00104.x>
- Horton, W. S., & Gerrig, R. J. (2002). Speakers' experiences and audience design: Knowing when and knowing how to adjust utterances to addressees. *Journal of Memory and Language*, 47(4), 589–606. [https://doi.org/10.1016/S0749-596X\(02\)00019-0](https://doi.org/10.1016/S0749-596X(02)00019-0)
- Horton, W. S., & Gerrig, R. J. (2005). The impact of memory demands on audience design during language production. *Cognition*, 96(2), 127–142. <https://doi.org/10.1016/j.cognition.2004.07.001>
- Horton, W. S., & Keysar, B. (1996). When do speakers take into account common ground? *Cognition*, 59(1), 91–117. [https://doi.org/10.1016/0010-0277\(96\)81418-1](https://doi.org/10.1016/0010-0277(96)81418-1)
- Horton, W. S., & Spieler, D. H. (2007). Age-related differences in communication and audience design. *Psychology and Aging*, 22(2), 281–290. <https://doi.org/10.1037/0882-7974.22.2.281>
- Hu, J., Small, H., Kean, H., Takahashi, A., Zekelman, L., Kleinman, D., Ryan, E., Nieto-Castañón, A., Ferreira, V., & Fedorenko, E. (2022). Precision fMRI reveals that the language-selective network supports both phrase-structure building and lexical access during language production. *Cerebral Cortex*, 33(8), 4384–4404. <https://doi.org/10.1093/cercor/bhac350>
- Huang, Y. T., & Snedeker, J. (2009). Semantic meaning and pragmatic interpretation in 5-year-olds: Evidence from real-time spoken language comprehension. *Developmental Psychology*, 45(6), 1723–1739. <https://doi.org/10.1037/a0016704>

- Hupet, M., Chantraine, Y., & Nef, F. (1993). References in conversation between young and old normal adults. *Psychology and Aging*, 8(3), 339–346. <https://doi.org/10.1037/0882-7974.8.3.339>
- Ibarra, A., & Tanenhaus, M. K. (2016). The Flexibility of Conceptual Pacts: Referring Expressions Dynamically Shift to Accommodate New Conceptualizations. *Frontiers in Psychology*, 7. <https://doi.org/10.3389/fpsyg.2016.00561>
- Ji, A., Kojima, N., Rush, N., Suhr, A., Vong, W. K., Hawkins, R., & Artzi, Y. (2022). Abstract Visual Reasoning with Tangram Shapes. In Y. Goldberg, Z. Kozareva, & Y. Zhang (Eds.), *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing* (pp. 582–601). Association for Computational Linguistics. <https://doi.org/10.18653/v1/2022.emnlp-main.38>
- Kang, Y., Wang, T., & de Melo, G. (2020). Incorporating Pragmatic Reasoning Communication into Emergent Language. *Advances in Neural Information Processing Systems*, 33, 10348–10359.
- Kao, J. T. (2014). Formalizing the Pragmatics of Metaphor Understanding. *Proceedings of the Annual Meeting of the Cognitive Science Society*, 36.
- Keen, R. (2003). Representation of Objects and Events. *Current Directions in Psychological Science*, 12(3), 79–83. <https://doi.org/10.1111/1467-8721.01234>
- Kemp, C., Xu, Y., & Regier, T. (2018). Semantic Typology and Efficient Communication. *Annual Review of Linguistics*, 4(1), 109–128. <https://doi.org/10.1146/annurev-linguistics-011817-045406>
- Keysar, B., Barr, D. J., Balin, J. A., & Brauner, J. S. (2000). Taking Perspective in Conversation: The Role of Mutual Knowledge in Comprehension. *Psychological Science*, 11(1), 32–38. <https://doi.org/10.1111/1467-9280.00211>
- Keysar, B., Lin, S., & Barr, D. J. (2003). Limits on theory of mind use in adults. *Cognition*, 89(1), 25–41. [https://doi.org/10.1016/S0010-0277\(03\)00064-7](https://doi.org/10.1016/S0010-0277(03)00064-7)
- Kirby, S., Tamariz, M., Cornish, H., & Smith, K. (2015). Compression and communication in the cultural evolution of linguistic structure. *Cognition*, 141, 87–102. <https://doi.org/10.1016/j.cognition.2015.03.016>
- Knutsen, D., & Le Bigot, L. (2018). The influence of conceptual (mis)match on collaborative referring in dialogue. *Psychological Research*, 84(2), 514–527. <https://doi.org/10.1007/s00426-018-1060-1>

- Köymen, B., Schmerse, D., Lieven, E., & Tomasello, M. (2014). Young children create partner-specific referential pacts with peers. *Developmental Psychology*, 50(10), 2334–2342. <https://doi.org/10.1037/a0037837>
- Krauss, R. M., & Bricker, P. D. (1967). Effects of Transmission Delay and Access Delay on the Efficiency of Verbal Communication. *The Journal of the Acoustical Society of America*, 41(2), 286–292. <https://doi.org/10.1121/1.1910338>
- Krauss, R. M., & Glucksberg, S. (1969). The Development of Communication: Competence as a Function of Age. *Child Development*, 40(1), 255. <https://doi.org/10.2307/1127172>
- Krauss, R. M., & Glucksberg, S. (1977). Social and Nonsocial Speech. *Scientific American*, 236(2), 100–105. <https://doi.org/10.1038/scientificamerican0277-100>
- Krauss, R. M., & Weinheimer, S. (1964). Changes in reference phrases as a function of frequency of usage in social interaction: A preliminary study. *Psychonomic Science*, 1(1-12), 113–114. <https://doi.org/10.3758/BF03342817>
- Krauss, R. M., & Weinheimer, S. (1966). Concurrent feedback, confirmation, and the encoding of referents in verbal communication. *Journal of Personality and Social Psychology*, 4(3), 343–346. <https://doi.org/10.1037/h0023705>
- Kraut, R. E., Lewis, S. H., & Swezey, L. W. (1982). Listener responsiveness and the coordination of conversation. *Journal of Personality and Social Psychology*, 43(4), 718–731. <https://doi.org/10.1037/0022-3514.43.4.718>
- Le, H., Daryanto, T., Zhafransyah, F., Wijaya, D., Coppock, E., & Chin, S. (2022). *Referring Expressions with Rational Speech Act Framework: A Probabilistic Approach* (No. arXiv:2205.07795). arXiv. <https://doi.org/10.48550/arXiv.2205.07795>
- Leung, A., Yurovsky, D., & Hawkins, R. D. (2024). Parents spontaneously scaffold the formation of conversational pacts with their children. *Child Development*, 96(2), 546–561. <https://doi.org/10.1111/cdev.14186>
- Levy, R. (2008). A Noisy-Channel Model of Human Sentence Comprehension under Uncertain Input. *Proceedings of the 2008 Conference on Empirical Methods in Natural Language Processing*, 234–243.
- Lewis, D. (1969). *Convention: A philosophical study*. John Wiley & Sons.

- Lysander, K., & Horton, W. S. (2012). Conversational Grounding in Younger and Older Adults: The Effect of Partner Visibility and Referent Abstractness in Task-Oriented Dialogue. *Discourse Processes*, 49(1), 29–60. <https://doi.org/10.1080/0163853X.2011.625547>
- MacDonald, M. C. (2013). How language production shapes language form and comprehension. *Frontiers in Psychology*, 4. <https://doi.org/10.3389/fpsyg.2013.00226>
- MacMillan, J., Entin, E. E., & Serfaty, D. (2004). Communication overhead: The hidden cost of team cognition. In *Team cognition: Understanding the factors that drive process and performance*. (pp. 61–82). American Psychological Association. <https://doi.org/10.1037/10690-004>
- Mahowald, K., Fedorenko, E., Piantadosi, S. T., & Gibson, E. (2013). Info/information theory: Speakers choose shorter words in predictive contexts. *Cognition*, 126(2), 313–318. <https://doi.org/10.1016/j.cognition.2012.09.010>
- Mankewitz, J., Boyce, V., Waldon, B., Loukatou, G., Yu, D., Mu, J., Goodman, N. D., & Frank, M. C. (2021). *Multi-party referential communication in complex strategic games*. <https://doi.org/10.31234/osf.io/tfb3d>
- Matthews, D., Butcher, J., Lieven, E., & Tomasello, M. (2012). Two- and Four-Year-Olds Learn to Adapt Referring Expressions to Context: Effects of Distracters and Feedback on Referential Communication. *Topics in Cognitive Science*, 4(2), 184–210. <https://doi.org/10.1111/j.1756-8765.2012.01181.x>
- MATTHEWS, D., LIEVEN, E., THEAKSTON, A., & TOMASELLO, M. (2006). The effect of perceptual availability and prior discourse on young children's use of referring expressions. *Applied Psycholinguistics*, 27(3), 403–422. <https://doi.org/10.1017/S0142716406060334>
- Matthews, D., Lieven, E., & Tomasello, M. (2010). What's in a manner of speaking? Children's sensitivity to partner-specific referential precedents. *Developmental Psychology*, 46(4), 749–760. <https://doi.org/10.1037/a0019657>
- McKinley, G. L., Brown-Schmidt, S., & Benjamin, A. S. (2017). Memory for conversation and the development of common ground. *Memory & Cognition*, 45(8), 1281–1294. <https://doi.org/10.3758/s13421-017-0730-3>

- Metzing, C., & Brennan, S. (2003). When conceptual pacts are broken: Partner-specific effects on the comprehension of referring expressions. *Journal of Memory and Language*, 49(2), 201–213. [https://doi.org/10.1016/S0749-596X\(03\)00028-7](https://doi.org/10.1016/S0749-596X(03)00028-7)
- Misyak, J., Noguchi, T., & Chater, N. (2016). Instantaneous Conventions. *Psychological Science*, 27(12), 1550–1561. <https://doi.org/10.1177/0956797616661199>
- Mollica, F., Bacon, G., Zaslavsky, N., Xu, Y., Regier, T., & Kemp, C. (2021). The forms and meanings of grammatical markers support efficient communication. *Proceedings of the National Academy of Sciences*, 118(49), 135–138. <https://doi.org/10.1073/pnas.2025993118>
- Momma, S., Slevc, L. R., & Phillips, C. (2016). The timing of verb selection in Japanese sentence production. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 42(5), 813–824. <https://doi.org/10.1037/xlm0000195>
- Monroe, W., & Potts, C. (2015). Learning in the Rational Speech Acts Model. *arXiv:1510.06807 [Cs]*. <https://doi.org/10.48550/arxiv.1510.06807>
- Morisseau, T., Davies, C., & Matthews, D. (2013). How do 3- and 5-year-olds respond to under- and over-informative utterances? *Journal of Pragmatics*, 59, 26–39. <https://doi.org/10.1016/j.pragma.2013.03.007>
- Murfitt, T., & McAllister, J. (2001). The Effect of Production Variables in Monolog and Dialog on Comprehension by Novel Listeners. *Language and Speech*, 44(3), 325–350. <https://doi.org/10.1177/00238309010440030201>
- Muttenthaler, L., & Hebart, M. N. (2021). THINGSvision: A Python Toolbox for Streamlining the Extraction of Activations From Deep Neural Networks. *Frontiers in Neuroinformatics*, 15. <https://doi.org/10.3389/fninf.2021.679838>
- Nadig, A. S., & Sedivy, J. C. (2002). Evidence of Perspective-Taking Constraints in Children’s On-Line Reference Resolution. *Psychological Science*, 13(4), 329–336. <https://doi.org/10.1111/j.0956-7976.2002.00460.x>
- Nilsen, E., & Graham, S. (2009). The relations between children’s communicative perspective-taking and executive functioning. *Cognitive Psychology*, 58(2), 220–249. <https://doi.org/10.1016/j.cogpsych.2008.07.002>
- Noveck, I. A. (2001). When children are more logical than adults: Experimental investigations of scalar implicature. *Cognition*, 78(2), 165–188. [https://doi.org/10.1016/S0010-0285\(01\)00050-0](https://doi.org/10.1016/S0010-0285(01)00050-0)

- org/10.1016/S0010-0277(00)00114-1
- Ohmer, X., Franke, M., & König, P. (2022). Mutual Exclusivity in Pragmatic Agents. *Cognitive Science*, 46(1), e13069. <https://doi.org/10.1111/cogs.13069>
- Parisi, J. A., & Brungart, D. S. (2005). Evaluating communication effectiveness in team collaboration. *Ninth European Conference on Speech Communication and Technology (INTERSPEECH)*. <https://doi.org/10.21437/Interspeech.2005-459>
- Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Müller, A., Nothman, J., Louppe, G., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau, D., Brucher, M., Perrot, M., & Duchesnay, É. (2011). Scikit-learn: Machine Learning in Python. *arXiv*, 12, 2825–2830. <https://doi.org/10.48550/arxiv.1201.0490>
- Piantadosi, S. T. (2014). Zipf’s word frequency law in natural language: A critical review and future directions. *Psychonomic Bulletin & Review*, 21(5), 1112–1130. <https://doi.org/10.3758/s13423-014-0585-6>
- Piantadosi, S. T., Tily, H., & Gibson, E. (2012). The communicative function of ambiguity in language. *Cognition*, 122(3), 280–291. <https://doi.org/10.1016/j.cognition.2011.10.004>
- Pickering, M. J., & Ferreira, V. S. (2008). Structural priming: A critical review. *Psychological Bulletin*, 134(3), 427–459. <https://doi.org/10.1037/0033-2909.134.3.427>
- Pickering, M. J., & Garrod, S. (2004). Toward a mechanistic psychology of dialogue. *Behavioral and Brain Sciences*, 27(02). <https://doi.org/10.1017/S0140525X04000056>
- Plotkin, J. B., & Nowak, M. A. (2000). Language Evolution and Information Theory. *Journal of Theoretical Biology*, 205(1), 147–159. <https://doi.org/10.1006/jtbi.2000.2053>
- Potts, C., Lassiter, D., Levy, R., & Frank, M. C. (2015). Embedded Implicatures as Pragmatic Inferences under Compositional Lexical Uncertainty. *Journal of Semantics*, ffv012. <https://doi.org/10.1093/jos/ffv012>
- Qing, C., Goodman, N. D., & Lassiter, D. (2016). A rational speech-act model of projective content. *Proceedings of the Annual Meeting of the Cognitive Science*

Society, 38.

- Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., Sas-
try, G., Askell, A., Mishkin, P., Clark, J., Krueger, G., & Sutskever, I. (2021).
Learning Transferable Visual Models From Natural Language Supervision (No.
arXiv:2103.00020). arXiv. <https://doi.org/10.48550/arXiv.2103.00020>
- Radford, A., Kim, J. W., Xu, T., Brockman, G., McLeavey, C., & Sutskever, I.
(2022). *Robust Speech Recognition via Large-Scale Weak Supervision*. <https://doi.org/10.48550/arxiv.2212.04356>
- Rakoczy, H. (2022). Foundations of theory of mind and its development in early
childhood. *Nature Reviews Psychology*, 1(4), 223–235. <https://doi.org/10.1038/s44159-022-00037-z>
- Reimers, N., & Gurevych, I. (2019). *Sentence-BERT: Sentence Embeddings using
Siamese BERT-Networks* (No. arXiv:1908.10084). arXiv. <https://doi.org/10.48550/arXiv.1908.10084>
- Rogers, S. L., Fay, N., & Maybery, M. (2013). Audience Design through Social
Interaction during Group Discussion. *PLoS ONE*, 8(2), e57211. <https://doi.org/10.1371/journal.pone.0057211>
- Rubio-Fernandez, P., Mollica, F., & Jara-Ettinger, J. (2021). Speakers and listeners
exploit word order for communicative efficiency: A cross-linguistic investigation.
Journal of Experimental Psychology: General, 150(3), 583–594. <https://doi.org/10.1037/xge0000963>
- Rubio-Fernández, P. (2008). On the automaticity of egocentricity: A review of the
Egocentric Anchoring and Adjustment model of perspective taking. *UCL Working
Papers in Linguistics*, 20.
- San Juan, V., Khu, M., & Graham, S. A. (2015). A New Perspective on Children's
Communicative Perspective Taking: When and How Do Children Use Perspective
Inferences to Inform Their Comprehension of Spoken Language? *Child Develop-
ment Perspectives*, 9(4), 245–249. <https://doi.org/10.1111/cdep.12141>
- Schober, M. F., & Clark, H. H. (1989). Understanding by addressees and overhear-
ers. *Cognitive Psychology*, 21(2), 211–232. [https://doi.org/10.1016/0010-0285\(89\)90008-X](https://doi.org/10.1016/0010-0285(89)90008-X)
- Schuster, S., & Degen, J. (2020). I know what you're probably going to say: Listener

- adaptation to variable use of uncertainty expressions. *Cognition*, 203, 104285. <https://doi.org/10.1016/j.cognition.2020.104285>
- Seaman, C. B., & Basili, V. R. (1997). Communication and organization in software development: An empirical study. *IBM Systems Journal*, 36(4), 550–563. <https://doi.org/10.1147/sj.364.0550>
- Shain, C., Blank, I. A., van Schijndel, M., Schuler, W., & Fedorenko, E. (2020). fMRI reveals language-specific predictive coding during naturalistic sentence comprehension. *Neuropsychologia*, 138, 107307. <https://doi.org/10.1016/j.neuropsychologia.2019.107307>
- Smith, N. J., & Levy, R. (2013). The effect of word predictability on reading time is logarithmic. *Cognition*, 128(3), 302–319. <https://doi.org/10.1016/j.cognition.2013.02.013>
- Swaab, R. I., Galinsky, A. D., Medvec, V., & Diermeier, D. A. (2011). The Communication Orientation Model. *Personality and Social Psychology Review*, 16(1), 25–53. <https://doi.org/10.1177/1088868311417186>
- Tanenhaus, M. K., Spivey-Knowlton, M. J., Eberhard, K. M., & Sedivy, J. C. (1995). Integration of Visual and Linguistic Information in Spoken Language Comprehension. *Science*, 268(5217), 1632–1634. <https://doi.org/10.1126/science.7777863>
- Tannen, D. (2005). *Conversational style: Analyzing talk among friends*. Oxford University Press.
- Tolins, J., & Fox Tree, J. E. (2015). Overhearers Use Addressee Backchannels in Dialog Comprehension. *Cognitive Science*, 40(6), 1412–1434. <https://doi.org/10.1111/cogs.12278>
- Tolins, J., Zeamer, C., & Fox Tree, J. E. (2017). Overhearing Dialogues and Monologues: How Does Entrainment Lead to More Comprehensible Referring Expressions? *Discourse Processes*, 55(7), 545–565. <https://doi.org/10.1080/0163853X.2017.1279516>
- Tomasello, M. (2008). *Origins of human communication*. MIT Press.
- Tooley, K. M., & Traxler, M. J. (2010). Syntactic Priming Effects in Comprehension: A Critical Review. *Language and Linguistics Compass*, 4(10), 925–937. <https://doi.org/10.1111/j.1749-818X.2010.00249.x>

- Traum, D. (2004). Issues in Multiparty Dialogues. In G. Goos, J. Hartmanis, J. van Leeuwen, & F. Dignum (Eds.), *Advances in Agent Communication* (Vol. 2922, pp. 201–211). Springer Berlin Heidelberg. https://doi.org/10.1007/978-3-540-24608-4_12
- Turan-Küçük, E. N., & Kibbe, M. M. (2024). Three-year-olds’ ability to plan for mutually exclusive future possibilities is limited primarily by their representations of possible plans, not possible events. *Cognition*, 244, 105712. <https://doi.org/10.1016/j.cognition.2023.105712>
- Turner, C., & Knutsen, D. (2021). Audience Design in Collaborative Dialogue between Teachers and Students. *Discourse Processes*, 58(8), 703–725. <https://doi.org/10.1080/0163853X.2021.1904768>
- van Dongen, N., van Bork, R., Finnemann, A., Haslbeck, J. M. B., van der Maas, H. L. J., Robinaugh, D. J., de Ron, J., Sprenger, J., & Borsboom, D. (2025). Productive explanation: A framework for evaluating explanations in psychological science. *Psychological Review*, 132(2), 311–329. <https://doi.org/10.1037/rev0000479>
- Vanlangendonck, F., Takashima, A., Willems, R. M., & Hagoort, P. (2018). Distinguishable memory retrieval networks for collaboratively and non-collaboratively learned information. *Neuropsychologia*, 111, 123–132. <https://doi.org/10.1016/j.neuropsychologia.2017.12.008>
- Weber, R. A., & Camerer, C. F. (2003). Cultural Conflict and Merger Failure: An Experimental Approach. *Management Science*, 49(4), 400–415. <https://doi.org/10.1287/mnsc.49.4.400.14430>
- White, J., Mu, J., & Goodman, N. D. (2020). Learning to refer informatively by amortizing pragmatic reasoning. *arXiv:2006.00418 [Cs]*. <https://doi.org/10.48550/arxiv.2006.00418>
- Wilcox, E., Vani, P., & Levy, R. (2021). A Targeted Assessment of Incremental Processing in Neural Language Models and Humans. *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, 939–952. <https://doi.org/10.18653/v1/2021.acl-long.76>
- Wilkes-Gibbs, D., & Clark, H. H. (1992). Coordinating beliefs in conversation. *Journal of Memory and Language*, 31(2), 183–194. <https://doi.org/10.1016/0749->

- 596X(92)90010-U
- Wittgenstein, L. (1953). *Philosophical investigations*. Wiley-Blackwell.
- Yarkoni, T., & Westfall, J. (2017). Choosing Prediction Over Explanation in Psychology: Lessons From Machine Learning. *Perspectives on Psychological Science*, 12(6), 1100–1122. <https://doi.org/10.1177/1745691617693393>
- Yoon, E. J., Tessler, M. H., Goodman, N. D., & Frank, M. C. (2020). Polite Speech Emerges From Competing Social Goals. *Open Mind*, 4, 71–87. https://doi.org/10.1162/opmi_a_00035
- Yoon, S. O., & Brown-Schmidt, S. (2014). Adjusting conceptual pacts in three-party conversation. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 40(4), 919–937. <https://doi.org/10.1037/a0036161>
- Yoon, S. O., & Brown-Schmidt, S. (2017). Aim Low: Mechanisms of Audience Design in Multiparty Conversation. *Discourse Processes*, 55(7), 566–592. <https://doi.org/10.1080/0163853X.2017.1286225>
- Yoon, S. O., & Brown-Schmidt, S. (2019a). Audience Design in Multiparty Conversation. *Cognitive Science*, 43(8), e12774. <https://doi.org/10.1111/cogs.12774>
- Yoon, S. O., & Brown-Schmidt, S. (2019b). Contextual Integration in Multiparty Audience Design. *Cognitive Science*, 43(12), e12807. <https://doi.org/10.1111/cogs.12807>
- Yoon, S. O., Duff, M. C., & Brown-Schmidt, S. (2017). Learning and using knowledge about what other people do and don't know despite amnesia. *Cortex*, 94, 164–175. <https://doi.org/10.1016/j.cortex.2017.06.020>
- Yurovsky, D., & Frank, M. C. (2017). Beyond naïve cue combination: Salience and social cues in early word learning. *Developmental Science*, 20(2). <https://doi.org/10.1111/desc.12349>
- Zack, M. H. (1993). Interactivity and Communication Mode Choice in Ongoing Management Groups. *Information Systems Research*, 4(3), 207–239. <https://doi.org/10.1287/isre.4.3.207>
- Zaslavsky, N., Hu, J., & Levy, R. P. (2020). *A Rate-Distortion view of human pragmatic reasoning* (No. arXiv:2005.06641). arXiv. <https://doi.org/10.48550/arXiv.2005.06641>
- Zaslavsky, N., Kemp, C., Regier, T., & Tishby, N. (2018). Efficient compression in

- color naming and its evolution. *Proceedings of the National Academy of Sciences*, 115(31), 7937–7942. <https://doi.org/10.1073/pnas.1800521115>
- Zipf, G. (1949). *Human behavior and the principle of least effort: An introduction to human ecology*.